

Distributed Search Engine Backend

□ Table of Contents

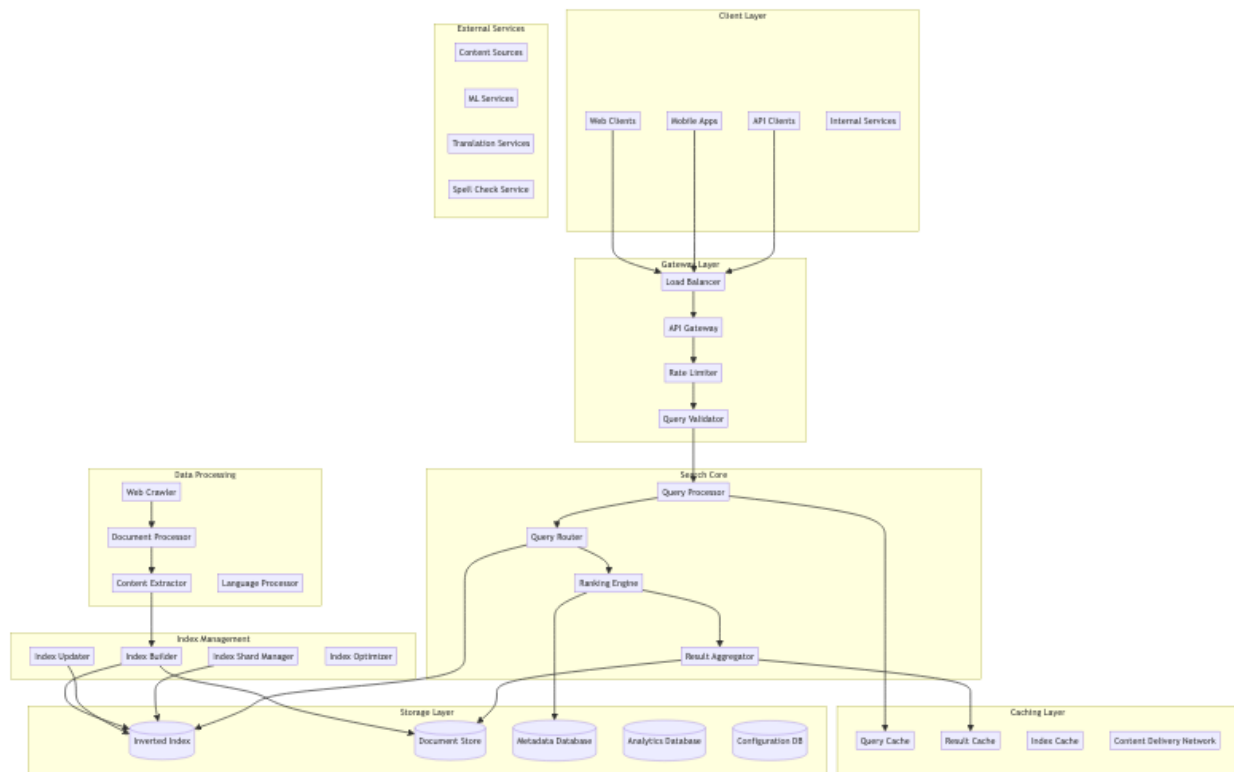
- Distributed Search Engine Backend
 - High-Level Design (HLD)
 - * System Architecture Overview
 - * Search Processing Flow
 - Low-Level Design (LLD)
 - * Query Processing Engine
 - * Distributed Index Architecture
 - * Real-time Ranking System
 - Core Algorithms
 - * 1. Distributed Query Processing Algorithm
 - * 2. Advanced Ranking Algorithm
 - * 3. Real-time Index Update Algorithm
 - * 4. Auto-complete and Suggestion Algorithm
 - * 5. Search Analytics and Performance Monitoring
 - Performance Optimizations
 - * Index Optimization
 - * Caching Strategy
 - * Database Optimization
 - Security Considerations
 - * Search Security Framework
 - Testing Strategy
 - * Performance Testing
 - * Search Quality Testing
 - Trade-offs and Considerations
 - * Relevance vs Performance
 - * Storage vs Speed
 - * Scalability vs Consistency
-

High-Level Design (HLD)

□ [Back to Top](#)




System Architecture Overview

□ [Back to Top](#)



Search Processing Flow

□ [Back to Top](#)

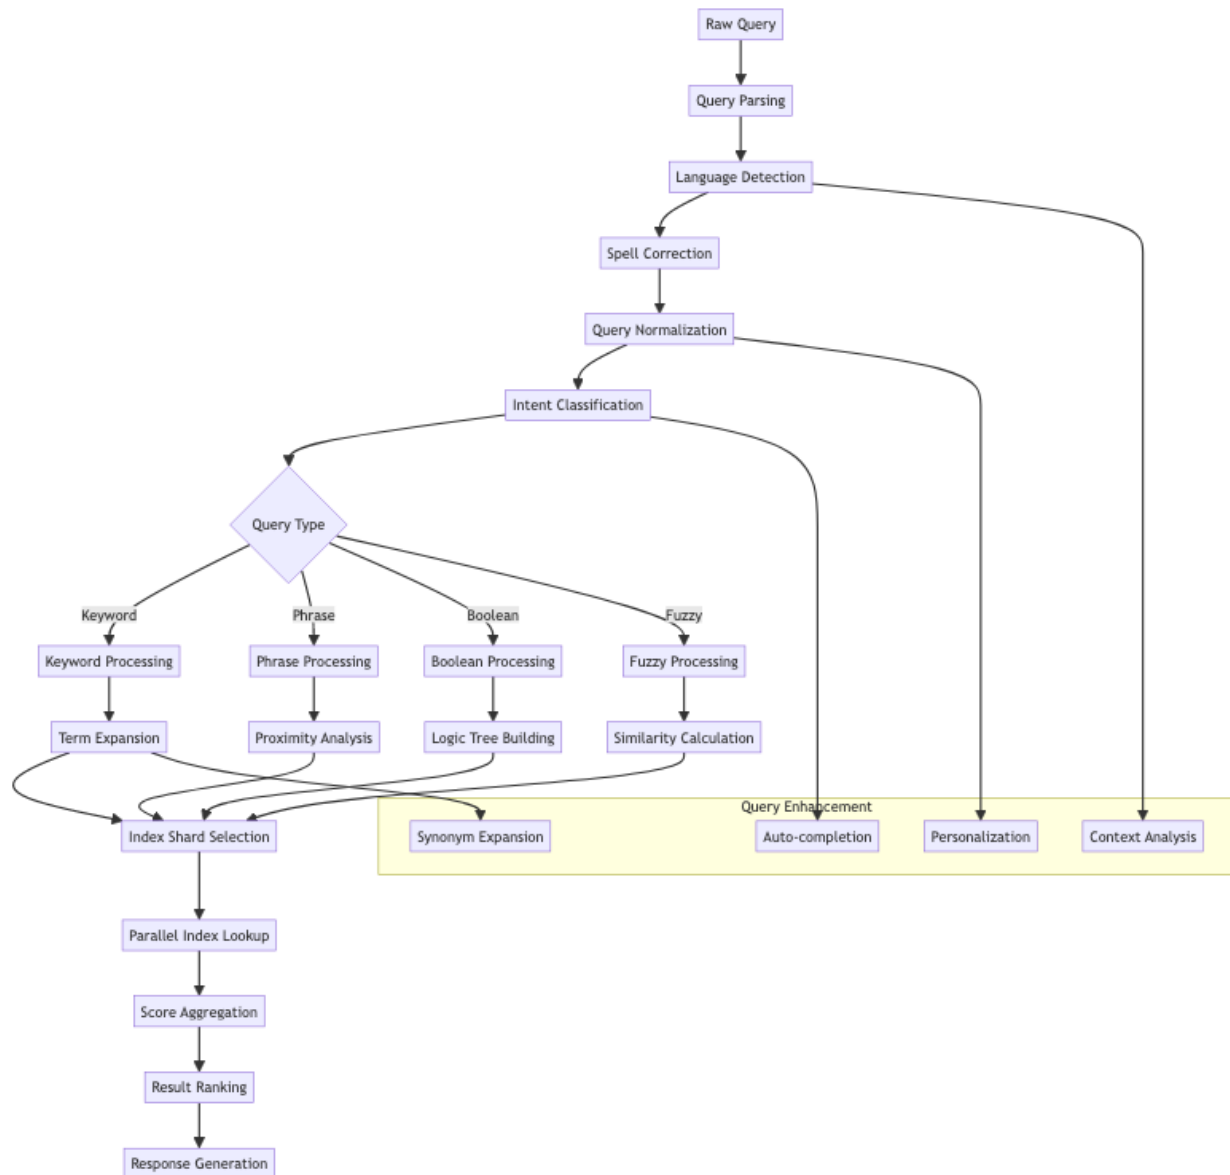


Low-Level Design (LLD)

□ [Back to Top](#)

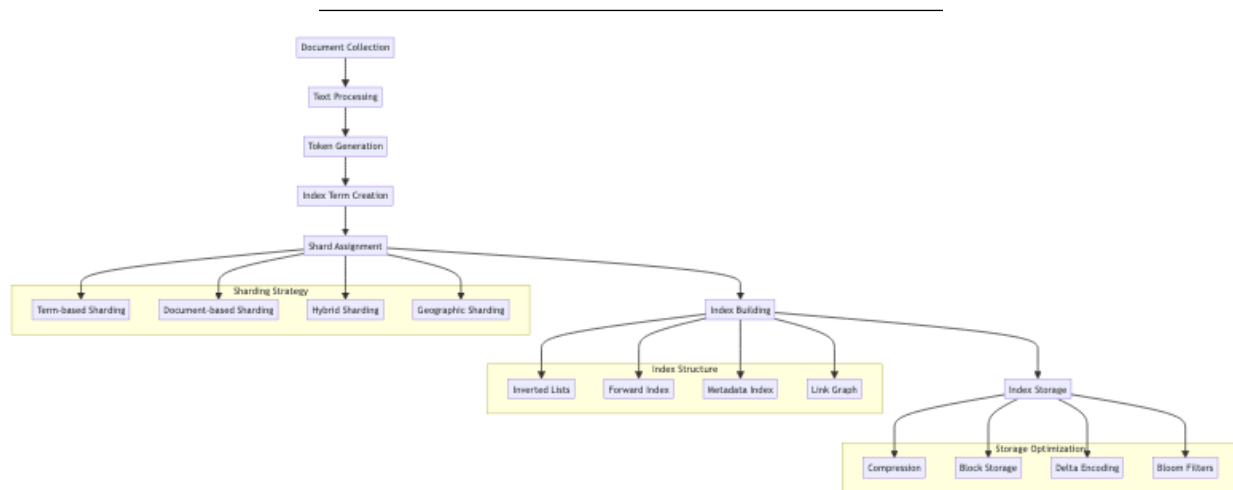
Query Processing Engine

□ [Back to Top](#)



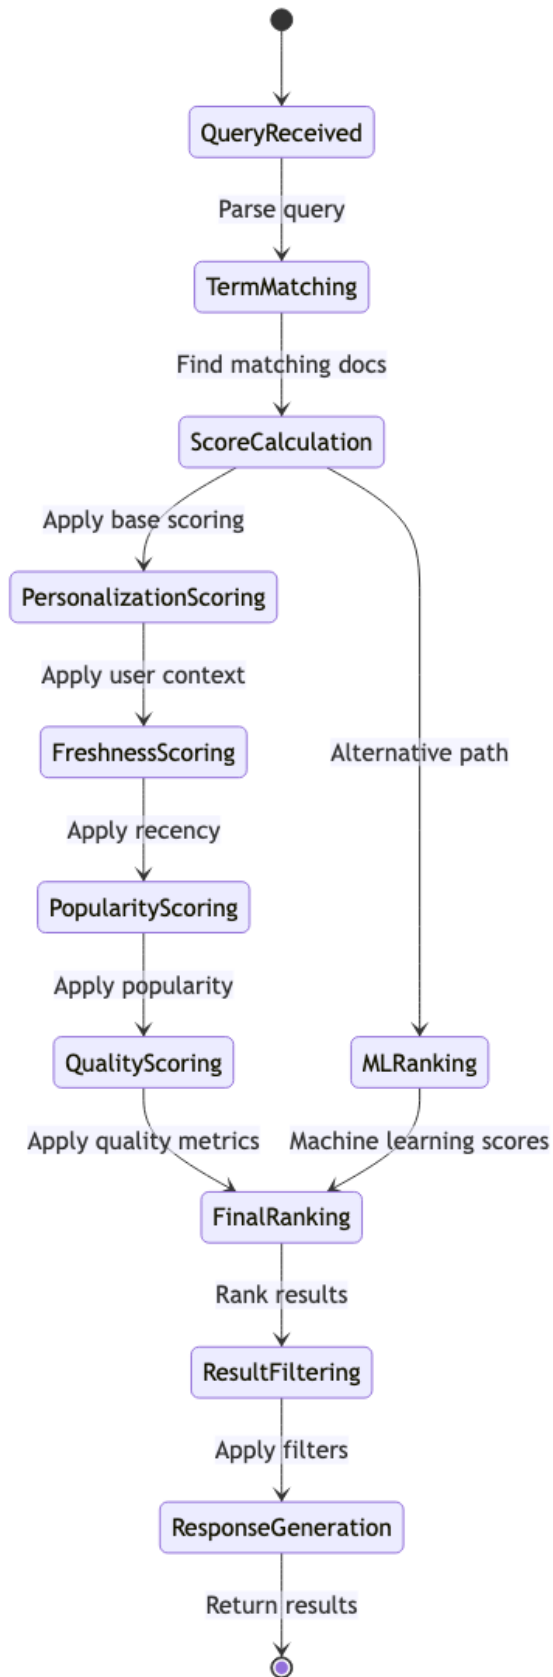
Distributed Index Architecture

[□ Back to Top](#)



Real-time Ranking System

[□ Back to Top](#)



Core Algorithms

□ [Back to Top](#)

1. Distributed Query Processing Algorithm

□ [Back to Top](#)

Purpose: Process search queries across distributed index shards with optimal performance and relevance.

Query Distribution Strategy:

```
SearchConfig = {
  shardingStrategy: 'hybrid',      // term-based, document-based, or hybrid
  replicationFactor: 3,           // Number of replicas per shard
  maxShardsPerQuery: 50,          // Limit concurrent shard queries
  queryTimeout: 5000,             // 5 seconds max query time
  minScoreThreshold: 0.1,         // Minimum relevance score
  maxResultsPerShard: 1000        // Results per shard before aggregation
}

function processDistributedQuery(query, searchParams):
  // Parse and analyze query
  processedQuery = analyzeQuery(query, searchParams)

  // Determine relevant shards
  relevantShards = selectShards(processedQuery, SearchConfig.shardingStrategy)

  if relevantShards.length > SearchConfig.maxShardsPerQuery:
    // Prioritize shards by relevance
    relevantShards = prioritizeShards(relevantShards, processedQuery)
    relevantShards = relevantShards.slice(0, SearchConfig.maxShardsPerQuery)

  // Execute parallel queries across shards
  shardPromises = relevantShards.map(shard =>
    executeShardQuery(shard, processedQuery, searchParams)
  )

  // Collect results with timeout
  shardResults = await Promise.allSettled(shardPromises, SearchConfig.queryTimeout)

  // Filter successful results
```

```

validResults = shardResults
  .filter(result => result.status === 'fulfilled')
  .map(result => result.value)

// Aggregate and rank results
aggregatedResults = aggregateShardResults(validResults, processedQuery)

// Apply global ranking
finalResults = applyGlobalRanking(aggregatedResults, searchParams)

return {
  results: finalResults,
  totalHits: calculateTotalHits(validResults),
  queryTime: Date.now() - processedQuery.startTime,
  shardsQueried: validResults.length,
  partialResults: shardResults.some(r => r.status === 'rejected')
}

function selectShards(processedQuery, strategy):
  switch strategy:
    case 'term-based':
      return selectShardsByTerms(processedQuery.terms)
    case 'document-based':
      return selectShardsByCategories(processedQuery.categories)
    case 'hybrid':
      termShards = selectShardsByTerms(processedQuery.terms)
      categoryShards = selectShardsByCategories(processedQuery.categories)
      return mergeShardSelections(termShards, categoryShards)
    default:
      return getAllShards()

```

Parallel Shard Query Execution:

```

function executeShardQuery(shard, query, params):
  return new Promise((resolve, reject) => {
    const timeout = setTimeout(() => {
      reject(new Error('Shard query timeout'))
    }, SearchConfig.queryTimeout)

    try:
      // Connect to shard
      connection = getShardConnection(shard.id)

      // Execute query on shard
      shardResults = connection.search({
        query: query.optimizedQuery,

```



```

        filters: params.filters,
        sorting: params.sorting,
        maxResults: SearchConfig.maxResultsPerShard,
        includeMetadata: true
    })

    clearTimeout(timeout)

    // Apply shard-level scoring
    scoredResults = applyShardScoring(shardResults, query, shard)

    resolve({
        shardId: shard.id,
        results: scoredResults,
        metadata: {
            totalHits: shardResults.totalHits,
            queryTime: shardResults.queryTime,
            shardLoad: shard.currentLoad
        }
    })

    catch error:
        clearTimeout(timeout)
        reject(error)
    })

```

2. Advanced Ranking Algorithm

□ [Back to Top](#)

Purpose: Rank search results using multiple relevance signals and machine learning models.

Multi-Factor Ranking System:

```

RankingConfig = {
    textRelevanceWeight: 0.4,           // TF-IDF and BM25 scores
    popularityWeight: 0.2,              // Click-through rates, links
    freshnessWeight: 0.15,              // Recency of content
    personalizedWeight: 0.15,          // User preferences and history
    qualityWeight: 0.1,                 // Content quality signals

    bm25Parameters: {
        k1: 1.2,                        // Term frequency saturation
        b: 0.75                         // Field length normalization
    }
}

```

```

    },

    freshnessDecay: 0.1,           // Daily decay factor for old content
    minPopularityScore: 0.01       // Minimum popularity threshold
}

function calculateRelevanceScore(document, query, userContext):
    score = 0

    // Text relevance (TF-IDF + BM25)
    textScore = calculateTextRelevance(document, query)
    score += textScore * RankingConfig.textRelevanceWeight

    // Popularity signals
    popularityScore = calculatePopularityScore(document)
    score += popularityScore * RankingConfig.popularityWeight

    // Freshness score
    freshnessScore = calculateFreshnessScore(document.publishedDate)
    score += freshnessScore * RankingConfig.freshnessWeight

    // Personalization
    personalizedScore = calculatePersonalizedScore(document, userContext)
    score += personalizedScore * RankingConfig.personalizedWeight

    // Quality signals
    qualityScore = calculateQualityScore(document)
    score += qualityScore * RankingConfig.qualityWeight

    return Math.max(score, 0)

function calculateTextRelevance(document, query):
    bm25Score = 0

    for term in query.terms:
        // Calculate BM25 score for each term
        tf = getTermFrequency(term, document)
        df = getDocumentFrequency(term)
        idf = Math.log((totalDocuments - df + 0.5) / (df + 0.5))

        // BM25 formula
        termScore = idf * (tf * (RankingConfig.bm25Parameters.k1 + 1)) /
            (tf + RankingConfig.bm25Parameters.k1 *
            (1 - RankingConfig.bm25Parameters.b +
            RankingConfig.bm25Parameters.b * (document.length / avgDocumentLength)))

```

```

    bm25Score += termScore

    // Apply field boosting
    titleBoost = calculateFieldBoost(document.title, query.terms, 2.0)
    headingBoost = calculateFieldBoost(document.headings, query.terms, 1.5)

    return bm25Score + titleBoost + headingBoost
}

function calculatePopularityScore(document):
    // Combine multiple popularity signals
    clickThroughRate = document.analytics.ctr || 0
    linkCount = Math.log(document.inboundLinks + 1)
    socialSignals = Math.log(document.socialShares + 1)

    // Normalize and combine
    normalizedCTR = Math.min(clickThroughRate / 0.1, 1.0) // Cap at 10% CTR
    normalizedLinks = Math.min(linkCount / 10, 1.0) // Cap at 10 log scale
    normalizedSocial = Math.min(socialSignals / 5, 1.0) // Cap at 5 log scale

    return (normalizedCTR * 0.5 + normalizedLinks * 0.3 + normalizedSocial * 0.2)

```

Machine Learning Ranking Integration:

```

function applyMLRanking(results, query, userContext):
    // Prepare features for ML model
    features = results.map(result => extractFeatures(result, query, userContext))

    // Get ML model predictions
    mlScores = mlRankingModel.predict(features)

    // Combine with traditional ranking
    combinedResults = results.map((result, index) => ({
        ...result,
        mlScore: mlScores[index],
        combinedScore: (result.relevanceScore * 0.7) + (mlScores[index] * 0.3)
    }))

    // Re-sort by combined score
    return combinedResults.sort((a, b) => b.combinedScore - a.combinedScore)

function extractFeatures(document, query, userContext):
    return {
        // Text features
        titleTermMatches: countTermMatches(document.title, query.terms),
        bodyTermMatches: countTermMatches(document.body, query.terms),
    }

```

```

    exactPhraseMatches: countPhraseMatches(document, query.phrases),

    // Document features
    documentLength: document.body.length,
    readabilityScore: document.metrics.readability,
    authorityScore: document.metrics.authority,

    // Popularity features
    clickThroughRate: document.analytics.ctr,
    bounceRate: document.analytics.bounceRate,
    timeOnPage: document.analytics.avgTimeOnPage,

    // Freshness features
    daysSincePublished: (Date.now() - document.publishedDate) / (24 * 60 * 60 * 1000),
    daysSinceUpdated: (Date.now() - document.updatedDate) / (24 * 60 * 60 * 1000),

    // User context features
    userLanguage: userContext.language,
    userLocation: userContext.location,
    userInterests: userContext.interests,
    previousInteractions: userContext.previousClicks.includes(document.id)
}

```

3. Real-time Index Update Algorithm

□ [Back to Top](#)

Purpose: Maintain search index consistency with real-time document updates and minimal search disruption.

Incremental Index Update:

```

IndexUpdateConfig = {
  batchSize: 1000,           // Documents per batch
  maxBatchWaitTime: 30000,   // 30 seconds max wait
  updateStrategy: 'incremental', // 'incremental' or 'rebuild'
  mergeFactor: 10,          // Segments to merge at once
  maxSegments: 100,         // Max segments before forced merge

  priorities: {
    add: 1,      // New documents
    update: 2,   // Document changes
    delete: 3    // Document removals
  }
}

```

```

function processDocumentUpdate(document, operation, priority):
    // Create update operation
    updateOperation = {
        documentId: document.id,
        operation: operation, // 'add', 'update', 'delete'
        priority: priority,
        timestamp: Date.now(),
        shardId: determineShardForDocument(document),
        indexTerms: operation !== 'delete' ? extractIndexTerms(document) : null
    }

    // Add to appropriate priority queue
    addToUpdateQueue(updateOperation)

    // Trigger batch processing if needed
    if shouldProcessBatch():
        processBatchUpdates()

function processBatchUpdates():
    // Get batch of updates sorted by priority and shard
    updateBatch = getNextUpdateBatch(IndexUpdateConfig.batchSize)

    if updateBatch.length === 0:
        return

    // Group by shard for efficient processing
    shardGroups = groupUpdatesByShard(updateBatch)

    // Process each shard group
    for shardId in Object.keys(shardGroups):
        processShardUpdates(shardId, shardGroups[shardId])

function processShardUpdates(shardId, updates):
    shard = getIndexShard(shardId)

    // Apply updates in transaction
    transaction = shard.beginTransaction()

    try:
        for update in updates:
            switch update.operation:
                case 'add':
                    addDocumentToIndex(shard, update.documentId, update.indexTerms)
                    break

```

```

        case 'update':
            updateDocumentInIndex(shard, update.documentId, update.indexTerms)
            break
        case 'delete':
            removeDocumentFromIndex(shard, update.documentId)
            break

    // Commit transaction
    transaction.commit()

    // Update metadata
    updateShardMetadata(shardId, updates.length)

    // Check if merge is needed
    if shouldMergeSegments(shard):
        scheduleSegmentMerge(shardId)

catch error:
    transaction.rollback()
    // Re-queue failed updates
    requeueUpdates(updates)
    throw error

```

Index Segment Merging:

```

function mergeIndexSegments(shardId):
    shard = getIndexShard(shardId)
    segments = shard.getSegments()

    // Select segments to merge based on size and age
    segmentsToMerge = selectMergeSegments(segments)

    if segmentsToMerge.length < 2:
        return // Nothing to merge

    // Create new merged segment
    mergedSegment = createNewSegment(shardId)

    // Merge inverted lists
    termMerger = new TermMerger()

    for segment in segmentsToMerge:
        for term in segment.getTerms():
            postingList = segment.getPostingList(term)
            termMerger.addPostingList(term, postingList)

```

```

// Write merged data
for term in termMerger.getTerms():
    mergedPostingList = termMerger.getMergedPostingList(term)
    mergedSegment.writePostingList(term, mergedPostingList)

// Atomically replace old segments
shard.replaceSegments(segmentsToMerge, mergedSegment)

// Clean up old segment files
cleanupSegments(segmentsToMerge)

// Update search caches
invalidateSearchCaches(shardId)

```

4. Auto-complete and Suggestion Algorithm

□ [Back to Top](#)

Purpose: Provide real-time query suggestions and auto-completion with high relevance and performance.

Trie-based Auto-completion:

```

AutocompleteConfig = {
    maxSuggestions: 10,
    minQueryLength: 2,
    suggestionSources: ['queries', 'documents', 'entities'],
    popularityWeight: 0.6,
    recencyWeight: 0.3,
    personalizedWeight: 0.1
}

```

```

class AutocompleteTrie:
    constructor():
        this.root = new TrieNode()
        this.queryFrequency = new Map()
        this.recentQueries = new CircularBuffer(10000)

    function insertQuery(query, frequency, timestamp):
        node = this.root

        // Insert query character by character
        for char in query.toLowerCase():
            if not node.children.has(char):
                node.children.set(char, new TrieNode())

```

```

    node = node.children.get(char)

    // Update node statistics
    node.passthrough += frequency

    // Mark end of query
    node.isEndOfQuery = true
    node.queries.add({
        text: query,
        frequency: frequency,
        timestamp: timestamp
    })

    // Update global frequency tracking
    this.queryFrequency.set(query, frequency)
    this.recentQueries.add({ query, timestamp })

function getSuggestions(prefix, userContext):
    if prefix.length < AutocompleteConfig.minQueryLength:
        return []

    // Navigate to prefix node
    node = this.root
    for char in prefix.toLowerCase():
        if not node.children.has(char):
            return [] // No suggestions for this prefix
        node = node.children.get(char)

    // Collect all completions from this node
    completions = []
    this.collectCompletions(node, prefix, completions, userContext)

    // Score and rank suggestions
    scoredSuggestions = completions.map(completion => ({
        ...completion,
        score: this.calculateSuggestionScore(completion, userContext)
    }))

    // Sort by score and return top suggestions
    return scoredSuggestions
        .sort((a, b) => b.score - a.score)
        .slice(0, AutocompleteConfig.maxSuggestions)

function calculateSuggestionScore(suggestion, userContext):
    score = 0

```



```

// Popularity component
popularityScore = Math.log(suggestion.frequency + 1) / Math.log(maxFrequency + 1)
score += popularityScore * AutocompleteConfig.popularityWeight

// Recency component
daysSinceUsed = (Date.now() - suggestion.lastUsed) / (24 * 60 * 60 * 1000)
recencyScore = Math.exp(-daysSinceUsed / 30) // 30-day half-life
score += recencyScore * AutocompleteConfig.recencyWeight

// Personalization component
if userContext.searchHistory:
    personalizedScore = calculatePersonalizedScore(suggestion, userContext)
    score += personalizedScore * AutocompleteConfig.personalizedWeight

return score

```

Contextual Query Suggestions:

```

function generateContextualSuggestions(query, searchResults, userContext):
    suggestions = []

    // Extract entities from current results
    entities = extractEntitiesFromResults(searchResults)

    // Generate entity-based expansions
    for entity in entities:
        relatedQueries = getRelatedQueries(entity, userContext.language)
        suggestions.push(...relatedQueries)

    // Generate topic-based suggestions
    topics = extractTopicsFromResults(searchResults)
    for topic in topics:
        topicSuggestions = getTopicSuggestions(topic, query)
        suggestions.push(...topicSuggestions)

    // Generate corrective suggestions
    if searchResults.length < 5:
        corrections = generateSpellCorrections(query)
        reformulations = generateQueryReformulations(query)
        suggestions.push(...corrections, ...reformulations)

    // Score and rank suggestions
    return rankSuggestions(suggestions, query, userContext)

function generateQueryReformulations(originalQuery):

```

```

reformulations = []

// Synonym-based reformulations
synonyms = getSynonyms(originalQuery.terms)
for synonym in synonyms:
    reformulatedQuery = replaceTermWithSynonym(originalQuery, synonym)
    reformulations.push({
        type: 'synonym',
        original: originalQuery.text,
        suggestion: reformulatedQuery,
        confidence: synonym.confidence
    })

// Broader/narrower reformulations
broader = generateBroaderQueries(originalQuery)
narrower = generateNarrowerQueries(originalQuery)

reformulations.push(...broader, ...narrower)

return reformulations

```

5. Search Analytics and Performance Monitoring

□ [Back to Top](#)

Purpose: Track search performance, user behavior, and system health for continuous optimization.

Real-time Analytics Collection:

```

AnalyticsConfig = {
    trackingEvents: [
        'query_submitted',
        'result_clicked',
        'result_viewed',
        'query_abandoned',
        'autocomplete_used',
        'filter_applied',
        'sort_changed'
    ],
    samplingRate: 0.1,           // Sample 10% of queries for detailed analysis
    aggregationWindows: [60, 300, 3600, 86400], // 1m, 5m, 1h, 1d
    retentionPeriod: 90         // Keep detailed data for 90 days
}

```

```

function trackSearchEvent(eventType, eventData, userContext):
    event = {
        type: eventType,
        timestamp: Date.now(),
        sessionId: userContext.sessionId,
        userId: userContext.userId,
        query: eventData.query,
        results: eventData.results,
        metadata: eventData.metadata
    }

    // Store raw event
    storeSearchEvent(event)

    // Update real-time metrics
    updateRealTimeMetrics(event)

    // Sample for detailed analysis
    if Math.random() < AnalyticsConfig.samplingRate:
        performDetailedAnalysis(event)

function calculateSearchMetrics(timeWindow):
    events = getSearchEvents(timeWindow)

    metrics = {
        totalQueries: 0,
        uniqueUsers: new Set(),
        clickThroughRate: 0,
        averageResultsPerQuery: 0,
        averageQueryTime: 0,
        zeroResultQueries: 0,
        popularQueries: new Map(),
        performanceMetrics: {
            p50QueryTime: 0,
            p95QueryTime: 0,
            p99QueryTime: 0
        }
    }

    queryTimes = []
    clickEvents = 0
    totalResults = 0

    for event in events:
        switch event.type:

```

```

    case 'query_submitted':
      metrics.totalQueries++
      metrics.uniqueUsers.add(event.userId)
      queryTimes.push(event.metadata.queryTime)

      resultCount = event.results ? event.results.length : 0
      totalResults += resultCount

      if resultCount === 0:
        metrics.zeroResultQueries++

      // Track popular queries
      queryText = event.query.toLowerCase()
      currentCount = metrics.popularQueries.get(queryText) || 0
      metrics.popularQueries.set(queryText, currentCount + 1)
      break

    case 'result_clicked':
      clickEvents++
      break

  // Calculate derived metrics
  if metrics.totalQueries > 0:
    metrics.clickThroughRate = clickEvents / metrics.totalQueries
    metrics.averageResultsPerQuery = totalResults / metrics.totalQueries
    metrics.averageQueryTime = queryTimes.reduce((a, b) => a + b, 0) / queryTimes.length

  // Calculate percentiles
  sortedTimes = queryTimes.sort((a, b) => a - b)
  metrics.performanceMetrics.p50QueryTime = getPercentile(sortedTimes, 0.5)
  metrics.performanceMetrics.p95QueryTime = getPercentile(sortedTimes, 0.95)
  metrics.performanceMetrics.p99QueryTime = getPercentile(sortedTimes, 0.99)

  return metrics

```

Performance Optimizations

□ [Back to Top](#)

Index Optimization

□ [Back to Top](#)

Compression Strategies:

```
CompressionConfig = {  
  postingListCompression: 'variable_byte', // Variable-byte encoding  
  documentCompression: 'lz4',             // Fast compression for documents  
  blockSize: 128,                         // Compression block size  
  enableDeltaEncoding: true               // Delta encoding for doc IDs  
}
```

Caching Strategy

[□ Back to Top](#)

Multi-Layer Caching: - Query result caching (Redis, 15-minute TTL) - Index segment caching (Application memory) - Autocomplete caching (Edge CDN) - Analytics data caching (Time-series database)

Database Optimization

[□ Back to Top](#)

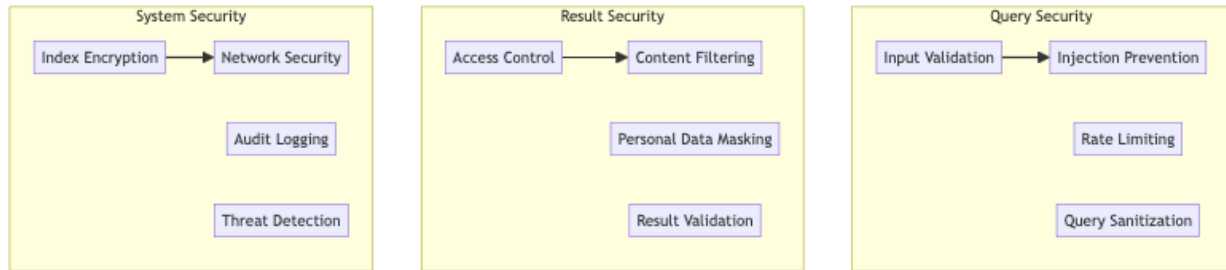
Index Sharding Strategy: - Term-based sharding for text search - Geographic sharding for location-based queries - Time-based sharding for temporal data - Hybrid approach for complex queries

Security Considerations

[□ Back to Top](#)

Search Security Framework

[□ Back to Top](#)



Testing Strategy

[□ Back to Top](#)

Performance Testing

[□ Back to Top](#)

Load Testing Scenarios: - Peak traffic simulation (10x normal load) - Concurrent user testing (100K+ simultaneous queries) - Index update stress testing - Cache invalidation testing

Search Quality Testing

[□ Back to Top](#)

Relevance Testing: - Human relevance evaluation - A/B testing for ranking changes - Click-through rate analysis - Search result diversity measurement

Trade-offs and Considerations

[□ Back to Top](#)

Relevance vs Performance

[□ Back to Top](#)

- **Complex ranking:** Accuracy vs query response time
- **Real-time updates:** Freshness vs index consistency

- **Personalization:** Relevance vs privacy concerns
- **Comprehensive results:** Completeness vs response speed

Storage vs Speed

□ [Back to Top](#)

- **Index compression:** Storage efficiency vs decompression overhead
- **Caching strategy:** Memory usage vs cache hit rates
- **Replication:** Availability vs storage costs
- **Sharding granularity:** Load distribution vs query complexity

Scalability vs Consistency

□ [Back to Top](#)

- **Distributed indexing:** Scalability vs update consistency
- **Cross-shard queries:** Completeness vs performance
- **Index synchronization:** Consistency vs availability
- **Cache coherence:** Performance vs data freshness

This distributed search engine backend provides a comprehensive foundation for large-scale search with features like intelligent query processing, advanced ranking algorithms, real-time indexing, and robust performance monitoring while maintaining high relevance, scalability, and security standards.