

Distributed Notification System

□ Table of Contents

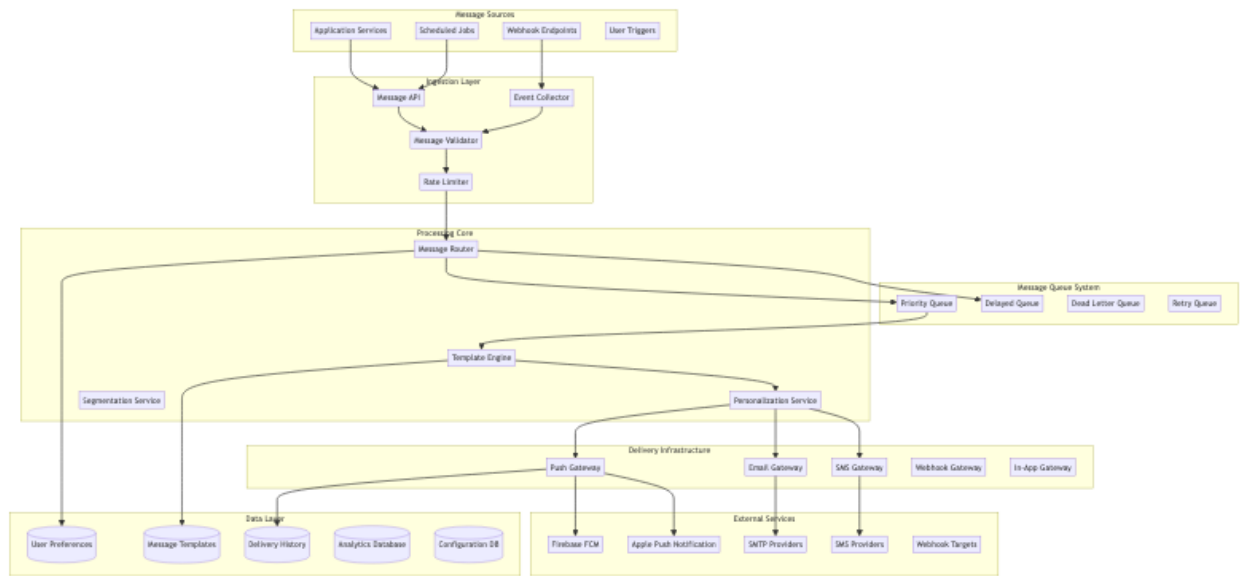
- Distributed Notification System
 - High-Level Design (HLD)
 - * System Architecture Overview
 - * Notification Flow Architecture
 - Low-Level Design (LLD)
 - * Message Processing Pipeline
 - * Smart Routing Engine
 - * Delivery Reliability System
 - Core Algorithms
 - * 1. Intelligent Message Routing Algorithm
 - * 2. Frequency Capping and Rate Limiting Algorithm
 - * 3. Template Engine and Personalization Algorithm
 - * 4. Delivery Reliability and Retry Logic
 - * 5. Analytics and Performance Monitoring
 - Performance Optimizations
 - * Message Queue Optimization
 - * Database Optimization
 - * Caching Strategy
 - Security Considerations
 - * Privacy and Data Protection
 - Testing Strategy
 - * Load Testing Scenarios
 - * Reliability Testing
 - Trade-offs and Considerations
 - * Delivery Speed vs Reliability
 - * Personalization vs Privacy
 - * Scalability vs Cost
-

High-Level Design (HLD)

□ [Back to Top](#)

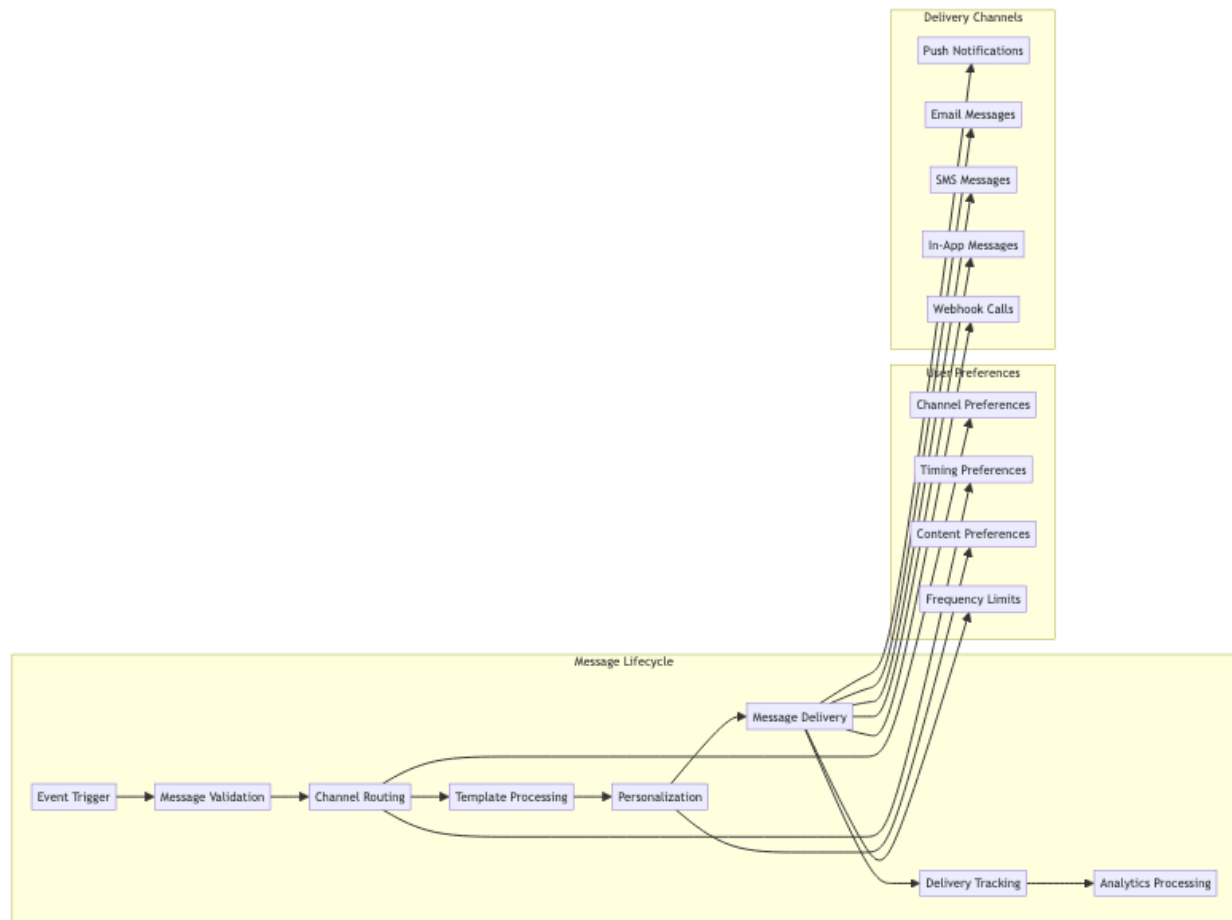
System Architecture Overview

□ [Back to Top](#)



Notification Flow Architecture

□ [Back to Top](#)

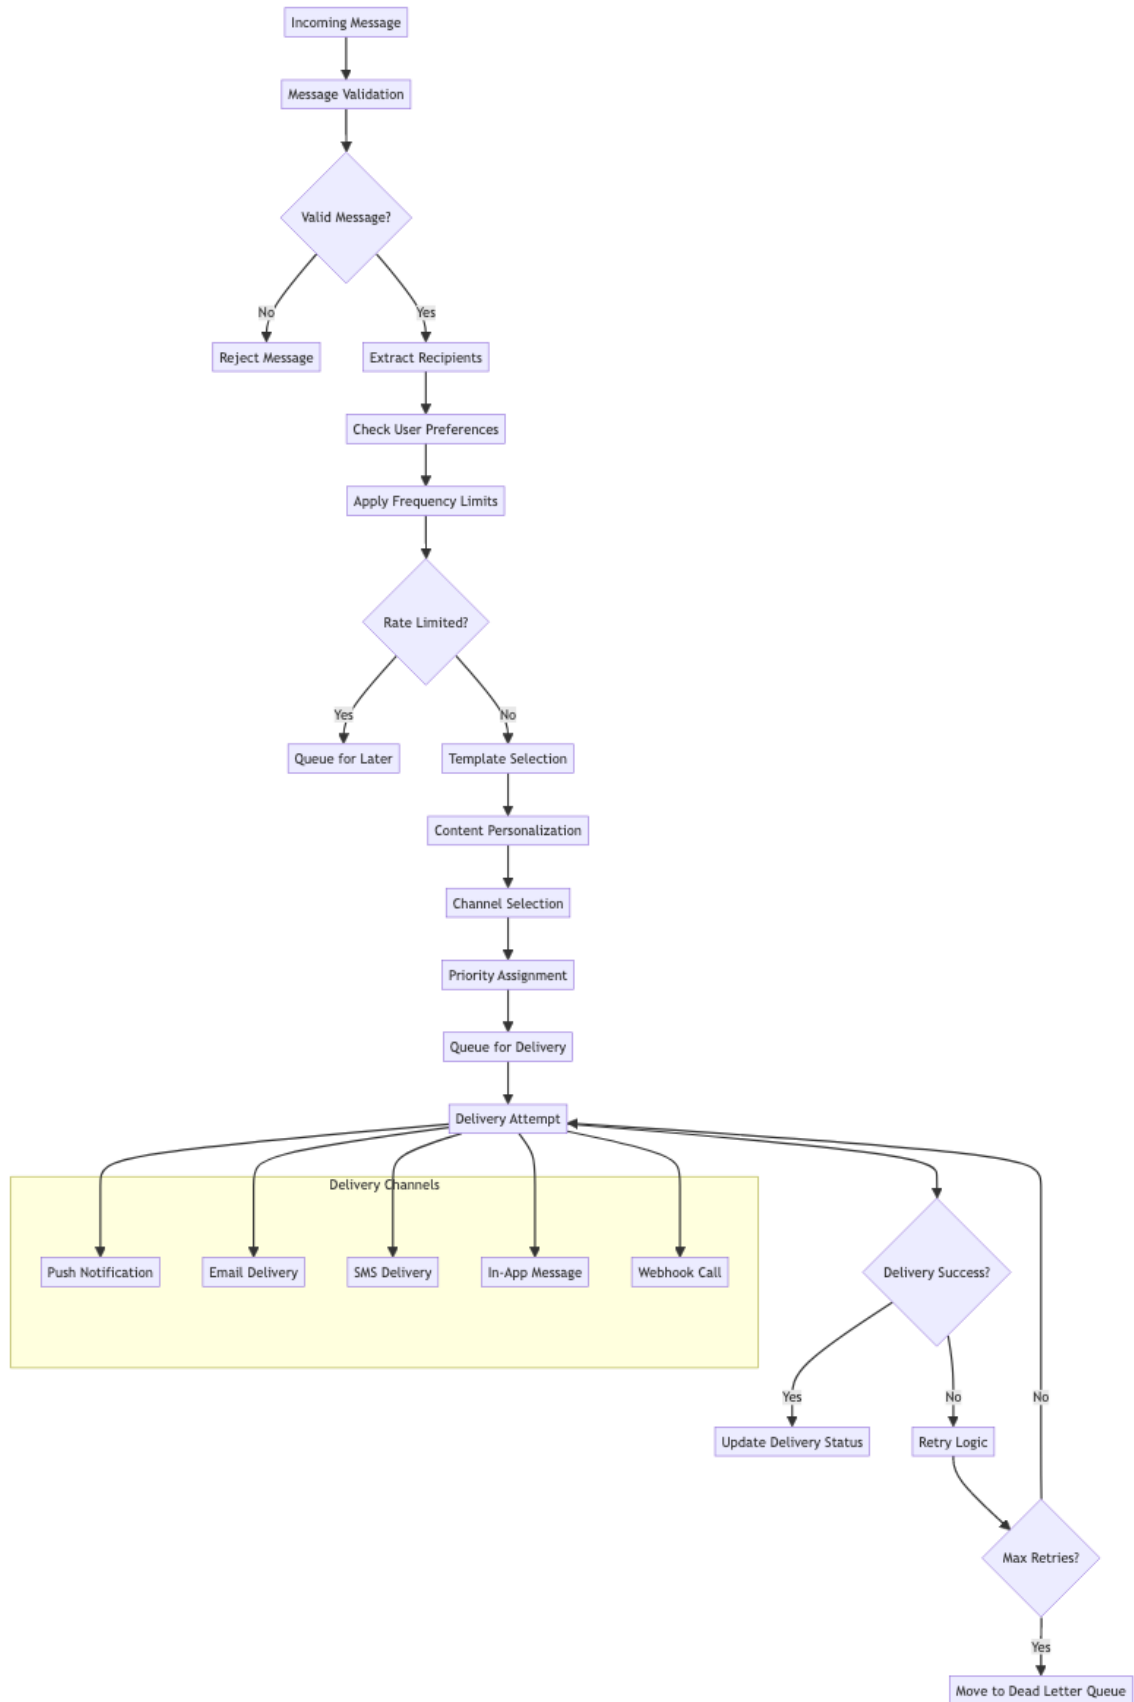


Low-Level Design (LLD)

□ [Back to Top](#)

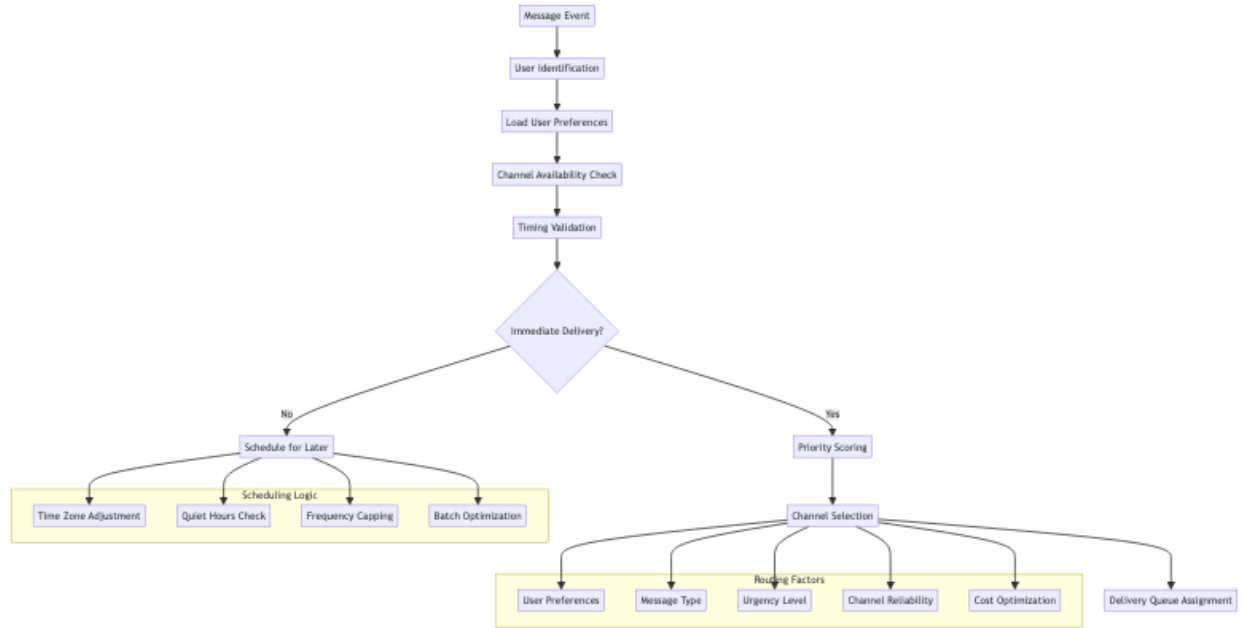
Message Processing Pipeline

□ [Back to Top](#)



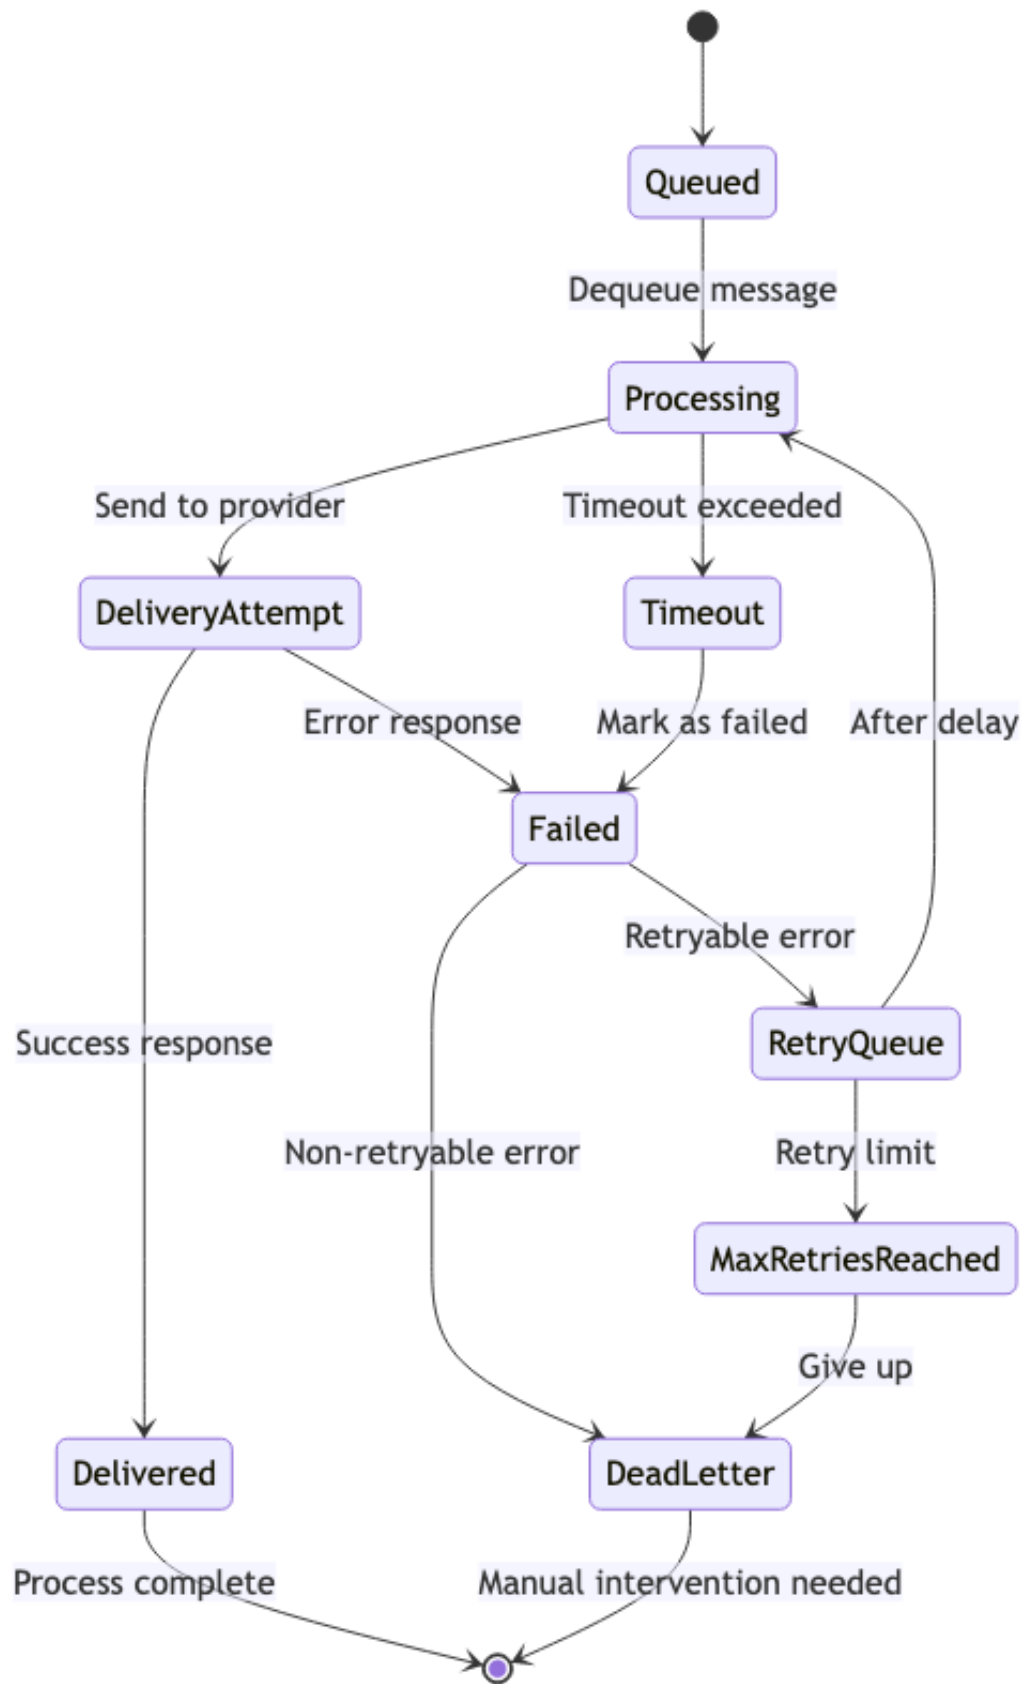
Smart Routing Engine

[Back to Top](#)



Delivery Reliability System

[Back to Top](#)



Core Algorithms

□ [Back to Top](#)

1. Intelligent Message Routing Algorithm

□ [Back to Top](#)

Purpose: Route messages through optimal channels based on user preferences, message urgency, and delivery success rates.

Channel Selection Algorithm:

```
ChannelConfig = {
  channels: {
    push: { reliability: 0.95, cost: 0.001, latency: 1000 },
    email: { reliability: 0.98, cost: 0.01, latency: 5000 },
    sms: { reliability: 0.99, cost: 0.05, latency: 3000 },
    inapp: { reliability: 0.90, cost: 0.0001, latency: 500 },
    webhook: { reliability: 0.85, cost: 0.002, latency: 2000 }
  },
  routingStrategy: 'reliability_optimized', // or 'cost_optimized', 'latency_optimized'
  fallbackEnabled: true,
  maxChannelsPerMessage: 3
}

function selectOptimalChannels(message, user, context):
  // Get user's enabled channels
  userChannels = getUserEnabledChannels(user.id, message.type)

  if userChannels.length === 0:
    return { success: false, reason: 'no_enabled_channels' }

  // Score each channel
  scoredChannels = userChannels.map(channel => ({
    channel: channel,
    score: calculateChannelScore(channel, message, user, context)
  }))

  // Sort by score (descending)
  rankedChannels = scoredChannels.sort((a, b) => b.score - a.score)

  // Select primary and fallback channels
```

```

selectedChannels = rankedChannels.slice(0, ChannelConfig.maxChannelsPerMessage)

return {
  success: true,
  primary: selectedChannels[0],
  fallbacks: selectedChannels.slice(1)
}

function calculateChannelScore(channel, message, user, context):
  channelInfo = ChannelConfig.channels[channel]
  score = 0

  // Reliability factor (40% weight)
  reliabilityScore = channelInfo.reliability
  score += reliabilityScore * 0.4

  // User preference factor (30% weight)
  userPreference = getUserChannelPreference(user.id, channel, message.type)
  score += userPreference * 0.3

  // Message urgency alignment (20% weight)
  urgencyAlignment = calculateUrgencyAlignment(channel, message.urgency)
  score += urgencyAlignment * 0.2

  // Cost factor (10% weight) - lower cost = higher score
  costScore = 1 - (channelInfo.cost / MAX_CHANNEL_COST)
  score += costScore * 0.1

  // Apply channel-specific penalties
  if isChannelOverloaded(channel):
    score *= 0.8

  if hasRecentFailures(user.id, channel):
    score *= 0.7

  return score

```

Smart Fallback Logic:

```

function executeDeliveryWithFallback(message, channelSelection, user):
  deliveryAttempts = []

  // Try primary channel
  primaryResult = attemptDelivery(message, channelSelection.primary, user)
  deliveryAttempts.push(primaryResult)

```



```

if primaryResult.success:
    return {
        success: true,
        finalChannel: channelSelection.primary.channel,
        attempts: deliveryAttempts
    }

// Try fallback channels
for fallbackChannel in channelSelection.fallbacks:
    // Check if fallback is appropriate
    if not shouldRetryWithFallback(primaryResult, fallbackChannel, message):
        continue

    fallbackResult = attemptDelivery(message, fallbackChannel, user)
    deliveryAttempts.push(fallbackResult)

    if fallbackResult.success:
        return {
            success: true,
            finalChannel: fallbackChannel.channel,
            attempts: deliveryAttempts
        }

// All channels failed
return {
    success: false,
    attempts: deliveryAttempts,
    reason: 'all_channels_failed'
}

```

2. Frequency Capping and Rate Limiting Algorithm

□ [Back to Top](#)

Purpose: Prevent notification fatigue by intelligently limiting message frequency per user.

Frequency Limit Configuration:

```

FrequencyLimits = {
  global: {
    maxPerHour: 10,
    maxPerDay: 50,
    maxPerWeek: 200
  },
  byType: {

```

```

    marketing: { maxPerHour: 2, maxPerDay: 5, maxPerWeek: 20 },
    transactional: { maxPerHour: 20, maxPerDay: 100, maxPerWeek: 500 },
    system: { maxPerHour: 5, maxPerDay: 20, maxPerWeek: 100 },
    social: { maxPerHour: 8, maxPerDay: 30, maxPerWeek: 150 }
  },
  byChannel: {
    push: { maxPerHour: 5, maxPerDay: 20, maxPerWeek: 100 },
    email: { maxPerHour: 3, maxPerDay: 10, maxPerWeek: 50 },
    sms: { maxPerHour: 2, maxPerDay: 5, maxPerWeek: 25 }
  }
}

function checkFrequencyLimits(userId, message, channel):
  currentTime = Date.now()

  // Get time windows
  hourStart = currentTime - 3600000 // 1 hour
  dayStart = currentTime - 86400000 // 24 hours
  weekStart = currentTime - 604800000 // 7 days

  // Check global limits
  globalCounts = getMessageCounts(userId, hourStart, dayStart, weekStart)
  if isGlobalLimitExceeded(globalCounts):
    return {
      allowed: false,
      reason: 'global_frequency_limit',
      retryAfter: calculateRetryTime(globalCounts, FrequencyLimits.global)
    }

  // Check type-specific limits
  typeCounts = getMessageCountsByType(userId, message.type, hourStart, dayStart, weekStart)
  typeLimit = FrequencyLimits.byType[message.type]
  if isTypeLimitExceeded(typeCounts, typeLimit):
    return {
      allowed: false,
      reason: 'type_frequency_limit',
      retryAfter: calculateRetryTime(typeCounts, typeLimit)
    }

  // Check channel-specific limits
  channelCounts = getMessageCountsByChannel(userId, channel, hourStart, dayStart, weekStart)
  channelLimit = FrequencyLimits.byChannel[channel]
  if isChannelLimitExceeded(channelCounts, channelLimit):
    return {
      allowed: false,

```

```

        reason: 'channel_frequency_limit',
        retryAfter: calculateRetryTime(channelCounts, channelLimit)
    }

```

```

return { allowed: true }

```

Adaptive Frequency Management:

```

function adaptFrequencyLimits(userId, engagementMetrics):
    userProfile = getUserProfile(userId)
    baseFrequencyLimits = FrequencyLimits.global

    // Adjust based on engagement rate
    engagementRate = calculateEngagementRate(userId, engagementMetrics)

    if engagementRate > 0.8:
        // High engagement - allow more notifications
        adjustmentFactor = 1.3
    else if engagementRate > 0.5:
        // Medium engagement - keep standard limits
        adjustmentFactor = 1.0
    else if engagementRate > 0.2:
        // Low engagement - reduce notifications
        adjustmentFactor = 0.7
    else:
        // Very low engagement - significantly reduce
        adjustmentFactor = 0.4

    // Adjust based on user activity patterns
    activityScore = calculateActivityScore(userId)
    adjustmentFactor *= (0.8 + activityScore * 0.4) // 0.8 to 1.2 range

    // Adjust based on subscription tier
    if userProfile.tier === 'premium':
        adjustmentFactor *= 1.2
    else if userProfile.tier === 'basic':
        adjustmentFactor *= 0.8

    // Apply adjustments
    adaptedLimits = {
        maxPerHour: Math.floor(baseFrequencyLimits.maxPerHour * adjustmentFactor),
        maxPerDay: Math.floor(baseFrequencyLimits.maxPerDay * adjustmentFactor),
        maxPerWeek: Math.floor(baseFrequencyLimits.maxPerWeek * adjustmentFactor)
    }

    // Store adapted limits for user

```

```

storeUserFrequencyLimits(userId, adaptedLimits)

return adaptedLimits

```

3. Template Engine and Personalization Algorithm

□ [Back to Top](#)

Purpose: Generate personalized message content using templates and user data.

Dynamic Template Processing:

```

TemplateConfig = {
    supportedFormats: ['html', 'text', 'markdown', 'json'],
    personalizationTokens: {
        user: ['firstName', 'lastName', 'email', 'timezone'],
        context: ['timestamp', 'action', 'location'],
        content: ['subject', 'body', 'ctaText', 'imageUrl']
    },
    fallbackValues: {
        firstName: 'there',
        lastName: '',
        timezone: 'UTC'
    }
}

function processTemplate(templateId, userData, contextData, messageType):
    // Retrieve template
    template = getTemplate(templateId, messageType)

    if not template:
        return { success: false, error: 'template_not_found' }

    // Prepare personalization data
    personalizationData = {
        user: extractUserData(userData),
        context: processContextData(contextData),
        timestamp: formatTimestamp(Date.now(), userData.timezone || 'UTC'),
        metadata: extractMetadata(contextData)
    }

    // Process template for each format
    processedContent = {}

    for format in template.formats:

```

```

    try:
        processedContent[format] = renderTemplate(
            template.content[format],
            personalizationData
        )
    catch error:
        // Use fallback template
        fallbackTemplate = getFallbackTemplate(messageType, format)
        processedContent[format] = renderTemplate(fallbackTemplate, personalizationData)

// Apply content optimization
optimizedContent = optimizeContent(processedContent, userData.preferences)

return {
    success: true,
    content: optimizedContent,
    metadata: {
        templateId: templateId,
        personalizationApplied: true,
        fallbackUsed: template.fallbackUsed || false
    }
}
}

function renderTemplate(templateContent, data):
    // Replace personalization tokens
    renderedContent = templateContent

    // Simple token replacement
    for key in Object.keys(data):
        if typeof data[key] === 'object':
            for subKey in Object.keys(data[key]):
                token = `{{${key}.${subKey}}}`
                value = data[key][subKey] || TemplateConfig.fallbackValues[subKey] || ''
                renderedContent = renderedContent.replace(new RegExp(token, 'g'), value)
            else:
                token = `{{${key}}}`
                value = data[key] || TemplateConfig.fallbackValues[key] || ''
                renderedContent = renderedContent.replace(new RegExp(token, 'g'), value)

    // Process conditional blocks
    renderedContent = processConditionalBlocks(renderedContent, data)

    // Process loops
    renderedContent = processLoops(renderedContent, data)

```

```
return renderedContent
```

A/B Testing Framework:

```
function selectTemplateVariant(templateId, userId, experimentConfig):
    // Get active experiments for this template
    activeExperiments = getActiveExperiments(templateId)

    if activeExperiments.length === 0:
        return { templateId: templateId, variant: 'control' }

    // Select experiment based on user cohort
    selectedExperiment = selectExperimentForUser(userId, activeExperiments)

    if not selectedExperiment:
        return { templateId: templateId, variant: 'control' }

    // Determine variant assignment
    userHash = hashUserId(userId, selectedExperiment.id)
    variantIndex = userHash % selectedExperiment.variants.length
    selectedVariant = selectedExperiment.variants[variantIndex]

    // Track experiment assignment
    trackExperimentAssignment(userId, selectedExperiment.id, selectedVariant.id)

    return {
        templateId: selectedVariant.templateId,
        variant: selectedVariant.name,
        experimentId: selectedExperiment.id
    }
```

4. Delivery Reliability and Retry Logic

□ [Back to Top](#)

Purpose: Ensure reliable message delivery with intelligent retry mechanisms.

Retry Strategy Configuration:

```
RetryConfig = {
    maxRetries: 5,
    baseDelay: 1000,          // 1 second
    maxDelay: 300000,         // 5 minutes
    backoffMultiplier: 2,
    jitterRange: 0.1,         // 10% jitter
    retryableErrors: [
```

```

        'network_timeout',
        'rate_limit_exceeded',
        'temporary_service_unavailable',
        'gateway_timeout'
    ],
    nonRetryableErrors: [
        'invalid_token',
        'invalid_recipient',
        'message_too_large',
        'content_blocked'
    ]
}

function executeDeliveryWithRetry(message, channel, user):
    attempt = 0
    lastError = null

    while attempt <= RetryConfig.maxRetries:
        try:
            // Attempt delivery
            result = attemptDelivery(message, channel, user)

            if result.success:
                // Update success metrics
                updateDeliveryMetrics(channel, 'success', attempt)
                return {
                    success: true,
                    attempts: attempt + 1,
                    finalResult: result
                }
            else:
                lastError = result.error

                // Check if error is retryable
                if not isRetryableError(lastError):
                    updateDeliveryMetrics(channel, 'permanent_failure', attempt)
                    return {
                        success: false,
                        attempts: attempt + 1,
                        error: lastError,
                        retryable: false
                    }

        catch error:
            lastError = error

```

```

    // Calculate next retry delay
    if attempt < RetryConfig.maxRetries:
        delay = calculateRetryDelay(attempt)
        scheduleRetry(message, channel, user, delay, attempt + 1)

    attempt++

// Max retries exceeded
updateDeliveryMetrics(channel, 'max_retries_exceeded', attempt)
return {
    success: false,
    attempts: attempt,
    error: lastError,
    retryable: true,
    maxRetriesExceeded: true
}

function calculateRetryDelay(attemptNumber):
    // Exponential backoff with jitter
    exponentialDelay = RetryConfig.baseDelay * Math.pow(RetryConfig.backoffMultiplier, attemptNumber)
    cappedDelay = Math.min(exponentialDelay, RetryConfig.maxDelay)

    // Add jitter to prevent thundering herd
    jitter = cappedDelay * RetryConfig.jitterRange * (Math.random() - 0.5)
    finalDelay = cappedDelay + jitter

    return Math.max(finalDelay, RetryConfig.baseDelay)

```

Circuit Breaker Pattern:

```

CircuitBreakerConfig = {
    failureThreshold: 5,           // Failures before opening
    timeoutThreshold: 10000,       // 10 seconds
    recoveryTimeout: 60000,        // 1 minute recovery time
    halfOpenMaxCalls: 3            // Test calls in half-open state
}

function executeWithCircuitBreaker(deliveryFunction, channel):
    circuitBreaker = getCircuitBreaker(channel)

    if circuitBreaker.state === 'open':
        if Date.now() - circuitBreaker.lastFailureTime > CircuitBreakerConfig.recoveryTimeout:
            // Move to half-open state
            circuitBreaker.state = 'half_open'
            circuitBreaker.halfOpenAttempts = 0

```



```

else:
    // Circuit is still open
    return {
        success: false,
        error: 'circuit_breaker_open',
        nextRetryTime: circuitBreaker.lastFailureTime + CircuitBreakerConfig.recoveryTime
    }

try:
    // Execute delivery function
    result = deliveryFunction()

    if result.success:
        if circuitBreaker.state === 'half_open':
            circuitBreaker.halfOpenAttempts++
            if circuitBreaker.halfOpenAttempts >= CircuitBreakerConfig.halfOpenMaxCalls:
                // Close circuit - service recovered
                circuitBreaker.state = 'closed'
                circuitBreaker.failureCount = 0
            else:
                // Reset failure count on success
                circuitBreaker.failureCount = 0

        return result
    else:
        // Handle failure
        return handleCircuitBreakerFailure(circuitBreaker, result)

catch error:
    return handleCircuitBreakerFailure(circuitBreaker, { success: false, error: error })

```

5. Analytics and Performance Monitoring

□ [Back to Top](#)

Purpose: Track delivery performance and user engagement for optimization.

Real-time Metrics Collection:

```

MetricsConfig = {
  trackingEvents: [
    'message_sent',
    'message_delivered',
    'message_opened',
    'message_clicked',
  ]
}

```

```

        'message_failed',
        'user_opted_out'
    ],
    aggregationWindows: [60000, 300000, 3600000, 86400000], // 1m, 5m, 1h, 1d
    retentionPeriod: 90 * 24 * 3600 * 1000 // 90 days
}

function trackDeliveryMetrics(event, messageId, userId, channel, metadata):
    eventData = {
        messageId: messageId,
        userId: userId,
        channel: channel,
        event: event,
        timestamp: Date.now(),
        metadata: metadata
    }

    // Store raw event
    storeEvent(eventData)

    // Update real-time counters
    updateRealtimeCounters(event, channel, metadata)

    // Update user engagement profile
    updateUserEngagementProfile(userId, event, channel)

    // Trigger real-time alerts if needed
    checkAlertThresholds(event, channel, metadata)

function calculateDeliveryMetrics(timeWindow, filters = {}):
    events = getEvents(timeWindow, filters)

    metrics = {
        totalSent: 0,
        totalDelivered: 0,
        totalOpened: 0,
        totalClicked: 0,
        totalFailed: 0,
        deliveryRate: 0,
        openRate: 0,
        clickRate: 0,
        failureRate: 0,
        channelBreakdown: {},
        errorBreakdown: {}
    }
}

```

```

for event in events:
    switch event.event:
        case 'message_sent':
            metrics.totalSent++
            break
        case 'message_delivered':
            metrics.totalDelivered++
            break
        case 'message_opened':
            metrics.totalOpened++
            break
        case 'message_clicked':
            metrics.totalClicked++
            break
        case 'message_failed':
            metrics.totalFailed++
            break

    // Update channel breakdown
    if not metrics.channelBreakdown[event.channel]:
        metrics.channelBreakdown[event.channel] = { sent: 0, delivered: 0, failed: 0 }

    if event.event === 'message_sent':
        metrics.channelBreakdown[event.channel].sent++
    else if event.event === 'message_delivered':
        metrics.channelBreakdown[event.channel].delivered++
    else if event.event === 'message_failed':
        metrics.channelBreakdown[event.channel].failed++

    // Calculate rates
    if metrics.totalSent > 0:
        metrics.deliveryRate = metrics.totalDelivered / metrics.totalSent
        metrics.failureRate = metrics.totalFailed / metrics.totalSent

    if metrics.totalDelivered > 0:
        metrics.openRate = metrics.totalOpened / metrics.totalDelivered

    if metrics.totalOpened > 0:
        metrics.clickRate = metrics.totalClicked / metrics.totalOpened

return metrics

```

Performance Optimizations

[❏ Back to Top](#)

Message Queue Optimization

[❏ Back to Top](#)

Priority Queue Management:

```
QueueConfig = {
  priorities: {
    critical: 1,    // System alerts, security notifications
    high: 2,       // Transactional messages
    medium: 3,     // User-generated notifications
    low: 4,        // Marketing messages
    bulk: 5        // Batch notifications
  },
  batchSize: {
    critical: 1,
    high: 10,
    medium: 50,
    low: 100,
    bulk: 500
  }
}
```

Database Optimization

[❏ Back to Top](#)

Sharding Strategy: - User preferences: Shard by user ID - Message templates: Replicate across all shards - Delivery history: Partition by time and shard by user ID - Analytics data: Time-series partitioning

Caching Strategy

[❏ Back to Top](#)

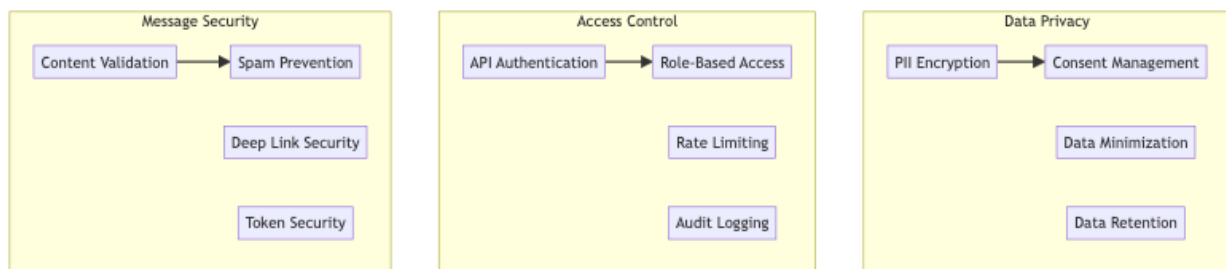
Multi-Layer Caching: - User preferences in Redis (1-hour TTL) - Message templates in application cache - Delivery provider configurations in memory - Analytics data in time-series cache

Security Considerations

□ [Back to Top](#)

Privacy and Data Protection

□ [Back to Top](#)



Testing Strategy

□ [Back to Top](#)

Load Testing Scenarios

□ [Back to Top](#)

High Volume Testing: - Black Friday notification bursts - Breaking news alerts - System-wide outage notifications - Batch campaign processing

Reliability Testing

□ [Back to Top](#)

Failure Scenarios: - Provider outage handling - Database failover testing - Network partition recovery - Queue overflow management

Trade-offs and Considerations

□ [Back to Top](#)

Delivery Speed vs Reliability

□ [Back to Top](#)

- **Immediate delivery:** Speed vs delivery guarantees
- **Retry mechanisms:** Reliability vs resource usage
- **Circuit breakers:** Fault tolerance vs message throughput
- **Queue processing:** Throughput vs message ordering

Personalization vs Privacy

□ [Back to Top](#)

- **User data usage:** Personalization vs data minimization
- **Tracking:** Analytics vs user privacy
- **Content optimization:** Relevance vs data collection
- **Behavioral targeting:** Effectiveness vs consent requirements

Scalability vs Cost

□ [Back to Top](#)

- **Multiple providers:** Reliability vs operational complexity
- **Real-time processing:** Speed vs computational resources
- **Global delivery:** Coverage vs infrastructure cost
- **Analytics granularity:** Insights vs storage expenses

This distributed notification system provides a comprehensive foundation for large-scale messaging with features like intelligent routing, frequency management, template personalization, and robust delivery guarantees while maintaining high performance, security, and user privacy standards.