

# API Rate Limiter System

## □ Table of Contents

- API Rate Limiter System
    - High-Level Design (HLD)
      - \* System Architecture Overview
      - \* Rate Limiting Flow
    - Low-Level Design (LLD)
      - \* Rate Limiting Engine
      - \* Distributed Rate Limiting
      - \* Policy Hierarchy System
    - Core Algorithms
      - \* 1. Token Bucket Algorithm
      - \* 2. Sliding Window Algorithm
      - \* 3. Hierarchical Rate Limiting Algorithm
      - \* 4. Smart Quota Management Algorithm
      - \* 5. Adaptive Rate Limiting Algorithm
    - Performance Optimizations
      - \* Memory-Efficient Storage
      - \* Distributed Coordination
    - Security Considerations
      - \* Rate Limiting Security Framework
    - Testing Strategy
      - \* Load Testing
      - \* Accuracy Testing
    - Trade-offs and Considerations
      - \* Accuracy vs Performance
      - \* Flexibility vs Simplicity
      - \* Cost vs Reliability
- 

## High-Level Design (HLD)

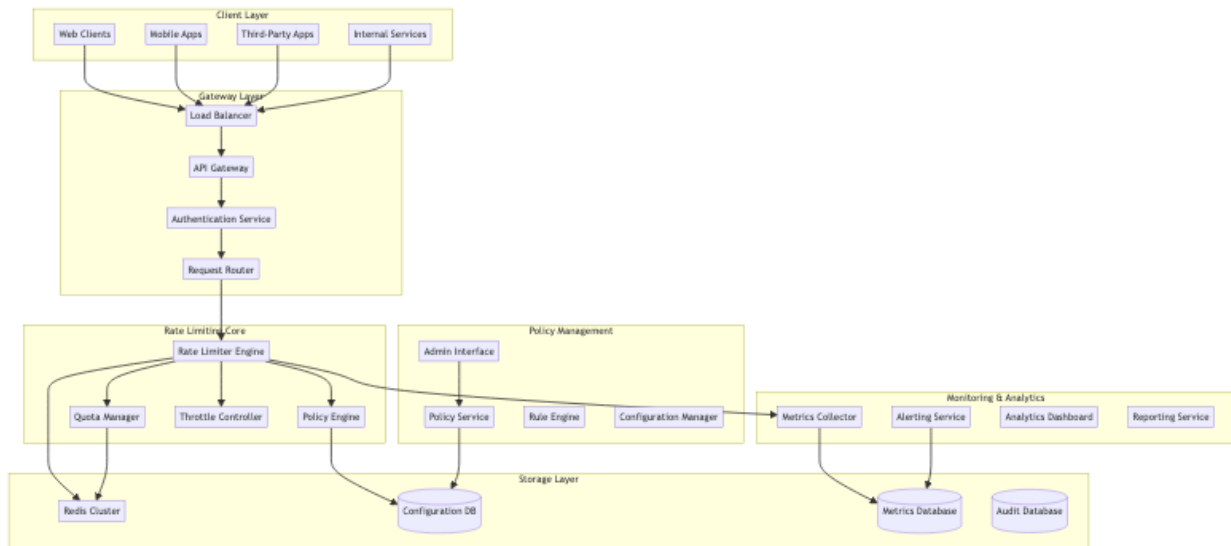
□ [Back to Top](#)

---

## System Architecture Overview

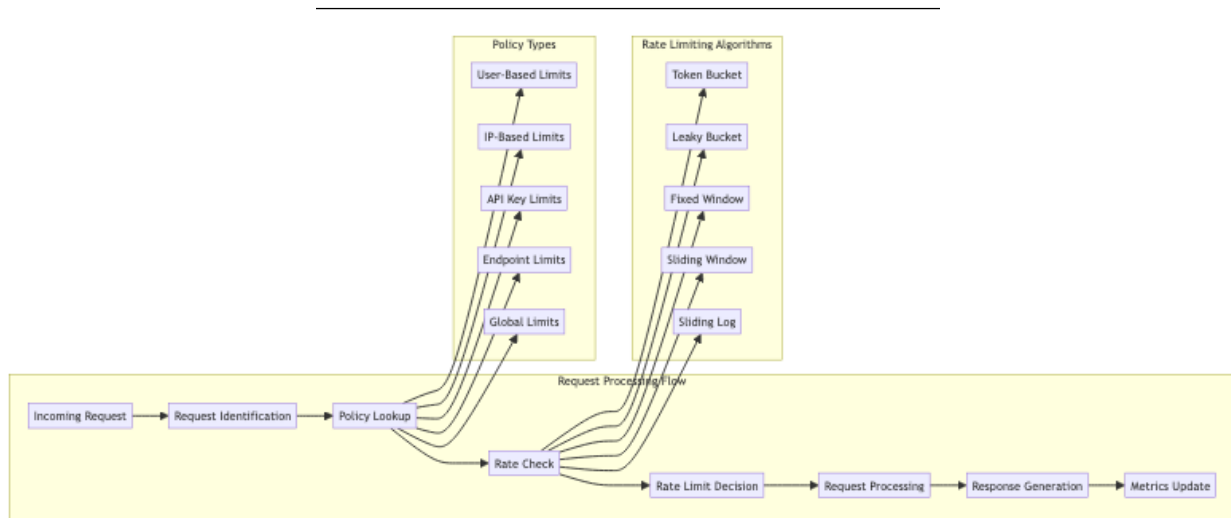
□ [Back to Top](#)

---



## Rate Limiting Flow

□ [Back to Top](#)

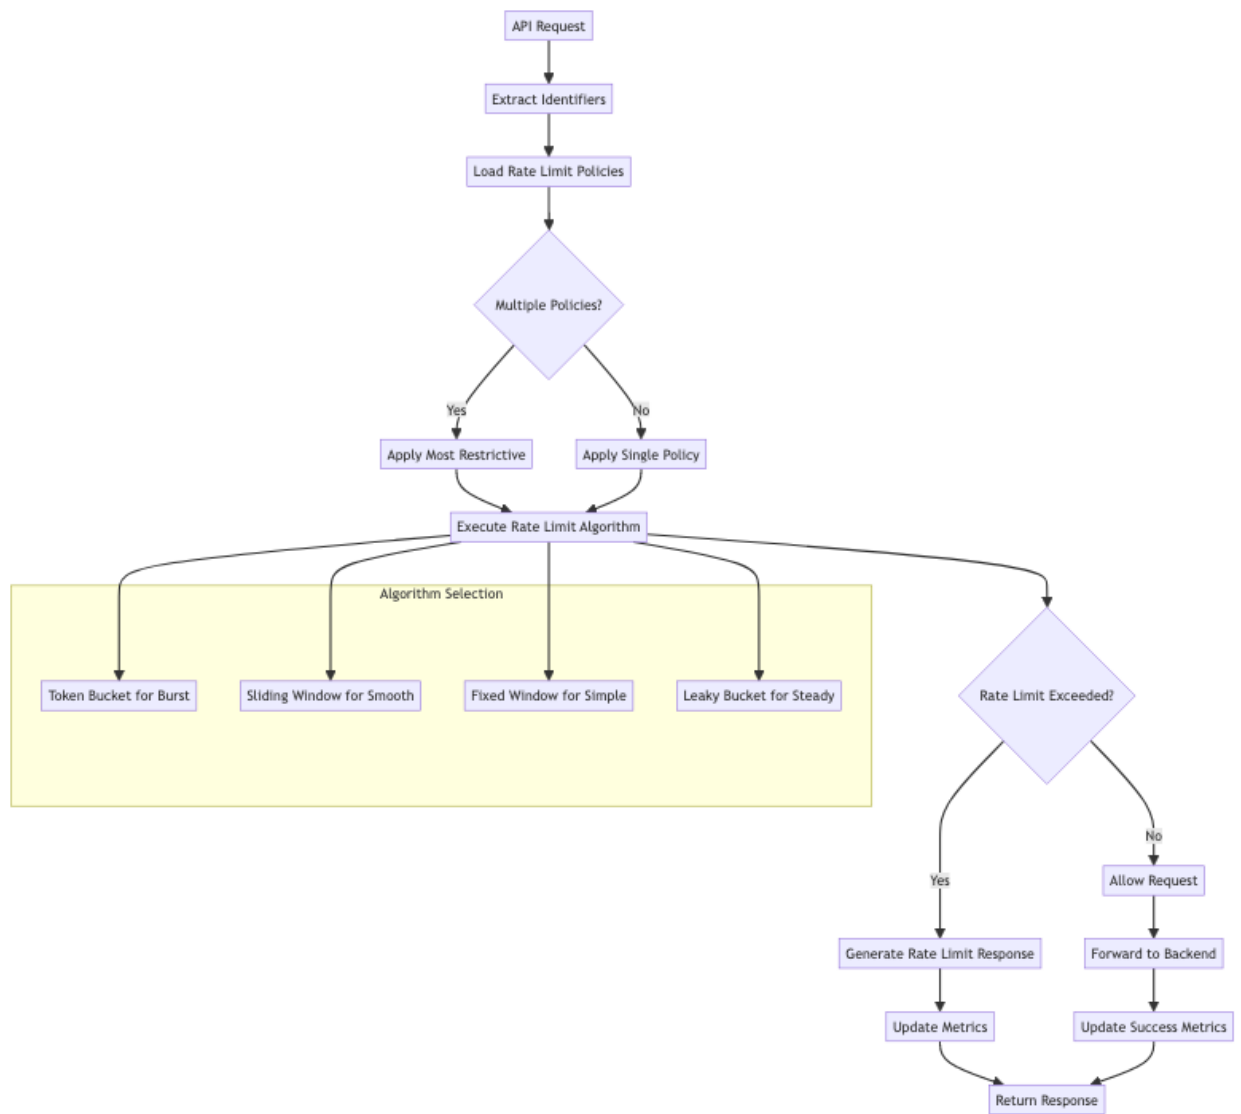


## Low-Level Design (LLD)

□ [Back to Top](#)

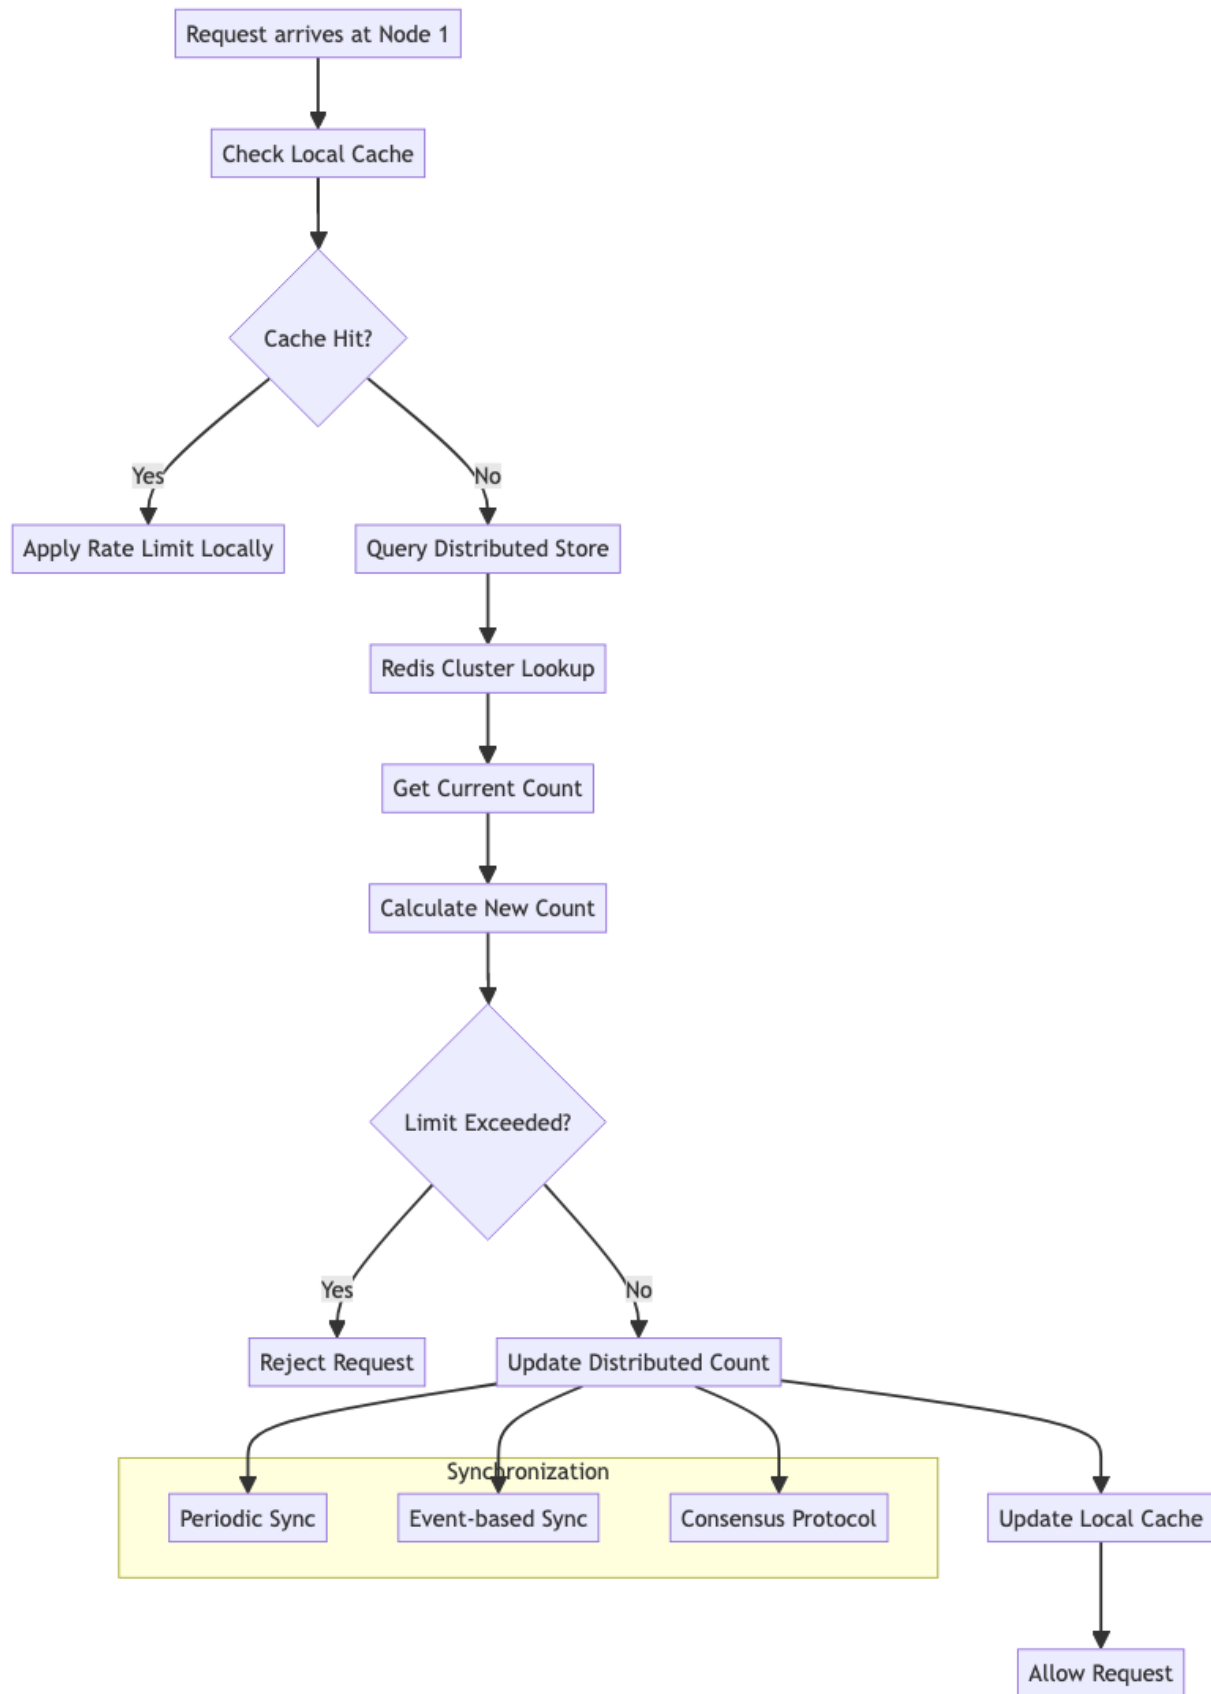
## Rate Limiting Engine

□ [Back to Top](#)



## Distributed Rate Limiting

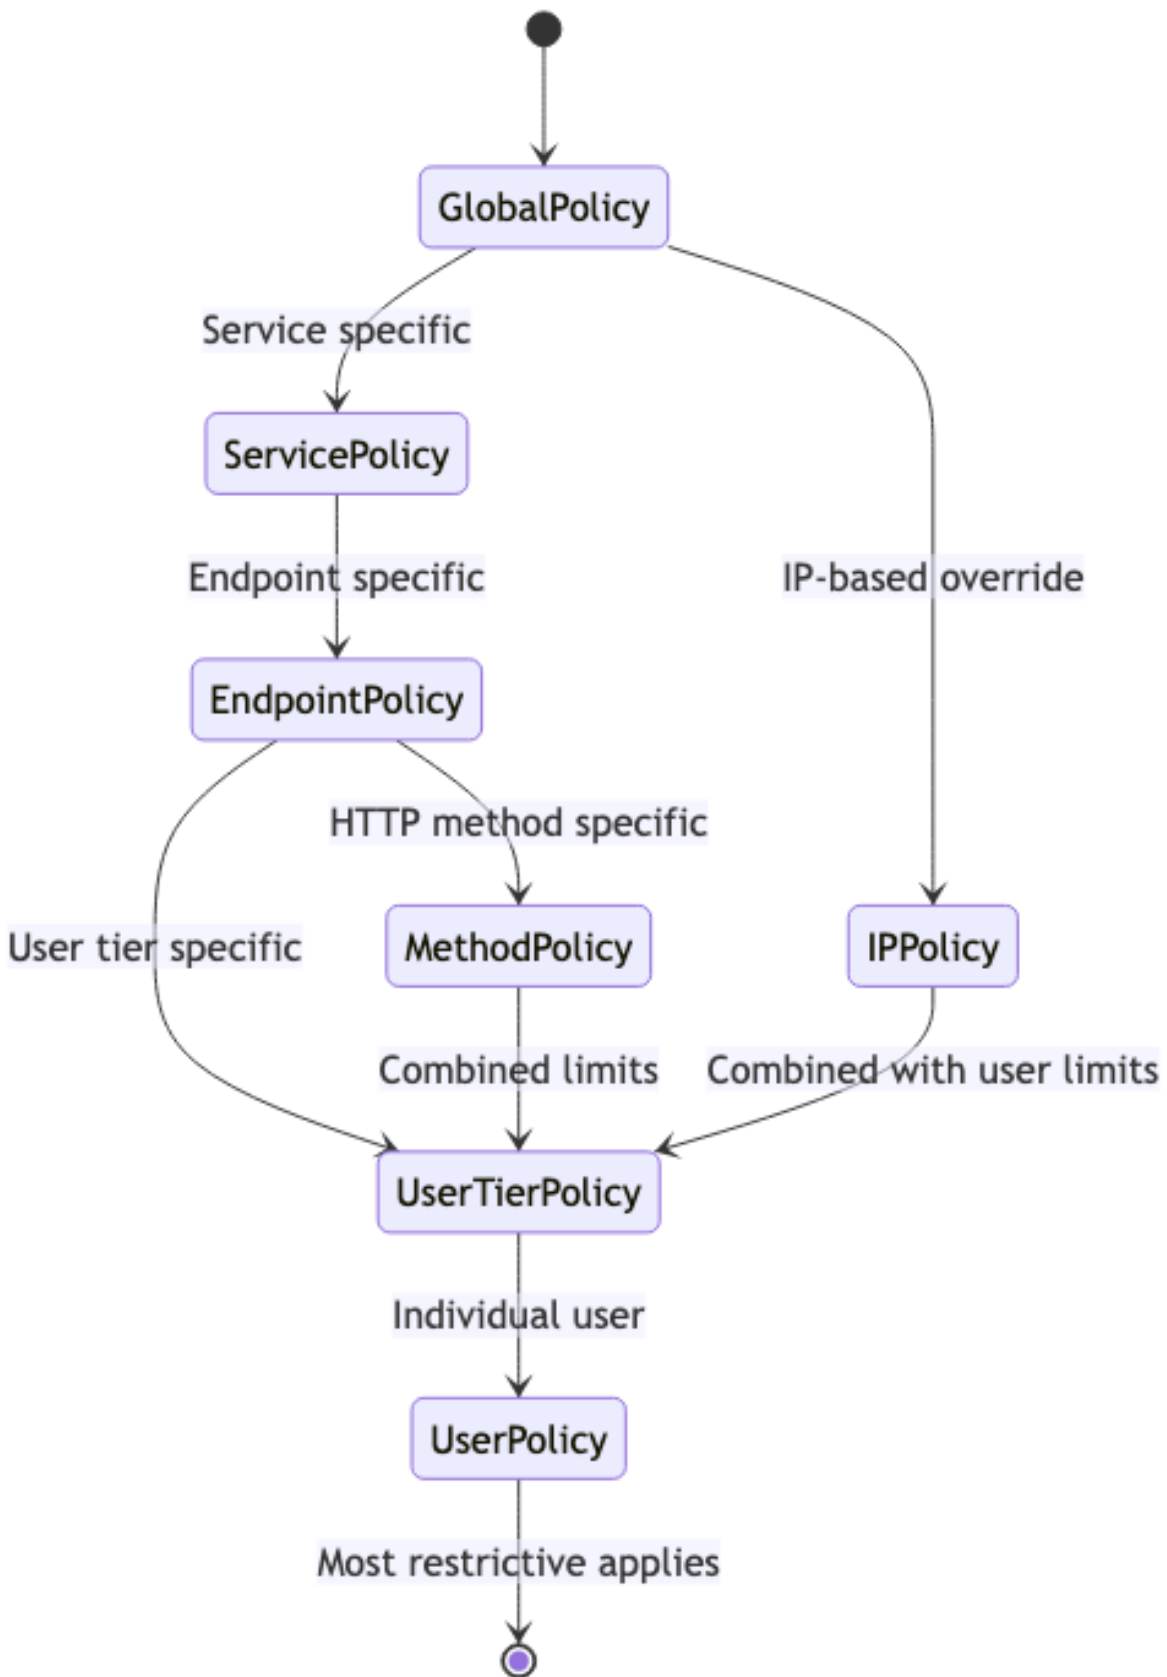
□ [Back to Top](#)



## Policy Hierarchy System

[□ Back to Top](#)

---



## Core Algorithms

□ [Back to Top](#)

---

### 1. Token Bucket Algorithm

□ [Back to Top](#)

---

**Purpose:** Allow burst traffic while maintaining average rate limits with token-based allowance system.

#### Token Bucket Implementation:

```
TokenBucketConfig = {  
  capacity: 100,           // Maximum tokens in bucket  
  refillRate: 10,          // Tokens added per second  
  refillInterval: 1000,    // Refill interval in milliseconds  
  initialTokens: 100       // Starting token count  
}
```

```
class TokenBucket:  
  constructor(config):  
    this.capacity = config.capacity  
    this.tokens = config.initialTokens  
    this.refillRate = config.refillRate  
    this.refillInterval = config.refillInterval  
    this.lastRefillTime = Date.now()  
  
  function checkRateLimit(requestWeight = 1):  
    currentTime = Date.now()  
  
    // Refill tokens based on elapsed time  
    this.refillTokens(currentTime)  
  
    // Check if enough tokens available  
    if this.tokens >= requestWeight:  
      this.tokens -= requestWeight  
      return {  
        allowed: true,  
        remainingTokens: this.tokens,  
        resetTime: this.calculateResetTime(),  
        retryAfter: null  
      }  
    }
```

```

else:
    return {
        allowed: false,
        remainingTokens: this.tokens,
        resetTime: this.calculateResetTime(),
        retryAfter: this.calculateRetryAfter(requestWeight)
    }

function refillTokens(currentTime):
    timeSinceLastRefill = currentTime - this.lastRefillTime

    if timeSinceLastRefill >= this.refillInterval:
        intervalsElapsed = Math.floor(timeSinceLastRefill / this.refillInterval)
        tokensToAdd = intervalsElapsed * this.refillRate

        this.tokens = Math.min(this.capacity, this.tokens + tokensToAdd)
        this.lastRefillTime = currentTime

function calculateRetryAfter(requestWeight):
    tokensNeeded = requestWeight - this.tokens
    timeToGetTokens = Math.ceil(tokensNeeded / this.refillRate) * this.refillInterval
    return timeToGetTokens

```

### **Distributed Token Bucket:**

```

function distributedTokenBucket(key, requestWeight, config):
    // Use Lua script for atomic operations in Redis
    luaScript = `
        local key = KEYS[1]
        local capacity = tonumber(ARGV[1])
        local refillRate = tonumber(ARGV[2])
        local refillInterval = tonumber(ARGV[3])
        local requestWeight = tonumber(ARGV[4])
        local currentTime = tonumber(ARGV[5])

        -- Get current state or initialize
        local bucket = redis.call('HMGET', key, 'tokens', 'lastRefillTime')
        local tokens = tonumber(bucket[1]) or capacity
        local lastRefillTime = tonumber(bucket[2]) or currentTime

        -- Calculate tokens to add
        local timeSinceRefill = currentTime - lastRefillTime
        local intervalsElapsed = math.floor(timeSinceRefill / refillInterval)
        local tokensToAdd = intervalsElapsed * refillRate

        -- Update token count
    `

```



```

tokens = math.min(capacity, tokens + tokensToAdd)
local newLastRefillTime = lastRefillTime + (intervalsElapsed * refillInterval)

-- Check if request can be allowed
if tokens >= requestWeight then
    tokens = tokens - requestWeight
    redis.call('HMSET', key, 'tokens', tokens, 'lastRefillTime', newLastRefillTime)
    redis.call('EXPIRE', key, 3600) -- 1 hour expiry
    return {1, tokens, newLastRefillTime} -- Allowed
else
    redis.call('HMSET', key, 'tokens', tokens, 'lastRefillTime', newLastRefillTime)
    redis.call('EXPIRE', key, 3600)
    return {0, tokens, newLastRefillTime} -- Denied
end
,

result = redis.eval(luaScript, [key], [
    config.capacity,
    config.refillRate,
    config.refillInterval,
    requestWeight,
    Date.now()
])

return {
    allowed: result[0] == 1,
    remainingTokens: result[1],
    lastRefillTime: result[2]
}

```

## 2. Sliding Window Algorithm

□ [Back to Top](#)

---

**Purpose:** Provide smooth rate limiting by tracking requests over a moving time window with precise time-based calculations.

### Sliding Window Counter:

```

SlidingWindowConfig = {
    windowSizeMs: 60000,      // 1 minute window
    maxRequests: 100,         // Maximum requests per window
    subWindowCount: 10,       // Number of sub-windows for precision
    cleanupInterval: 30000     // Cleanup old entries every 30 seconds
}

```

```

class SlidingWindowCounter:
    constructor(config):
        this.windowSizeMs = config.windowSizeMs
        this.maxRequests = config.maxRequests
        this.subWindowSize = config.windowSizeMs / config.subWindowCount
        this.requestCounts = new Map() // subWindowId -> count
        this.lastCleanup = Date.now()

function checkRateLimit(requestWeight = 1):
    currentTime = Date.now()

    // Clean up old sub-windows
    this.cleanupOldWindows(currentTime)

    // Calculate current request count in sliding window
    currentCount = this.getCurrentWindowCount(currentTime)

    if currentCount + requestWeight <= this.maxRequests:
        // Record the request
        this.recordRequest(currentTime, requestWeight)

        return {
            allowed: true,
            currentCount: currentCount + requestWeight,
            resetTime: this.calculateResetTime(currentTime),
            retryAfter: null
        }
    else:
        return {
            allowed: false,
            currentCount: currentCount,
            resetTime: this.calculateResetTime(currentTime),
            retryAfter: this.calculateRetryAfter(currentTime, requestWeight)
        }

function getCurrentWindowCount(currentTime):
    windowStartTime = currentTime - this.windowSizeMs
    totalCount = 0

    for [subWindowId, count] of this.requestCounts:
        subWindowTime = subWindowId * this.subWindowSize

        if subWindowTime > windowStartTime:
            // Calculate partial count for overlapping sub-window

```

```

        if subWindowTime < windowStartTime + this.subWindowSize:
            overlapRatio = (subWindowTime + this.subWindowSize - windowStartTime) / this.s
            totalCount += count * overlapRatio
        else:
            totalCount += count

    return Math.floor(totalCount)

function recordRequest(currentTime, requestWeight):
    subWindowId = Math.floor(currentTime / this.subWindowSize)
    currentCount = this.requestCounts.get(subWindowId) || 0
    this.requestCounts.set(subWindowId, currentCount + requestWeight)

function cleanupOldWindows(currentTime):
    if currentTime - this.lastCleanup < SlidingWindowConfig.cleanupInterval:
        return

    cutoffTime = currentTime - this.windowSizeMs
    cutoffSubWindowId = Math.floor(cutoffTime / this.subWindowSize)

    for subWindowId of this.requestCounts.keys():
        if subWindowId < cutoffSubWindowId:
            this.requestCounts.delete(subWindowId)

    this.lastCleanup = currentTime

```

### **Distributed Sliding Window with Redis:**

```

function distributedSlidingWindow(key, requestWeight, config):
    luaScript = `
        local key = KEYS[1]
        local windowSizeMs = tonumber(ARGV[1])
        local maxRequests = tonumber(ARGV[2])
        local requestWeight = tonumber(ARGV[3])
        local currentTime = tonumber(ARGV[4])
        local subWindowSize = tonumber(ARGV[5])

        -- Calculate window boundaries
        local windowStart = currentTime - windowSizeMs
        local currentSubWindow = math.floor(currentTime / subWindowSize)

        -- Remove old entries
        redis.call('ZREMRANGEBYSCORE', key, '-inf', windowStart)

        -- Count current requests in window
        local currentCount = 0
    `

```

```

local entries = redis.call('ZRANGEBYSCORE', key, windowStart, '+inf', 'WITHSCORES')

for i = 1, #entries, 2 do
    local timestamp = tonumber(entries[i + 1])
    local weight = tonumber(entries[i])
    currentCount = currentCount + weight
end

-- Check if request can be allowed
if currentCount + requestWeight <= maxRequests then
    -- Add new request
    redis.call('ZADD', key, currentTime, requestWeight .. ':' .. currentTime)
    redis.call('EXPIRE', key, math.ceil(windowSizeMs / 1000) + 1)
    return {1, currentCount + requestWeight} -- Allowed
else
    return {0, currentCount} -- Denied
end
,

result = redis.eval(luaScript, [key], [
    config.windowSizeMs,
    config.maxRequests,
    requestWeight,
    Date.now(),
    config.windowSizeMs / 10 // 10 sub-windows
])

return {
    allowed: result[0] == 1,
    currentCount: result[1]
}

```

### 3. Hierarchical Rate Limiting Algorithm

□ [Back to Top](#)

---

**Purpose:** Apply multiple rate limiting policies in hierarchical order with proper precedence and combination rules.

#### Policy Hierarchy System:

```

PolicyHierarchy = {
    levels: [
        'global',      // System-wide limits
        'service',     // Per-service limits
    ]
}

```

```

    'endpoint',      // Per-endpoint limits
    'user_tier',     // User tier-based limits
    'user',          // Individual user limits
    'ip'             // IP-based limits
  ],

  combinationStrategy: 'most_restrictive', // or 'additive', 'multiplicative'
  inheritanceEnabled: true,
  overrideRules: new Map()
}

function evaluateHierarchicalLimits(request, context):
  applicablePolicies = []

  // Collect all applicable policies
  for level in PolicyHierarchy.levels:
    policy = findApplicablePolicy(level, request, context)
    if policy:
      applicablePolicies.push({
        level: level,
        policy: policy,
        priority: getPolicyPriority(level)
      })

  // Sort by priority
  applicablePolicies.sort((a, b) => b.priority - a.priority)

  // Apply combination strategy
  return combineRateLimitPolicies(applicablePolicies, request)

function combineRateLimitPolicies(policies, request):
  if policies.length === 0:
    return { allowed: true, reason: 'no_applicable_policies' }

  switch PolicyHierarchy.combinationStrategy:
    case 'most_restrictive':
      return applyMostRestrictivePolicy(policies, request)
    case 'additive':
      return applyAdditivePolicies(policies, request)
    case 'multiplicative':
      return applyMultiplicativePolicies(policies, request)
    default:
      return applyMostRestrictivePolicy(policies, request)

function applyMostRestrictivePolicy(policies, request):

```

```

mostRestrictive = null
minAllowedRate = Infinity

for policyEntry in policies:
    policy = policyEntry.policy
    rateLimit = calculateEffectiveRateLimit(policy, request)

    if rateLimit.requestsPerSecond < minAllowedRate:
        minAllowedRate = rateLimit.requestsPerSecond
        mostRestrictive = policyEntry

if mostRestrictive:
    return executeRateLimit(mostRestrictive.policy, request)
else:
    return { allowed: true, reason: 'no_restrictive_policy' }

function executeRateLimit(policy, request):
    // Select appropriate algorithm based on policy configuration
    switch policy.algorithm:
        case 'token_bucket':
            return executeTokenBucket(policy, request)
        case 'sliding_window':
            return executeSlidingWindow(policy, request)
        case 'fixed_window':
            return executeFixedWindow(policy, request)
        case 'leaky_bucket':
            return executeLeakyBucket(policy, request)
        default:
            throw new Error(`Unknown rate limiting algorithm: ${policy.algorithm}`)

```

### Dynamic Policy Adjustment:

```

function adjustPolicyBasedOnMetrics(policyId, metrics):
    policy = getPolicy(policyId)
    currentMetrics = metrics

    // Analyze key metrics
    errorRate = currentMetrics.errorRate
    avgResponseTime = currentMetrics.avgResponseTime
    cpuUtilization = currentMetrics.cpuUtilization

    // Calculate adjustment factors
    adjustmentFactor = 1.0

    // Adjust based on error rate
    if errorRate > 0.05: // 5% error rate threshold

```

```

    adjustmentFactor *= 0.8 // Reduce rate limit by 20%
else if errorRate < 0.01: // 1% error rate
    adjustmentFactor *= 1.1 // Increase rate limit by 10%

// Adjust based on response time
if avgResponseTime > 1000: // 1 second threshold
    adjustmentFactor *= 0.9 // Reduce rate limit by 10%
else if avgResponseTime < 200: // 200ms threshold
    adjustmentFactor *= 1.05 // Increase rate limit by 5%

// Adjust based on CPU utilization
if cpuUtilization > 0.8: // 80% CPU threshold
    adjustmentFactor *= 0.7 // Reduce rate limit by 30%
else if cpuUtilization < 0.4: // 40% CPU threshold
    adjustmentFactor *= 1.2 // Increase rate limit by 20%

// Apply bounds to prevent extreme adjustments
adjustmentFactor = Math.max(0.1, Math.min(2.0, adjustmentFactor))

// Update policy
newRateLimit = Math.floor(policy.baseRateLimit * adjustmentFactor)
updatePolicyRateLimit(policyId, newRateLimit)

// Log adjustment
logPolicyAdjustment(policyId, policy.baseRateLimit, newRateLimit, adjustmentFactor, me

```

## 4. Smart Quota Management Algorithm

□ [Back to Top](#)

---

**Purpose:** Manage API quotas across different time periods with intelligent allocation and usage tracking.

### Multi-Period Quota System:

```

QuotaConfig = {
  periods: {
    minute: { duration: 60000, allocation: 100 },
    hour: { duration: 3600000, allocation: 5000 },
    day: { duration: 86400000, allocation: 100000 },
    month: { duration: 2629746000, allocation: 2000000 }
  },

  allocationStrategy: 'proportional', // 'proportional', 'priority', 'waterfall'
  carryoverEnabled: true,             // Allow unused quota carryover

```

```

    burstAllowance: 0.2,                // 20% burst over period limit
    quotaRefreshStrategy: 'sliding'     // 'sliding' or 'fixed'
}

class SmartQuotaManager:
    constructor(config):
        this.config = config
        this.quotaUsage = new Map() // userId -> period -> usage
        this.quotaAllocations = new Map() // userId -> period -> allocation

    function checkQuotaAvailability(userId, requestWeight, period):
        currentUsage = this.getCurrentUsage(userId, period)
        allocatedQuota = this.getAllocatedQuota(userId, period)
        burstLimit = allocatedQuota * (1 + this.config.burstAllowance)

        // Check if request exceeds quota
        if currentUsage + requestWeight <= allocatedQuota:
            return {
                allowed: true,
                quotaType: 'normal',
                remainingQuota: allocatedQuota - (currentUsage + requestWeight),
                resetTime: this.calculateQuotaResetTime(period)
            }
        else if currentUsage + requestWeight <= burstLimit:
            return {
                allowed: true,
                quotaType: 'burst',
                remainingQuota: burstLimit - (currentUsage + requestWeight),
                resetTime: this.calculateQuotaResetTime(period),
                warning: 'using_burst_quota'
            }
        else:
            return {
                allowed: false,
                quotaType: 'exceeded',
                excessAmount: (currentUsage + requestWeight) - burstLimit,
                resetTime: this.calculateQuotaResetTime(period)
            }

    function allocateSmartQuota(userId, totalQuota, periods):
        allocations = {}

        // Get user's historical usage patterns
        usageHistory = this.getUserUsageHistory(userId)
        usagePatterns = this.analyzeUsagePatterns(usageHistory)

```



```

switch this.config.allocationStrategy:
  case 'proportional':
    allocations = this.allocateProportionally(totalQuota, periods, usagePatterns)
    break
  case 'priority':
    allocations = this.allocateByPriority(totalQuota, periods, usagePatterns)
    break
  case 'waterfall':
    allocations = this.allocateWaterfall(totalQuota, periods, usagePatterns)
    break

// Store allocations
this.quotaAllocations.set(userId, allocations)

return allocations

function allocateProportionally(totalQuota, periods, usagePatterns):
  allocations = {}

  // Calculate proportional allocation based on period duration and usage patterns
  totalWeight = 0
  periodWeights = {}

  for period in periods:
    historicalUsageRatio = usagePatterns[period]?.averageUsageRatio || 0.1
    durationWeight = periods[period].duration / 86400000 // Normalize to days

    periodWeights[period] = historicalUsageRatio * durationWeight
    totalWeight += periodWeights[period]

  // Allocate quota proportionally
  for period in periods:
    if totalWeight > 0:
      proportionalAllocation = (periodWeights[period] / totalWeight) * totalQuota
      allocations[period] = Math.max(
        proportionalAllocation,
        periods[period].duration / 86400000 * 100 // Minimum daily equivalent
      )
    else:
      allocations[period] = periods[period].allocation

  return allocations

```

### Quota Optimization Engine:

```

function optimizeQuotaDistribution(userId, currentAllocations, usageMetrics):
  optimization = {
    suggestions: [],
    projectedSavings: 0,
    implementationPriority: []
  }

  // Analyze quota utilization efficiency
  for period in Object.keys(currentAllocations):
    utilizationRate = usageMetrics[period].used / currentAllocations[period]

    if utilizationRate < 0.3: // Less than 30% utilization
      suggestion = {
        type: 'reduce_allocation',
        period: period,
        currentAllocation: currentAllocations[period],
        suggestedAllocation: Math.ceil(currentAllocations[period] * 0.7),
        reasoning: 'low_utilization',
        expectedSavings: currentAllocations[period] * 0.3
      }
      optimization.suggestions.push(suggestion)
    else if utilizationRate > 0.9: // More than 90% utilization
      suggestion = {
        type: 'increase_allocation',
        period: period,
        currentAllocation: currentAllocations[period],
        suggestedAllocation: Math.ceil(currentAllocations[period] * 1.3),
        reasoning: 'high_utilization',
        riskMitigation: 'prevent_quota_exhaustion'
      }
      optimization.suggestions.push(suggestion)

  // Identify cross-period optimization opportunities
  crossPeriodOptimizations = identifyCrossPeriodOptimizations(currentAllocations, usageMetrics)
  optimization.suggestions.push(...crossPeriodOptimizations)

  // Calculate projected savings
  optimization.projectedSavings = optimization.suggestions
    .filter(s => s.expectedSavings)
    .reduce((total, s) => total + s.expectedSavings, 0)

  return optimization

```

## 5. Adaptive Rate Limiting Algorithm

□ [Back to Top](#)

---

**Purpose:** Automatically adjust rate limits based on system performance, user behavior, and traffic patterns.

### Machine Learning-Based Adaptation:

```
AdaptiveConfig = {
  learningWindow: 3600000,           // 1 hour learning window
  adaptationFrequency: 300000,       // Adjust every 5 minutes
  minAdjustmentThreshold: 0.05,     // 5% minimum change
  maxAdjustmentFactor: 2.0,         // Maximum 2x adjustment
  stabilityPeriod: 1800000,          // 30 minutes stability before adaptation

  features: [
    'error_rate',
    'response_time',
    'cpu_utilization',
    'memory_usage',
    'request_pattern',
    'user_behavior',
    'time_of_day',
    'day_of_week'
  ]
}

class AdaptiveRateLimiter:
  constructor(config):
    this.config = config
    this.baseRateLimits = new Map()
    this.currentRateLimits = new Map()
    this.performanceHistory = []
    this.lastAdaptation = Date.now()

  function adaptRateLimits():
    currentTime = Date.now()

    if currentTime - this.lastAdaptation < this.config.adaptationFrequency:
      return // Too soon to adapt

    // Collect current system metrics
    systemMetrics = this.collectSystemMetrics()
    userMetrics = this.collectUserMetrics()
```

```

// Prepare features for ML model
features = this.prepareFeatures(systemMetrics, userMetrics)

// Get adaptation recommendations
recommendations = this.mlModel.predict(features)

// Apply adaptations
for recommendation in recommendations:
    this.applyAdaptation(recommendation)

this.lastAdaptation = currentTime

function prepareFeatures(systemMetrics, userMetrics):
    return {
        error_rate: systemMetrics.errorRate,
        response_time: systemMetrics.avgResponseTime,
        cpu_utilization: systemMetrics.cpuUsage,
        memory_usage: systemMetrics.memoryUsage,
        request_pattern: this.analyzeRequestPattern(userMetrics.requestTimestamps),
        user_behavior: this.analyzeUserBehavior(userMetrics.userSessions),
        time_of_day: new Date().getHours(),
        day_of_week: new Date().getDay(),
        current_rate_limit: this.getCurrentRateLimit(),
        historical_performance: this.getHistoricalPerformance()
    }

function applyAdaptation(recommendation):
    adaptationFactor = Math.max(
        1 / this.config.maxAdjustmentFactor,
        Math.min(this.config.maxAdjustmentFactor, recommendation.factor)
    )

    // Only apply if change is significant
    if Math.abs(adaptationFactor - 1.0) < this.config.minAdjustmentThreshold:
        return

    for [identifier, currentLimit] of this.currentRateLimits:
        newLimit = Math.floor(currentLimit * adaptationFactor)

        // Ensure new limit is reasonable
        baseLimit = this.baseRateLimits.get(identifier)
        newLimit = Math.max(baseLimit * 0.1, Math.min(baseLimit * 5, newLimit))

        this.currentRateLimits.set(identifier, newLimit)

```

```

        // Log adaptation
        this.logAdaptation(identifier, currentLimit, newLimit, recommendation.reason)

function analyzeRequestPattern(timestamps):
    if timestamps.length < 2:
        return { pattern: 'insufficient_data', score: 0 }

    // Calculate request intervals
    intervals = []
    for i in range(1, timestamps.length):
        intervals.push(timestamps[i] - timestamps[i-1])

    // Analyze pattern regularity
    avgInterval = intervals.reduce((a, b) => a + b) / intervals.length
    variance = intervals.reduce((sum, interval) => sum + Math.pow(interval - avgInterval, 2))

    coefficientOfVariation = Math.sqrt(variance) / avgInterval

    if coefficientOfVariation < 0.5:
        return { pattern: 'regular', score: 1 - coefficientOfVariation }
    else if coefficientOfVariation < 1.0:
        return { pattern: 'moderate', score: 0.5 }
    else:
        return { pattern: 'bursty', score: coefficientOfVariation / 2 }

```

## Performance Optimizations

□ [Back to Top](#)

---

## Memory-Efficient Storage

□ [Back to Top](#)

---

### Optimized Data Structures:

```

StorageOptimization = {
    useCompactCounters: true,           // Bit-packed counters
    enableCompression: true,           // Compress stored data
    circularBuffers: true,             // Fixed-size buffers for history
    lazyCleaning: true                 // Clean expired data lazily
}

```

## Distributed Coordination

□ [Back to Top](#)

---

**Consensus-Free Approach:** - Eventually consistent counters - Local rate limiting with periodic sync - Probabilistic counting for high-scale scenarios - Gossip protocol for configuration updates

## Security Considerations


□ [Back to Top](#)

---

## Rate Limiting Security Framework

□ [Back to Top](#)

---



## Testing Strategy

□ [Back to Top](#)

---

## Load Testing

□ [Back to Top](#)

---

**High Volume Scenarios:** - Burst traffic testing (10x normal load) - Sustained high load testing - Rate limit bypass attempts - Distributed coordination under load

## Accuracy Testing

□ [Back to Top](#)

---

**Rate Limiting Precision:** - Algorithm accuracy validation - Distributed consistency testing  
- Edge case handling verification - Performance regression testing

## Trade-offs and Considerations

□ [Back to Top](#)

---

### Accuracy vs Performance

□ [Back to Top](#)

---

- **Exact counting:** Perfect accuracy vs high computational cost
- **Approximate counting:** Good accuracy vs better performance
- **Local vs distributed:** Speed vs consistency
- **Memory vs disk:** Fast access vs persistent storage

### Flexibility vs Simplicity

□ [Back to Top](#)

---

- **Complex policies:** Fine-grained control vs operational complexity
- **Dynamic adaptation:** Smart behavior vs predictable performance
- **Multiple algorithms:** Optimal choice vs system complexity
- **Hierarchical limits:** Comprehensive control vs configuration complexity

### Cost vs Reliability

□ [Back to Top](#)

---

- **High availability:** Fault tolerance vs infrastructure cost
- **Global distribution:** Low latency vs operational overhead
- **Real-time monitoring:** Quick response vs resource usage
- **Historical analytics:** Deep insights vs storage expenses

This API rate limiting system provides a comprehensive foundation for protecting APIs with features like multiple rate limiting algorithms, hierarchical policy management, adaptive behavior, and robust monitoring while maintaining high performance, accuracy, and security standards.