

Content Delivery Network (CDN) System

□ Table of Contents

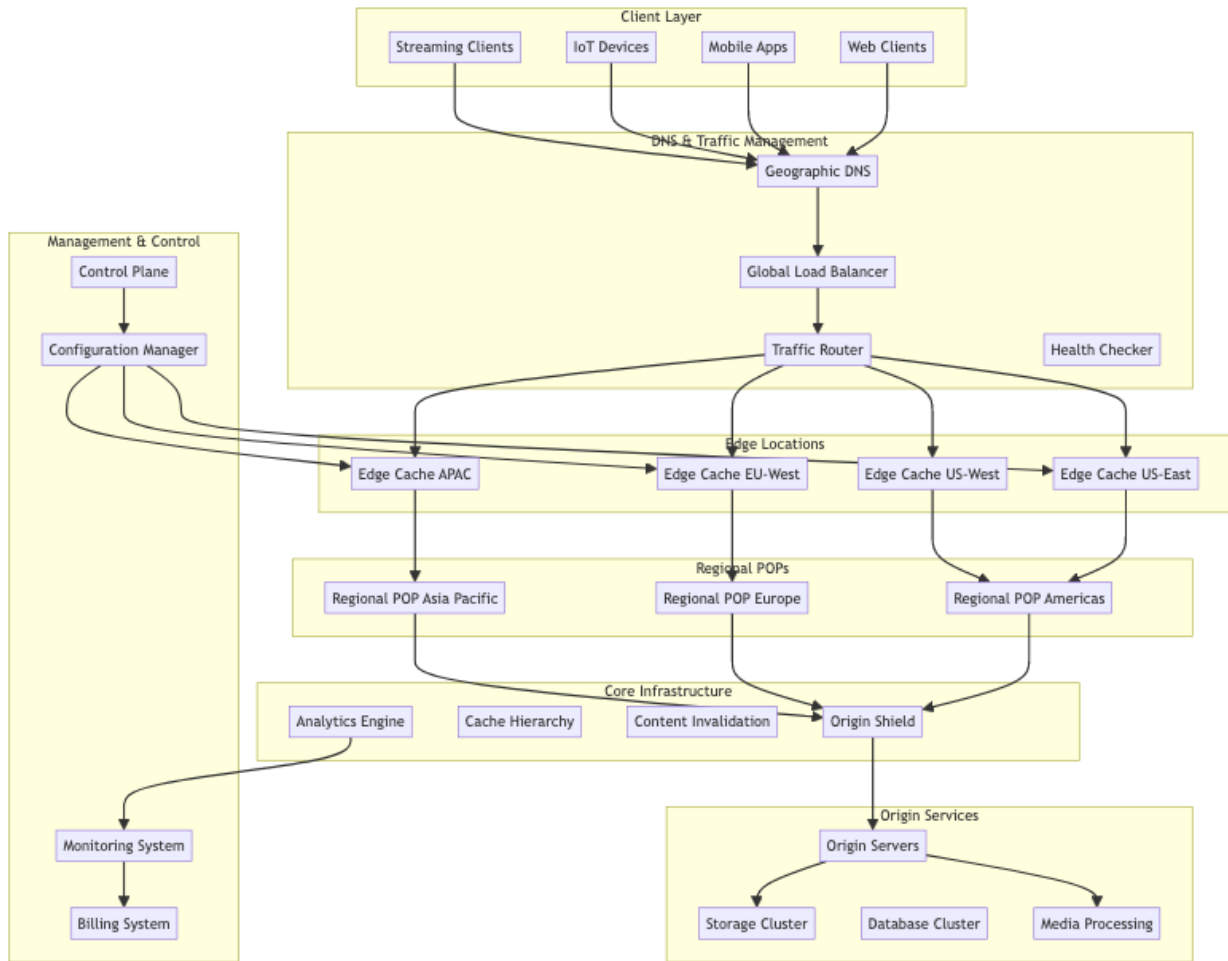
- Content Delivery Network (CDN) System
 - High-Level Design (HLD)
 - * System Architecture Overview
 - * Content Delivery Flow
 - Low-Level Design (LLD)
 - * Cache Architecture
 - * Geographic Routing Algorithm
 - * Content Invalidation Flow
 - Core Algorithms
 - * 1. Intelligent Content Caching Algorithm
 - * 2. Global Traffic Routing Algorithm
 - * 3. Content Optimization Algorithm
 - * 4. Content Invalidation Algorithm
 - * 5. Performance Analytics Algorithm
 - Performance Optimizations
 - * CDN Performance Strategies
 - * Global Infrastructure Optimization
 - Security Considerations
 - * CDN Security Framework
 - Testing Strategy
 - * CDN Testing Framework
 - * Load Testing
 - Trade-offs and Considerations
 - * Performance vs Cost
 - * Consistency vs Availability
 - * Security vs Performance
-

High-Level Design (HLD)

□ [Back to Top](#)

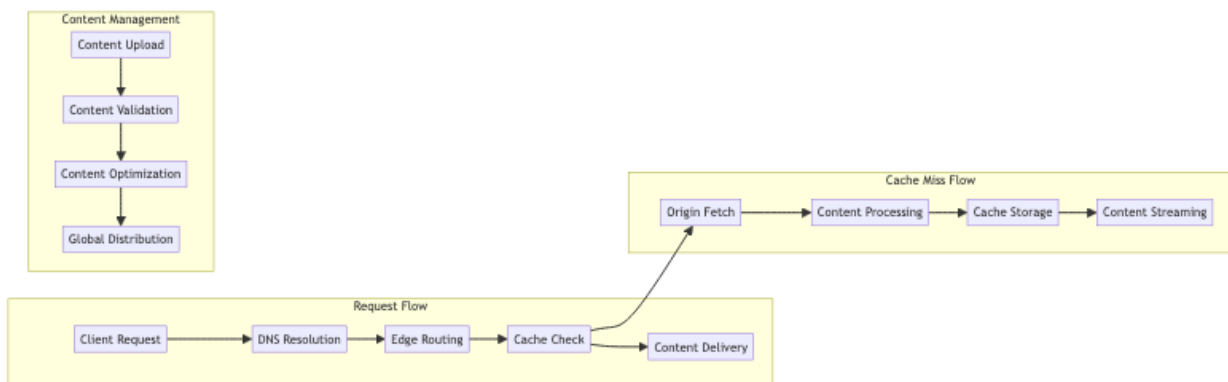
System Architecture Overview

□ [Back to Top](#)



Content Delivery Flow

□ Back to Top

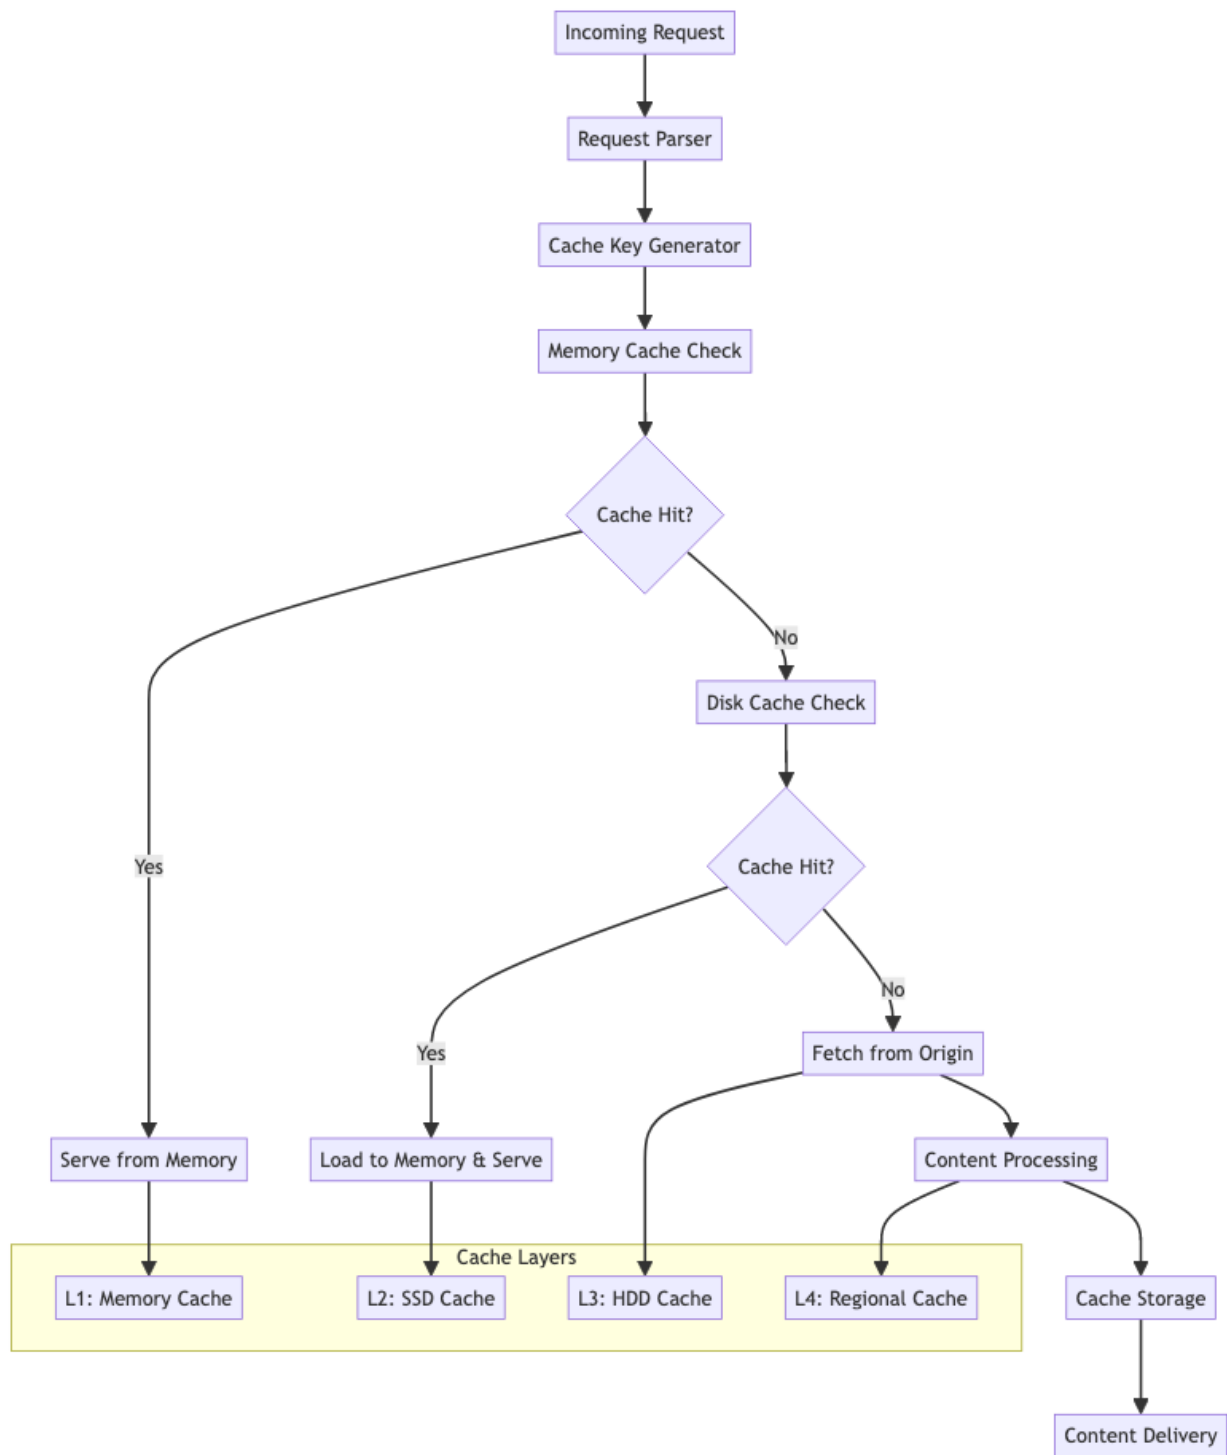


Low-Level Design (LLD)

□ [Back to Top](#)

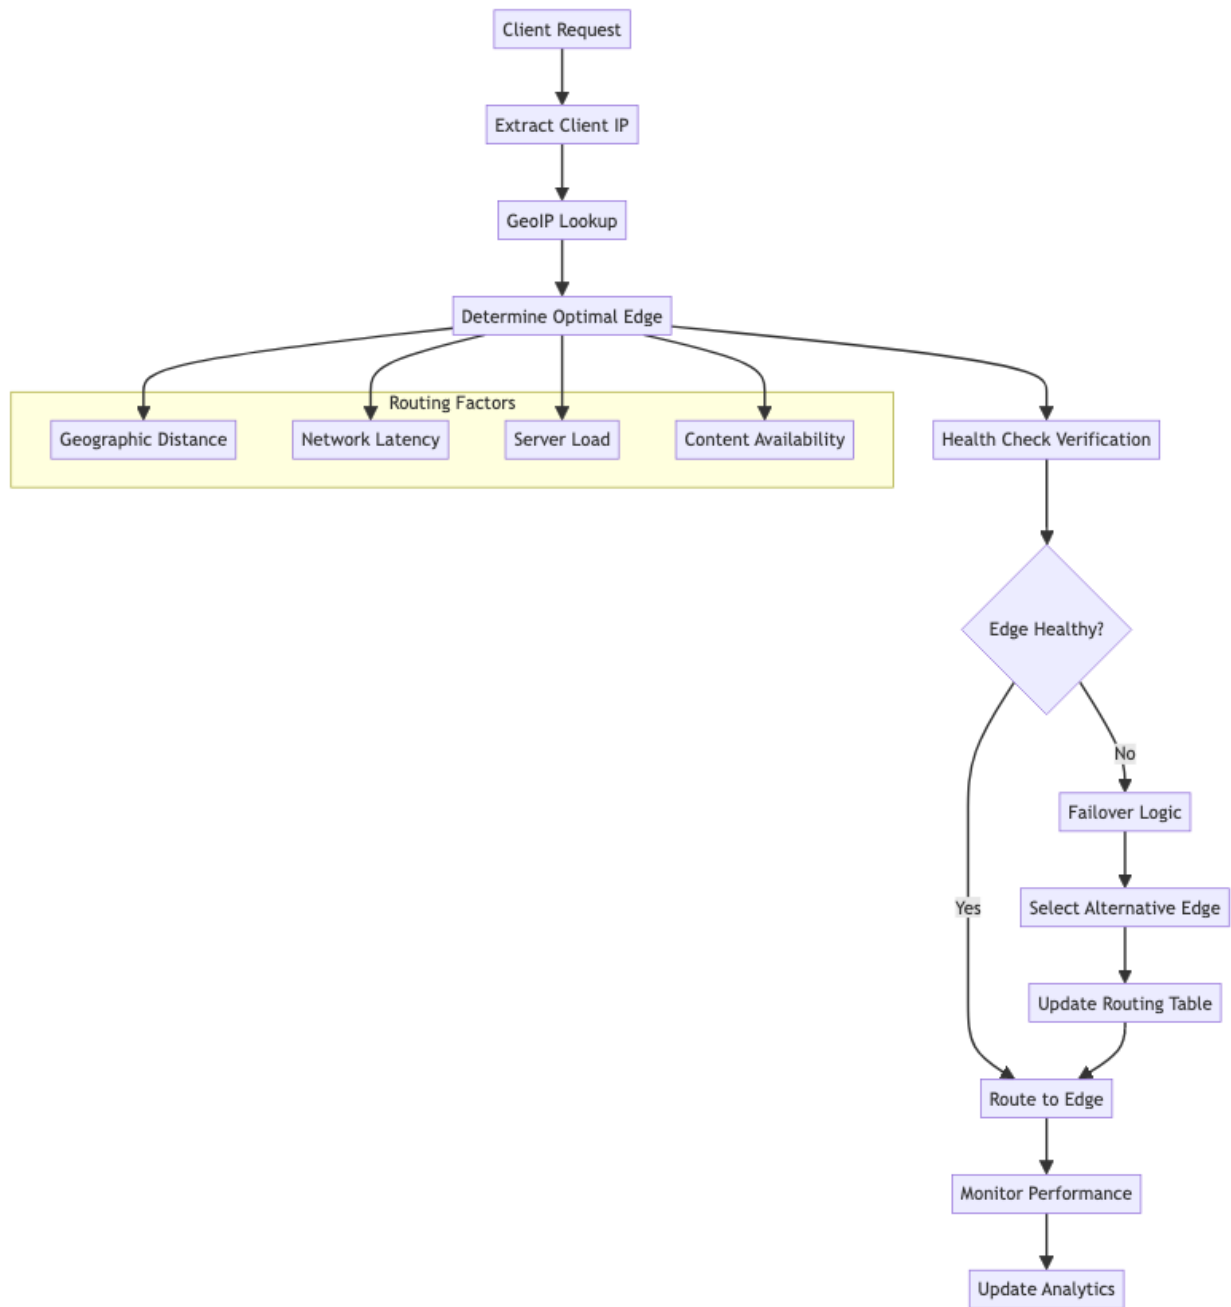
Cache Architecture

□ [Back to Top](#)



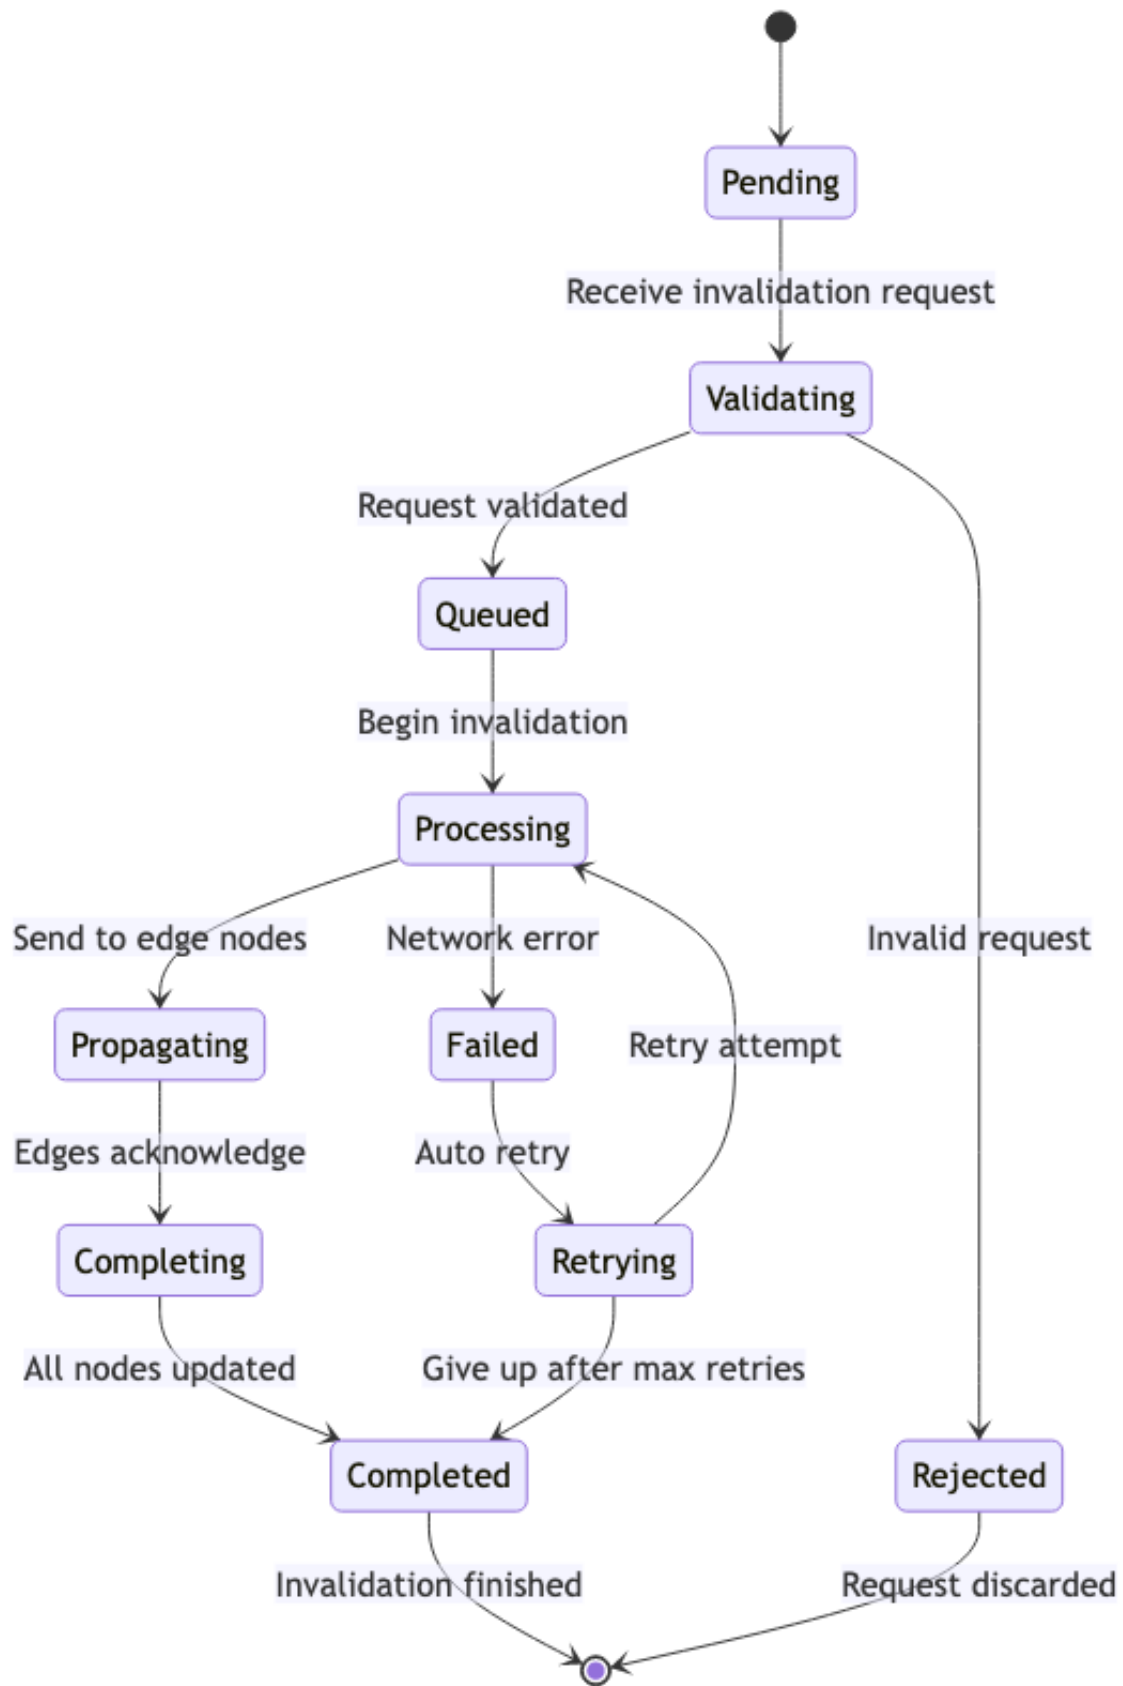
Geographic Routing Algorithm

□ [Back to Top](#)



Content Invalidation Flow

□ Back to Top



Core Algorithms

□ [Back to Top](#)

1. Intelligent Content Caching Algorithm

□ [Back to Top](#)

Purpose: Optimize cache hit rates by predicting content popularity and managing cache eviction based on multiple factors including geographic demand patterns.

Multi-Factor Cache Management:

```
CacheConfig = {
  caching: {
    levels: {
      memory: { size: '32GB', ttl: 3600000 },      # 1 hour
      ssd: { size: '1TB', ttl: 86400000 },         # 24 hours
      hdd: { size: '10TB', ttl: 604800000 }        # 7 days
    },

    evictionPolicy: 'adaptive_lru',                # 'lru', 'lfu', 'adaptive_lru'
    prefetchingEnabled: true,
    compressionEnabled: true,
    hotContentThreshold: 100                       # Requests per hour
  },

  prediction: {
    popularityWindow: 3600000,                    # 1 hour prediction window
    trendingThreshold: 5.0,                       # 5x increase in requests
    seasonalityDetection: true,
    geographicWeighting: true
  },

  optimization: {
    bandwidthThrottling: true,
    adaptiveQuality: true,
    contentCompression: true,
    imageOptimization: true
  }
}

class IntelligentCacheManager:
```

```

constructor(config):
    this.config = config
    this.memoryCache = new MemoryCache(config.caching.levels.memory)
    this.ssdCache = new SSDCache(config.caching.levels.ssd)
    this.hddCache = new HDDCache(config.caching.levels.hdd)
    this.popularityPredictor = new PopularityPredictor()
    this.geoAnalyzer = new GeographicAnalyzer()

function handleRequest(requestInfo):
    cacheKey = this.generateCacheKey(requestInfo)

    # Check cache hierarchy
    cacheResult = this.checkCacheHierarchy(cacheKey, requestInfo)

    if cacheResult.hit:
        # Update access patterns
        this.updateAccessPattern(cacheKey, requestInfo)
        return this.serveCachedContent(cacheResult)
    else:
        # Fetch from origin and cache
        return this.fetchAndCache(requestInfo, cacheKey)

function checkCacheHierarchy(cacheKey, requestInfo):
    # Level 1: Memory Cache
    memoryResult = this.memoryCache.get(cacheKey)
    if memoryResult:
        return {
            hit: true,
            level: 'memory',
            content: memoryResult.content,
            metadata: memoryResult.metadata,
            latency: 1 # ~1ms
        }

    # Level 2: SSD Cache
    ssdResult = this.ssdCache.get(cacheKey)
    if ssdResult:
        # Promote to memory cache if frequently accessed
        if this.shouldPromoteToMemory(cacheKey, requestInfo):
            this.memoryCache.set(cacheKey, ssdResult.content, ssdResult.metadata)

    return {
        hit: true,
        level: 'ssd',
        content: ssdResult.content,

```



```

        metadata: ssdResult.metadata,
        latency: 10 # ~10ms
    }

# Level 3: HDD Cache
hddResult = this.hddCache.get(cacheKey)
if hddResult:
    # Promote to SSD cache if warranted
    if this.shouldPromoteToSSD(cacheKey, requestInfo):
        this.ssdCache.set(cacheKey, hddResult.content, hddResult.metadata)

    return {
        hit: true,
        level: 'hdd',
        content: hddResult.content,
        metadata: hddResult.metadata,
        latency: 50 # ~50ms
    }

return { hit: false }

function fetchAndCache(requestInfo, cacheKey):
    fetchStart = Date.now()

    # Fetch from origin or upstream cache
    originResponse = this.fetchFromOrigin(requestInfo)

    if not originResponse.success:
        return this.handleOriginFailure(requestInfo)

    # Process and optimize content
    processedContent = this.processContent(originResponse.content, requestInfo)

    # Determine cache placement strategy
    cachePlacement = this.determineCachePlacement(cacheKey, processedContent, requestInfo)

    # Store in appropriate cache levels
    this.storeCachedContent(cacheKey, processedContent, cachePlacement)

    # Predict future demand and prefetch related content
    this.triggerPredictivePrefetch(cacheKey, requestInfo)

    return {
        content: processedContent.data,
        metadata: processedContent.metadata,

```

```

        cacheStatus: 'miss',
        origin: 'fetched',
        latency: Date.now() - fetchStart
    }
}

function determineCachePlacement(cacheKey, content, requestInfo):
    placement = {
        memory: false,
        ssd: false,
        hdd: false,
        priority: 'normal'
    }

    # Analyze content characteristics
    contentSize = content.size
    contentType = content.type

    # Analyze request patterns
    accessPattern = this.getAccessPattern(cacheKey)
    geographicDemand = this.geoAnalyzer.getGeographicDemand(cacheKey)

    # Predict popularity
    popularityScore = this.popularityPredictor.predict(cacheKey, requestInfo, accessPattern)

    # Small, frequently accessed content goes to memory
    if contentSize < 10485760 and popularityScore > 0.8: # 10MB and high popularity
        placement.memory = true
        placement.priority = 'high'

    # Medium-sized popular content goes to SSD
    if contentSize < 104857600 and popularityScore > 0.5: # 100MB and medium popularity
        placement.ssd = true
        placement.priority = popularityScore > 0.7 ? 'high' : 'normal'

    # All content goes to HDD cache (cold storage)
    placement.hdd = true

    # Geographic considerations
    if geographicDemand.isRegionallyPopular:
        placement.priority = 'high'
        placement.replicationFactor = geographicDemand.regions.length

    return placement

function storeCachedContent(cacheKey, content, placement):

```

```

metadata = {
    cacheKey: cacheKey,
    size: content.size,
    type: content.type,
    cachedAt: Date.now(),
    lastAccessed: Date.now(),
    accessCount: 1,
    etag: content.etag,
    expires: content.expires
}

# Store in appropriate cache levels
if placement.hdd:
    this.hddCache.set(cacheKey, content, metadata, placement.priority)

if placement.ssd:
    this.ssdCache.set(cacheKey, content, metadata, placement.priority)

if placement.memory:
    this.memoryCache.set(cacheKey, content, metadata, placement.priority)

# Log cache storage event
this.logCacheEvent('store', cacheKey, metadata, placement)

function shouldPromoteToMemory(cacheKey, requestInfo):
    accessPattern = this.getAccessPattern(cacheKey)

    # Promote if content is frequently accessed
    if accessPattern.hourlyAccess > this.config.caching.hotContentThreshold:
        return true

    # Promote if content is trending
    if accessPattern.trendingScore > this.config.prediction.trendingThreshold:
        return true

    # Promote if geographically concentrated demand
    geoDemand = this.geoAnalyzer.getGeographicDemand(cacheKey)
    if geoDemand.concentration > 0.8:
        return true

    return false

function triggerPredictivePrefetch(cacheKey, requestInfo):
    # Analyze related content that might be requested
    relatedContent = this.identifyRelatedContent(cacheKey, requestInfo)

```

```

for relatedItem in relatedContent:
    prefetchProbability = this.calculatePrefetchProbability(relatedItem, requestInfo)

    if prefetchProbability > 0.6: # 60% probability threshold
        this.schedulePrefetch(relatedItem, prefetchProbability)

function calculatePrefetchProbability(contentItem, requestInfo):
    # Base probability from historical patterns
    baseProbability = this.getHistoricalPrefetchSuccess(contentItem.id)

    # User behavior patterns
    userBehaviorScore = this.analyzeUserBehavior(requestInfo.userAgent, requestInfo.clientIp)

    # Content similarity
    similarityScore = this.calculateContentSimilarity(contentItem, requestInfo.requestedContent)

    # Time-based patterns
    temporalScore = this.getTemporalPrefetchScore(contentItem.id, Date.now())

    # Weighted combination
    probability = (
        baseProbability * 0.4 +
        userBehaviorScore * 0.3 +
        similarityScore * 0.2 +
        temporalScore * 0.1
    )

    return Math.min(1, Math.max(0, probability))

```

2. Global Traffic Routing Algorithm

□ [Back to Top](#)

Purpose: Route client requests to optimal edge locations based on geographic proximity, server load, content availability, and network conditions.

Anycast Routing with Load Balancing:

```

RoutingConfig = {
  routing: {
    algorithm: 'weighted_round_robin',      # 'round_robin', 'least_connections', 'weighted_round_robin'
    healthCheckInterval: 30000,              # 30 seconds
    failoverThreshold: 3,                   # 3 consecutive failures
    latencyThreshold: 200,                  # 200ms maximum acceptable latency
  }
}

```

```

        loadThreshold: 0.8                                # 80% load threshold
    },

    geographic: {
        regionPriority: true,
        continentFallback: true,
        crossOceanPenalty: 100,                            # Additional latency penalty for cross-ocean
        geoipAccuracy: 'city'                               # 'country', 'region', 'city'
    },

    optimization: {
        contentAwareRouting: true,
        networkConditionMonitoring: true,
        predictiveFailover: true,
        loadPrediction: true
    }
}

class GlobalTrafficRouter:
    constructor(config):
        this.config = config
        this.edgeLocations = new Map()
        this.geoipDatabase = new GeoIPDatabase()
        this.healthMonitor = new HealthMonitor()
        this.loadBalancer = new IntelligentLoadBalancer()
        this.networkMonitor = new NetworkConditionMonitor()

    function routeRequest(clientRequest):
        routingStart = Date.now()

        # Extract client information
        clientInfo = this.extractClientInfo(clientRequest)

        # Get available edge locations
        availableEdges = this.getAvailableEdgeLocations()

        # Score each edge location
        edgeScores = this.scoreEdgeLocations(availableEdges, clientInfo)

        # Select optimal edge
        selectedEdge = this.selectOptimalEdge(edgeScores, clientRequest)

        # Update routing metrics
        this.updateRoutingMetrics(selectedEdge, clientInfo, routingStart)

```

```

return {
    edgeLocation: selectedEdge,
    routingDecision: this.getRoutingDecision(edgeScores, selectedEdge),
    estimatedLatency: selectedEdge.estimatedLatency,
    routingTime: Date.now() - routingStart
}

function extractClientInfo(clientRequest):
    clientIP = this.extractClientIP(clientRequest)
    geoLocation = this.geoipDatabase.lookup(clientIP)

    return {
        ip: clientIP,
        country: geoLocation.country,
        region: geoLocation.region,
        city: geoLocation.city,
        latitude: geoLocation.latitude,
        longitude: geoLocation.longitude,
        isp: geoLocation.isp,
        asn: geoLocation.asn,
        userAgent: clientRequest.headers['user-agent'],
        acceptEncoding: clientRequest.headers['accept-encoding'],
        requestedContent: this.parseRequestedContent(clientRequest)
    }

function scoreEdgeLocations(edges, clientInfo):
    edgeScores = []

    for edge in edges:
        score = this.calculateEdgeScore(edge, clientInfo)
        edgeScores.push({
            edge: edge,
            score: score,
            factors: score.factors
        })

    # Sort by score (highest first)
    return edgeScores.sort((a, b) => b.score.total - a.score.total)

function calculateEdgeScore(edge, clientInfo):
    factors = {}

    # Geographic distance score
    factors.geographic = this.calculateGeographicScore(edge, clientInfo)

```

```

# Network latency score
factors.latency = this.calculateLatencyScore(edge, clientInfo)

# Server load score
factors.load = this.calculateLoadScore(edge)

# Content availability score
factors.content = this.calculateContentAvailabilityScore(edge, clientInfo.requestedC

# Network condition score
factors.network = this.calculateNetworkScore(edge, clientInfo)

# Health score
factors.health = this.calculateHealthScore(edge)

# Calculate weighted total score
totalScore = (
  factors.geographic * 0.25 +
  factors.latency * 0.20 +
  factors.load * 0.20 +
  factors.content * 0.15 +
  factors.network * 0.10 +
  factors.health * 0.10
)

return {
  total: totalScore,
  factors: factors
}

function calculateGeographicScore(edge, clientInfo):
  # Calculate great circle distance
  distance = this.calculateDistance(
    clientInfo.latitude, clientInfo.longitude,
    edge.latitude, edge.longitude
  )

  # Base score inversely proportional to distance
  baseScore = Math.max(0, 1 - (distance / 20003931)) # Normalize by half Earth circumf

  # Regional bonus
  regionalBonus = 0
  if edge.country === clientInfo.country:
    regionalBonus = 0.3
  else if edge.continent === clientInfo.continent:

```

```

        regionalBonus = 0.1

    # Cross-ocean penalty
    crossOceanPenalty = 0
    if this.isCrossOceanRouting(edge, clientInfo):
        crossOceanPenalty = 0.2

    return Math.min(1, Math.max(0, baseScore + regionalBonus - crossOceanPenalty))

function calculateLatencyScore(edge, clientInfo):
    # Get measured latency or estimate
    measuredLatency = this.networkMonitor.getLatency(edge.id, clientInfo.asn)

    if not measuredLatency:
        # Estimate latency based on distance and network path
        estimatedLatency = this.estimateLatency(edge, clientInfo)
        measuredLatency = estimatedLatency

    # Score based on latency threshold
    if measuredLatency <= 50:
        return 1.0 # Excellent latency
    else if measuredLatency <= 100:
        return 0.8 # Good latency
    else if measuredLatency <= this.config.routing.latencyThreshold:
        return 0.6 # Acceptable latency
    else:
        return 0.2 # Poor latency

function calculateLoadScore(edge):
    currentLoad = edge.currentLoad

    # Score inversely proportional to load
    if currentLoad <= 0.5:
        return 1.0 # Low load
    else if currentLoad <= 0.7:
        return 0.8 # Medium load
    else if currentLoad <= this.config.routing.loadThreshold:
        return 0.5 # High load
    else:
        return 0.1 # Overloaded

function calculateContentAvailabilityScore(edge, requestedContent):
    if not this.config.optimization.contentAwareRouting:
        return 1.0

```



```

# Check if content is already cached at edge
contentCached = edge.hasCachedContent(requestedContent.url)

if contentCached:
    return 1.0 # Content available locally

# Check cache hierarchy for content
hierarchyScore = this.calculateHierarchyAvailability(edge, requestedContent)

return hierarchyScore

function selectOptimalEdge(edgeScores, clientRequest):
    if edgeScores.length === 0:
        throw new Error('No available edge locations')

    # Primary selection: highest scored edge
    primaryEdge = edgeScores[0]

    # Validate primary edge can handle request
    if this.canHandleRequest(primaryEdge.edge, clientRequest):
        return primaryEdge.edge

    # Fallback selection
    for i in range(1, edgeScores.length):
        candidateEdge = edgeScores[i]

        if this.canHandleRequest(candidateEdge.edge, clientRequest):
            return candidateEdge.edge

    # Last resort: any available edge
    throw new Error('No edge location can handle the request')

function canHandleRequest(edge, clientRequest):
    # Check basic capacity
    if edge.currentLoad >= 0.95:
        return false

    # Check health status
    if not edge.healthy:
        return false

    # Check request size limits
    if clientRequest.contentSize > edge.maxRequestSize:
        return false

```

```

# Check supported content types
if not edge.supportedContentTypes.includes(clientRequest.contentType):
    return false

return true

function handleFailover(failedEdge, clientRequest):
    # Remove failed edge from available list temporarily
    this.markEdgeUnavailable(failedEdge.id)

    # Get alternative edges
    alternativeEdges = this.getAvailableEdgeLocations()
        .filter(edge => edge.id !== failedEdge.id)

    if alternativeEdges.length === 0:
        throw new Error('No alternative edge locations available')

    # Re-route to best alternative
    clientInfo = this.extractClientInfo(clientRequest)
    edgeScores = this.scoreEdgeLocations(alternativeEdges, clientInfo)

    selectedEdge = this.selectOptimalEdge(edgeScores, clientRequest)

    # Log failover event
    this.logFailoverEvent(failedEdge, selectedEdge, clientRequest)

    return selectedEdge

function updateRoutingMetrics(selectedEdge, clientInfo, routingStart):
    routingMetrics = {
        edgeId: selectedEdge.id,
        clientCountry: clientInfo.country,
        clientRegion: clientInfo.region,
        routingTime: Date.now() - routingStart,
        estimatedLatency: selectedEdge.estimatedLatency,
        timestamp: Date.now()
    }

    # Update edge load
    selectedEdge.requestCount++
    selectedEdge.currentLoad = this.calculateCurrentLoad(selectedEdge)

    # Update geographic routing statistics
    this.updateGeographicStats(selectedEdge.id, clientInfo)

```

```
# Store metrics for analysis
this.storeRoutingMetrics(routingMetrics)
```

3. Content Optimization Algorithm

□ [Back to Top](#)

Purpose: Optimize content delivery through compression, format conversion, image optimization, and adaptive quality selection based on client capabilities and network conditions.

Adaptive Content Optimization:

```
OptimizationConfig = {
  compression: {
    algorithms: ['gzip', 'brotli', 'zstd'],
    qualityLevels: [1, 6, 9],           # Compression levels
    minSizeThreshold: 1024,             # 1KB minimum for compression
    contentTypes: ['text/*', 'application/json', 'application/javascript']
  },

  images: {
    formats: ['webp', 'avif', 'jpeg', 'png'],
    qualityLevels: [50, 70, 85, 95],
    resolutions: [480, 720, 1080, 1440, 2160],
    enableAdaptive: true
  },

  video: {
    codecs: ['av1', 'h265', 'h264'],
    bitrates: [500, 1000, 2000, 4000, 8000], # kbps
    resolutions: ['480p', '720p', '1080p', '4k'],
    enableABR: true                      # Adaptive bitrate streaming
  }
}

class ContentOptimizationEngine:
  constructor(config):
    this.config = config
    this.imageProcessor = new ImageProcessor()
    this.videoProcessor = new VideoProcessor()
    this.compressionEngine = new CompressionEngine()
    this.formatConverter = new FormatConverter()
    this.qualityAnalyzer = new QualityAnalyzer()
```

```

function optimizeContent(content, clientInfo, networkConditions):
    optimizationStart = Date.now()

    # Analyze content characteristics
    contentAnalysis = this.analyzeContent(content)

    # Determine optimal transformations
    optimizations = this.planOptimizations(contentAnalysis, clientInfo, networkConditions)

    # Apply optimizations
    optimizedContent = this.applyOptimizations(content, optimizations)

    # Validate optimization results
    validationResult = this.validateOptimization(content, optimizedContent, optimizations)

    return {
        originalContent: content,
        optimizedContent: optimizedContent,
        optimizations: optimizations,
        metrics: {
            originalSize: content.size,
            optimizedSize: optimizedContent.size,
            compressionRatio: content.size / optimizedContent.size,
            optimizationTime: Date.now() - optimizationStart,
            qualityScore: validationResult.qualityScore
        }
    }

function analyzeContent(content):
    analysis = {
        type: content.type,
        size: content.size,
        format: this.detectFormat(content),
        characteristics: {}
    }

    if content.type.startsWith('image/'):
        analysis.characteristics = this.analyzeImage(content)
    else if content.type.startsWith('video/'):
        analysis.characteristics = this.analyzeVideo(content)
    else if content.type.startsWith('text/') or content.type.includes('javascript') or content.type.includes('application/'):
        analysis.characteristics = this.analyzeText(content)

    return analysis

```

```

function planOptimizations(contentAnalysis, clientInfo, networkConditions):
    optimizations = []

    # Compression optimization
    compressionPlan = this.planCompression(contentAnalysis, clientInfo, networkConditions)
    if compressionPlan:
        optimizations.push(compressionPlan)

    # Format optimization
    formatPlan = this.planFormatConversion(contentAnalysis, clientInfo)
    if formatPlan:
        optimizations.push(formatPlan)

    # Quality optimization
    qualityPlan = this.planQualityOptimization(contentAnalysis, clientInfo, networkConditions)
    if qualityPlan:
        optimizations.push(qualityPlan)

    # Resolution optimization
    resolutionPlan = this.planResolutionOptimization(contentAnalysis, clientInfo, networkConditions)
    if resolutionPlan:
        optimizations.push(resolutionPlan)

    return optimizations

function planCompression(contentAnalysis, clientInfo, networkConditions):
    # Skip compression for already compressed content
    if this.isCompressed(contentAnalysis.format):
        return null

    # Skip compression for small files
    if contentAnalysis.size < this.config.compression.minSizeThreshold:
        return null

    # Check if content type supports compression
    supportedType = this.config.compression.contentTypes.some(type =>
        contentAnalysis.type.match(new RegExp(type.replace('*', '.*')))
    )

    if not supportedType:
        return null

    # Select compression algorithm based on client support
    algorithm = this.selectCompressionAlgorithm(clientInfo.acceptEncoding)

```

```

# Select compression level based on network conditions
compressionLevel = this.selectCompressionLevel(networkConditions)

return {
  type: 'compression',
  algorithm: algorithm,
  level: compressionLevel,
  priority: this.calculateCompressionPriority(contentAnalysis, networkConditions)
}

function planFormatConversion(contentAnalysis, clientInfo):
  if contentAnalysis.type.startsWith('image/'):
    return this.planImageFormatConversion(contentAnalysis, clientInfo)
  else if contentAnalysis.type.startsWith('video/'):
    return this.planVideoFormatConversion(contentAnalysis, clientInfo)

  return null

function planImageFormatConversion(contentAnalysis, clientInfo):
  currentFormat = contentAnalysis.format

  # Determine best format based on client support
  supportedFormats = this.getSupportedImageFormats(clientInfo.userAgent)

  # Prefer modern formats for better compression
  preferredFormats = ['avif', 'webp', 'jpeg', 'png']

  for format in preferredFormats:
    if supportedFormats.includes(format) and format !== currentFormat:
      # Check if conversion would be beneficial
      expectedSavings = this.estimateFormatSavings(contentAnalysis, format)

      if expectedSavings > 0.1: # At least 10% savings
        return {
          type: 'format_conversion',
          sourceFormat: currentFormat,
          targetFormat: format,
          expectedSavings: expectedSavings,
          priority: 'medium'
        }

  return null

function planQualityOptimization(contentAnalysis, clientInfo, networkConditions):
  if not (contentAnalysis.type.startsWith('image/') or contentAnalysis.type.startsWith('video/')):
    return null

```

```

        return null

# Determine optimal quality based on network conditions
optimalQuality = this.calculateOptimalQuality(networkConditions, clientInfo)

# Check if quality reduction is needed
currentQuality = contentAnalysis.characteristics.quality || 85

if optimalQuality < currentQuality:
    return {
        type: 'quality_optimization',
        currentQuality: currentQuality,
        targetQuality: optimalQuality,
        expectedSavings: this.estimateQualitySavings(currentQuality, optimalQuality),
        priority: 'high'
    }

return null

function calculateOptimalQuality(networkConditions, clientInfo):
    baseQuality = 85 # Default quality

    # Adjust based on network speed
    if networkConditions.bandwidth < 1000000: # < 1 Mbps
        baseQuality = 60
    else if networkConditions.bandwidth < 5000000: # < 5 Mbps
        baseQuality = 70
    else if networkConditions.bandwidth > 25000000: # > 25 Mbps
        baseQuality = 95

    # Adjust based on device type
    if clientInfo.deviceType === 'mobile':
        baseQuality = Math.min(baseQuality, 80)
    else if clientInfo.deviceType === 'tablet':
        baseQuality = Math.min(baseQuality, 85)

    # Adjust based on screen resolution
    if clientInfo.screenResolution:
        if clientInfo.screenResolution.width <= 1920:
            baseQuality = Math.min(baseQuality, 85)

    return baseQuality

function applyOptimizations(content, optimizations):
    processedContent = { ...content }

```

```

# Sort optimizations by priority
sortedOptimizations = optimizations.sort((a, b) => {
  priorityOrder = { 'high': 3, 'medium': 2, 'low': 1 }
  return priorityOrder[b.priority] - priorityOrder[a.priority]
})

for optimization in sortedOptimizations:
  try:
    processedContent = this.applyOptimization(processedContent, optimization)
  catch error:
    this.logOptimizationError(optimization, error)
    continue

return processedContent

function applyOptimization(content, optimization):
  switch optimization.type:
    case 'compression':
      return this.applyCompression(content, optimization)

    case 'format_conversion':
      return this.applyFormatConversion(content, optimization)

    case 'quality_optimization':
      return this.applyQualityOptimization(content, optimization)

    case 'resolution_optimization':
      return this.applyResolutionOptimization(content, optimization)

    default:
      return content

function applyCompression(content, optimization):
  switch optimization.algorithm:
    case 'gzip':
      return this.compressionEngine.gzipCompress(content, optimization.level)

    case 'brotli':
      return this.compressionEngine.brotliCompress(content, optimization.level)

    case 'zstd':
      return this.compressionEngine.zstdCompress(content, optimization.level)

    default:

```



```

        throw new Error(`Unsupported compression algorithm: ${optimization.algorithm}`)
    }

    function applyImageFormatConversion(content, optimization):
        conversionOptions = {
            sourceFormat: optimization.sourceFormat,
            targetFormat: optimization.targetFormat,
            quality: optimization.quality || 85,
            preserveMetadata: false # Remove EXIF data for privacy and size
        }

        return this.formatConverter.convertImage(content, conversionOptions)

    function applyQualityOptimization(content, optimization):
        if content.type.startsWith('image/'):
            return this.imageProcessor.adjustQuality(content, optimization.targetQuality)
        else if content.type.startsWith('video/'):
            return this.videoProcessor.adjustQuality(content, optimization.targetQuality)

        return content

    function validateOptimization(originalContent, optimizedContent, optimizations):
        validation = {
            qualityScore: 1.0,
            sizeReduction: (originalContent.size - optimizedContent.size) / originalContent.size,
            errors: [],
            warnings: []
        }

        # Check file size reduction
        if optimizedContent.size >= originalContent.size:
            validation.warnings.push('Optimization did not reduce file size')

        # Check quality degradation for media files
        if originalContent.type.startsWith('image/') or originalContent.type.startsWith('video/'):
            qualityScore = this.calculateQualityScore(originalContent, optimizedContent)
            validation.qualityScore = qualityScore

            if qualityScore < 0.7:
                validation.warnings.push('Significant quality degradation detected')

        # Validate file integrity
        if not this.validateFileIntegrity(optimizedContent):
            validation.errors.push('Optimized content failed integrity check')

        return validation

```

```

function calculateQualityScore(originalContent, optimizedContent):
    if originalContent.type.startsWith('image/'):
        return this.calculateImageQualityScore(originalContent, optimizedContent)
    else if originalContent.type.startsWith('video/'):
        return this.calculateVideoQualityScore(originalContent, optimizedContent)

    return 1.0 # No quality metrics for non-media content

function generateAdaptiveVariants(content, clientProfiles):
    variants = []

    for profile in clientProfiles:
        try:
            # Generate optimized variant for this client profile
            optimizations = this.planOptimizations(
                this.analyzeContent(content),
                profile.clientInfo,
                profile.networkConditions
            )

            variant = this.applyOptimizations(content, optimizations)

            variants.push({
                profile: profile,
                content: variant,
                cacheKey: this.generateVariantCacheKey(content, profile),
                optimizations: optimizations
            })

        catch error:
            this.logVariantGenerationError(profile, error)
            continue

    return variants

```

4. Content Invalidation Algorithm

□ [Back to Top](#)

Purpose: Efficiently propagate content invalidation requests across the global CDN network while maintaining consistency and minimizing performance impact.

Distributed Cache Invalidation:

```

InvalidationConfig = {
  propagation: {
    method: 'cascade',          # 'broadcast', 'cascade', 'selective'
    batchSize: 1000,           # URLs per invalidation batch
    maxRetries: 3,
    retryDelay: 5000,          # 5 seconds
    timeout: 30000              # 30 seconds per node
  },

  priority: {
    critical: { ttl: 0, timeout: 5000 }, # Immediate invalidation
    high: { ttl: 60000, timeout: 15000 }, # 1 minute TTL
    normal: { ttl: 300000, timeout: 30000 }, # 5 minutes TTL
    low: { ttl: 3600000, timeout: 60000 } # 1 hour TTL
  },

  validation: {
    urlPatternValidation: true,
    authorizedSourceOnly: true,
    rateLimiting: { maxPerHour: 10000, maxPerMinute: 500 }
  }
}

class ContentInvalidationManager:
  constructor(config):
    this.config = config
    this.edgeNodes = new Map()
    this.invalidationQueue = new PriorityQueue()
    this.propagationTracker = new PropagationTracker()
    this.rateLimiter = new RateLimiter(config.validation.rateLimiting)

  function invalidateContent(invalidationRequest):
    invalidationStart = Date.now()

    # Validate invalidation request
    validation = this.validateInvalidationRequest(invalidationRequest)
    if not validation.isValid:
      return { success: false, errors: validation.errors }

    # Check rate limits
    rateLimitCheck = this.rateLimiter.checkLimit(invalidationRequest.source)
    if not rateLimitCheck.allowed:
      return { success: false, error: 'rate_limit_exceeded', retryAfter: rateLimitCheck.retryAfter }

    # Create invalidation job

```

```

invalidationJob = this.createInvalidationJob(invalidationRequest)

# Add to queue with appropriate priority
this.invalidationQueue.enqueue(invalidationJob, invalidationJob.priority)

# Process invalidation
result = this.processInvalidation(invalidationJob)

return {
  success: result.success,
  invalidationId: invalidationJob.id,
  affectedNodes: result.affectedNodes,
  propagationTime: Date.now() - invalidationStart,
  status: result.status
}

function createInvalidationJob(request):
  return {
    id: this.generateInvalidationId(),
    urls: this.normalizeUrls(request.urls),
    priority: this.determinePriority(request),
    source: request.source,
    reason: request.reason || 'manual',
    createdAt: Date.now(),

    # Invalidation strategy
    method: request.method || this.config.propagation.method,
    batchSize: request.batchSize || this.config.propagation.batchSize,

    # Tracking
    status: 'pending',
    progress: 0,
    affectedNodes: [],
    errors: []
  }

function processInvalidation(invalidationJob):
  invalidationJob.status = 'processing'

  try:
    # Determine target nodes for invalidation
    targetNodes = this.selectTargetNodes(invalidationJob.urls)

    # Group URLs into batches
    urlBatches = this.createUrlBatches(invalidationJob.urls, invalidationJob.batchSize

```

```

    # Execute invalidation across nodes
    propagationResult = this.propagateInvalidation(targetNodes, urlBatches, invalidationJob)

    # Update job status
    invalidationJob.status = propagationResult.success ? 'completed' : 'failed'
    invalidationJob.affectedNodes = propagationResult.affectedNodes
    invalidationJob.errors = propagationResult.errors

    return propagationResult

catch error:
    invalidationJob.status = 'failed'
    invalidationJob.errors.push(error.message)

    return {
        success: false,
        error: error.message,
        affectedNodes: [],
        status: 'failed'
    }

function selectTargetNodes(urls):
    targetNodes = new Set()

    for url in urls:
        # Find nodes that potentially have this content cached
        nodesWithContent = this.findNodesWithContent(url)

        for node in nodesWithContent:
            targetNodes.add(node)

    # Also include all edge nodes for pattern-based invalidations
    if this.hasWildcardPatterns(urls):
        allEdgeNodes = this.getAllEdgeNodes()
        for node in allEdgeNodes:
            targetNodes.add(node)

    return Array.from(targetNodes)

function propagateInvalidation(targetNodes, urlBatches, invalidationJob):
    propagationResults = {
        success: true,
        affectedNodes: [],
        errors: [],

```

```

    totalRequests: 0,
    successfulRequests: 0
  }

  # Propagate to nodes in parallel
  propagationPromises = targetNodes.map(node =>
    this.invalidateNode(node, urlBatches, invalidationJob)
  )

  # Wait for all propagations to complete
  results = await Promise.allSettled(propagationPromises)

  for i, result in results.entries():
    node = targetNodes[i]

    if result.status === 'fulfilled':
      propagationResults.affectedNodes.push(node)
      propagationResults.successfulRequests += result.value.successfulBatches
      propagationResults.totalRequests += result.value.totalBatches
    else:
      propagationResults.errors.push({
        node: node,
        error: result.reason
      })
      propagationResults.success = false

  return propagationResults

function invalidateNode(node, urlBatches, invalidationJob):
  return new Promise((resolve, reject) => {
    nodeInvalidation = {
      nodeId: node.id,
      totalBatches: urlBatches.length,
      successfulBatches: 0,
      errors: []
    }

    # Process batches sequentially to avoid overwhelming the node
    this.processBatchesSequentially(node, urlBatches, nodeInvalidation)
      .then(result => {
        # Track propagation success
        this.propagationTracker.recordNodeSuccess(invalidationJob.id, node.id, result)
        resolve(nodeInvalidation)
      })
      .catch(error => {

```

```

        # Track propagation failure
        this.propagationTracker.recordNodeFailure(invalidationJob.id, node.id, error)
        reject(error)
    })
})

function processBatchesSequentially(node, urlBatches, nodeInvalidation):
    return new Promise(async (resolve, reject) => {
        for batch in urlBatches:
            try:
                batchResult = await this.invalidateBatch(node, batch)

                if batchResult.success:
                    nodeInvalidation.successfulBatches++
                else:
                    nodeInvalidation.errors.push(batchResult.error)

                # Add delay between batches to prevent overwhelming
                if batch !== urlBatches[urlBatches.length - 1]:
                    await this.delay(100) # 100ms delay

            catch error:
                nodeInvalidation.errors.push({
                    batch: batch,
                    error: error.message
                })

        resolve(nodeInvalidation)
    })

function invalidateBatch(node, urlBatch):
    return new Promise((resolve, reject) => {
        invalidationRequest = {
            type: 'cache_invalidation',
            urls: urlBatch,
            timestamp: Date.now(),
            requestId: this.generateRequestId()
        }

        # Set timeout for node response
        timeout = setTimeout(() => {
            reject(new Error(`Node ${node.id} invalidation timeout`))
        }, this.config.propagation.timeout)

        # Send invalidation request to node

```

```

node.sendInvalidationRequest(invalidationRequest)
  .then(response => {
    clearTimeout(timeout)

    if response.success:
      resolve({
        success: true,
        invalidatedUrls: response.invalidatedUrls,
        responseTime: response.responseTime
      })
    else:
      reject(new Error(response.error))
  })
  .catch(error => {
    clearTimeout(timeout)
    reject(error)
  })
})

function validateInvalidationRequest(request):
  errors = []

  # Check required fields
  if not request.urls or request.urls.length === 0:
    errors.push('URLs array is required and cannot be empty')

  if not request.source:
    errors.push('Source identifier is required')

  # Validate URL patterns
  if this.config.validation.urlPatternValidation:
    for url in request.urls:
      if not this.isValidUrlPattern(url):
        errors.push(`Invalid URL pattern: ${url}`)

  # Check authorization
  if this.config.validation.authorizedSourceOnly:
    if not this.isAuthorizedSource(request.source):
      errors.push('Unauthorized invalidation source')

  # Check URL count limits
  if request.urls.length > this.config.propagation.batchSize * 10:
    errors.push(`Too many URLs in single request. Maximum: ${this.config.propagation.b

  return {

```



```

        isValid: errors.length === 0,
        errors: errors
    }
}

function findNodesWithContent(url):
    nodesWithContent = []

    # Check cache directories for URL presence
    for [nodeId, node] in this.edgeNodes:
        if node.hasCachedContent(url):
            nodesWithContent.push(node)

    # For pattern-based URLs, check more broadly
    if this.isWildcardPattern(url):
        matchingNodes = this.findNodesWithPatternMatch(url)
        nodesWithContent.push(...matchingNodes)

    return nodesWithContent

function determinePriority(request):
    # Check for explicit priority
    if request.priority:
        return request.priority

    # Determine priority based on URL patterns
    for url in request.urls:
        if this.isCriticalContent(url):
            return 'critical'

        if this.isHighPriorityContent(url):
            return 'high'

    # Check invalidation reason
    if request.reason === 'security' or request.reason === 'emergency':
        return 'critical'

    if request.reason === 'content_update':
        return 'high'

    return 'normal'

function createUrlBatches(urls, batchSize):
    batches = []

    for i in range(0, urls.length, batchSize):

```

```

        batch = urls.slice(i, i + batchSize)
        batches.push(batch)

    return batches

function getInvalidationStatus(invalidationId):
    # Get job from tracking system
    job = this.propagationTracker.getJob(invalidationId)

    if not job:
        return { status: 'not_found' }

    # Calculate progress
    totalNodes = job.targetNodes.length
    completedNodes = job.completedNodes.length
    progress = totalNodes > 0 ? (completedNodes / totalNodes) * 100 : 0

    return {
        invalidationId: invalidationId,
        status: job.status,
        progress: progress,
        createdAt: job.createdAt,
        completedAt: job.completedAt,
        affectedNodes: job.affectedNodes,
        errors: job.errors,

        # Statistics
        totalUrls: job.urls.length,
        totalNodes: totalNodes,
        successfulNodes: job.completedNodes.length,
        failedNodes: job.errors.length
    }

```

5. Performance Analytics Algorithm

□ [Back to Top](#)

Purpose: Collect, analyze, and provide insights on CDN performance metrics including cache hit rates, latency, bandwidth usage, and content popularity.

Real-time Analytics Engine:

```

AnalyticsConfig = {
    collection: {
        samplingRate: 0.1,
        # 10% sampling for detailed analytics
    }
}

```

```

    metricsRetention: 2592000000,          # 30 days retention
    realTimeWindow: 300000,               # 5 minutes real-time window
    aggregationIntervals: [60000, 300000, 3600000, 86400000] # 1min, 5min, 1hour, 1day
},

metrics: {
    performance: ['latency', 'throughput', 'cache_hit_rate', 'error_rate'],
    business: ['bandwidth_usage', 'request_count', 'unique_visitors', 'geographic_distribution'],
    technical: ['origin_load', 'edge_utilization', 'storage_usage', 'network_quality']
},

alerting: {
    thresholds: {
        cache_hit_rate: { warning: 0.8, critical: 0.7 },
        latency_p95: { warning: 200, critical: 500 },
        error_rate: { warning: 0.01, critical: 0.05 }
    }
}
}

class PerformanceAnalyticsEngine:
    constructor(config):
        this.config = config
        this.metricsCollector = new MetricsCollector()
        this.timeSeriesDB = new TimeSeriesDatabase()
        this.alertManager = new AlertManager()
        this.reportGenerator = new ReportGenerator()
        this.anomalyDetector = new AnomalyDetector()

    function collectRequestMetrics(requestInfo, responseInfo):
        # Sample request for detailed analytics
        if not this.shouldSampleRequest():
            return

        metrics = {
            timestamp: Date.now(),
            requestId: requestInfo.id,

            # Request characteristics
            url: requestInfo.url,
            method: requestInfo.method,
            contentType: requestInfo.contentType,
            contentSize: responseInfo.contentSize,

            # Client information

```

```

    clientIP: requestInfo.clientIP,
    userAgent: requestInfo.userAgent,
    country: requestInfo.country,
    region: requestInfo.region,

    # Performance metrics
    latency: responseInfo.latency,
    cacheStatus: responseInfo.cacheStatus, # hit, miss, stale
    edgeLocation: responseInfo.edgeLocation,

    # Network metrics
    bandwidth: responseInfo.bandwidth,
    transferTime: responseInfo.transferTime,

    # Error information
    statusCode: responseInfo.statusCode,
    errorType: responseInfo.errorType
}

# Store metrics
this.storeMetrics(metrics)

# Update real-time aggregations
this.updateRealTimeMetrics(metrics)

# Check for anomalies
this.checkAnomalies(metrics)

function generatePerformanceReport(timeRange, granularity = 'hour'):
    reportStart = Date.now()

    # Collect metrics for time range
    rawMetrics = this.timeSeriesDB.query({
        timeRange: timeRange,
        granularity: granularity,
        metrics: this.config.metrics.performance
    })

    # Calculate performance statistics
    performanceStats = this.calculatePerformanceStats(rawMetrics)

    # Generate geographic analysis
    geographicAnalysis = this.analyzeGeographicPerformance(rawMetrics)

    # Calculate cache efficiency

```

```

    cacheEfficiency = this.analyzeCacheEfficiency(rawMetrics)

    # Identify top content
    topContent = this.analyzeContentPopularity(rawMetrics)

    # Detect performance trends
    trends = this.detectPerformanceTrends(rawMetrics)

    # Generate insights and recommendations
    insights = this.generateInsights(performanceStats, trends)

    return {
      reportId: this.generateReportId(),
      timeRange: timeRange,
      granularity: granularity,
      generatedAt: Date.now(),
      generationTime: Date.now() - reportStart,

      performanceStats: performanceStats,
      geographicAnalysis: geographicAnalysis,
      cacheEfficiency: cacheEfficiency,
      topContent: topContent,
      trends: trends,
      insights: insights
    }
  }

function calculatePerformanceStats(rawMetrics):
  stats = {
    latency: this.calculateLatencyStats(rawMetrics),
    throughput: this.calculateThroughputStats(rawMetrics),
    cacheHitRate: this.calculateCacheHitRate(rawMetrics),
    errorRate: this.calculateErrorRate(rawMetrics),
    bandwidth: this.calculateBandwidthStats(rawMetrics)
  }

  return stats

function calculateLatencyStats(rawMetrics):
  latencyValues = rawMetrics
    .filter(m => m.latency != null)
    .map(m => m.latency)
    .sort((a, b) => a - b)

  if latencyValues.length === 0:
    return null

```

```

return {
  min: latencyValues[0],
  max: latencyValues[latencyValues.length - 1],
  mean: latencyValues.reduce((sum, val) => sum + val, 0) / latencyValues.length,
  median: this.calculatePercentile(latencyValues, 50),
  p95: this.calculatePercentile(latencyValues, 95),
  p99: this.calculatePercentile(latencyValues, 99),
  count: latencyValues.length
}

function calculateCacheHitRate(rawMetrics):
  cacheStatuses = rawMetrics.map(m => m.cacheStatus)

  hitCount = cacheStatuses.filter(status => status === 'hit').length
  missCount = cacheStatuses.filter(status => status === 'miss').length
  staleCount = cacheStatuses.filter(status => status === 'stale').length

  totalRequests = hitCount + missCount + staleCount

  if totalRequests === 0:
    return null

  return {
    hitRate: hitCount / totalRequests,
    missRate: missCount / totalRequests,
    staleRate: staleCount / totalRequests,
    totalRequests: totalRequests,
    breakdown: {
      hits: hitCount,
      misses: missCount,
      stale: staleCount
    }
  }

function analyzeGeographicPerformance(rawMetrics):
  # Group metrics by geographic region
  regionMetrics = new Map()

  for metric in rawMetrics:
    region = metric.region || 'unknown'

    if not regionMetrics.has(region):
      regionMetrics.set(region, [])

```

```

        regionMetrics.get(region).push(metric)

# Calculate performance stats for each region
regionPerformance = []

for [region, metrics] in regionMetrics:
    regionStats = {
        region: region,
        requestCount: metrics.length,
        latency: this.calculateLatencyStats(metrics),
        cacheHitRate: this.calculateCacheHitRate(metrics),
        errorRate: this.calculateErrorRate(metrics),
        topEdgeLocations: this.getTopEdgeLocations(metrics)
    }

    regionPerformance.push(regionStats)

# Sort by request count
regionPerformance.sort((a, b) => b.requestCount - a.requestCount)

return {
    regions: regionPerformance,
    totalRegions: regionPerformance.length,
    globalStats: this.calculatePerformanceStats(rawMetrics)
}

function analyzeCacheEfficiency(rawMetrics):
    # Analyze cache performance by content type
    contentTypeAnalysis = this.analyzeCacheByContentType(rawMetrics)

    # Analyze cache performance by URL patterns
    urlPatternAnalysis = this.analyzeCacheByUrlPattern(rawMetrics)

    # Identify cache optimization opportunities
    optimizationOpportunities = this.identifyOptimizationOpportunities(rawMetrics)

    return {
        overall: this.calculateCacheHitRate(rawMetrics),
        byContentType: contentTypeAnalysis,
        byUrlPattern: urlPatternAnalysis,
        optimizationOpportunities: optimizationOpportunities
    }

function analyzeContentPopularity(rawMetrics):
    # Group requests by URL

```

```

urlCounts = new Map()
urlMetrics = new Map()

for metric in rawMetrics:
    url = metric.url

    # Count requests
    urlCounts.set(url, (urlCounts.get(url) || 0) + 1)

    # Collect metrics
    if not urlMetrics.has(url):
        urlMetrics.set(url, [])
    urlMetrics.get(url).push(metric)

# Calculate statistics for each URL
urlStats = []

for [url, count] in urlCounts:
    metrics = urlMetrics.get(url)

    urlStats.push({
        url: url,
        requestCount: count,
        contentType: metrics[0].contentType,
        averageSize: metrics.reduce((sum, m) => sum + (m.contentSize || 0), 0) / metrics.length,
        cacheHitRate: this.calculateCacheHitRate(metrics).hitRate,
        averageLatency: this.calculateLatencyStats(metrics).mean,
        bandwidth: count * (metrics[0].contentSize || 0),
        uniqueClients: new Set(metrics.map(m => m.clientIP)).size
    })

# Sort by request count
urlStats.sort((a, b) => b.requestCount - a.requestCount)

return {
    topUrls: urlStats.slice(0, 100),
    totalUrls: urlStats.length,
    topContentTypes: this.getTopContentTypes(urlStats)
}

function detectPerformanceTrends(rawMetrics):
    # Group metrics by time intervals
    timeIntervals = this.groupByTimeInterval(rawMetrics, 3600000) # 1 hour intervals

    # Calculate trends for key metrics

```



```

trends = {
  latency: this.calculateTrend(timeIntervals, 'latency'),
  cacheHitRate: this.calculateTrend(timeIntervals, 'cacheHitRate'),
  throughput: this.calculateTrend(timeIntervals, 'throughput'),
  errorRate: this.calculateTrend(timeIntervals, 'errorRate')
}

return trends

function calculateTrend(timeIntervals, metric):
  # Extract metric values over time
  dataPoints = timeIntervals.map(interval => ({
    timestamp: interval.timestamp,
    value: this.calculateIntervalMetric(interval.metrics, metric)
  }))

  if dataPoints.length < 2:
    return { trend: 'insufficient_data' }

  # Calculate linear regression
  regression = this.linearRegression(dataPoints)

  # Determine trend direction
  trendDirection = 'stable'
  if regression.slope > 0.05:
    trendDirection = 'increasing'
  else if regression.slope < -0.05:
    trendDirection = 'decreasing'

  return {
    trend: trendDirection,
    slope: regression.slope,
    correlation: regression.correlation,
    dataPoints: dataPoints,
    prediction: this.predictNextValue(regression, dataPoints)
  }

function generateInsights(performanceStats, trends):
  insights = []

  # Cache hit rate insights
  if performanceStats.cacheHitRate.hitRate < 0.8:
    insights.push({
      type: 'performance',
      severity: 'warning',

```

```

        title: 'Low Cache Hit Rate',
        description: `Cache hit rate is ${performanceStats.cacheHitRate.hitRate * 100}.`,
        recommendation: 'Review cache TTL settings and consider preloading popular content'
    })

# Latency insights
if performanceStats.latency.p95 > 200:
    insights.push({
        type: 'performance',
        severity: 'warning',
        title: 'High Latency',
        description: `95th percentile latency is ${performanceStats.latency.p95}ms, above threshold`,
        recommendation: 'Investigate edge location optimization and origin server performance'
    })

# Trend insights
if trends.errorRate.trend === 'increasing':
    insights.push({
        type: 'reliability',
        severity: 'critical',
        title: 'Increasing Error Rate',
        description: 'Error rate is trending upward, indicating potential system issues',
        recommendation: 'Investigate error sources and implement additional monitoring'
    })

return insights

function checkAnomalies(metrics):
    # Check for latency anomalies
    if this.anomalyDetector.detectLatencyAnomaly(metrics.latency):
        this.alertManager.triggerAlert({
            type: 'performance_anomaly',
            severity: 'warning',
            metric: 'latency',
            value: metrics.latency,
            threshold: this.anomalyDetector.getLatencyThreshold(),
            edgeLocation: metrics.edgeLocation
        })

    # Check for cache hit rate anomalies
    currentHitRate = this.getCurrentCacheHitRate()
    if currentHitRate < this.config.alerting.thresholds.cache_hit_rate.critical:
        this.alertManager.triggerAlert({
            type: 'cache_performance',
            severity: 'critical',

```

```

        metric: 'cache_hit_rate',
        value: currentHitRate,
        threshold: this.config.alerting.thresholds.cache_hit_rate.critical
    })

    # Check for error rate spikes
    currentErrorRate = this.getCurrentErrorRate()
    if currentErrorRate > this.config.alerting.thresholds.error_rate.warning:
        this.alertManager.triggerAlert({
            type: 'error_rate_spike',
            severity: currentErrorRate > this.config.alerting.thresholds.error_rate.critical
            metric: 'error_rate',
            value: currentErrorRate,
            threshold: this.config.alerting.thresholds.error_rate.warning
        })

```

Performance Optimizations

[□ Back to Top](#)

CDN Performance Strategies

[□ Back to Top](#)

Multi-tier Optimization Framework:

```

PerformanceOptimization = {
  caching: {
    memoryTiers: ['L1', 'L2', 'L3'],
    compressionRatio: 0.7,
    prefetchingEnabled: true,
    hotDataPromotion: true
  },

  networking: {
    httpVersion: 'http2',
    connectionReuse: true,
    multiplexing: true,
    serverPush: true
  },

  content: {

```

```

    imageOptimization: true,
    videoAdaptiveStreaming: true,
    minification: true,
    bundling: true
  }
}

```

Global Infrastructure Optimization

□ [Back to Top](#)


Edge Location Strategy: - Geographic distribution optimization - Network topology aware placement - Dynamic capacity scaling - Cross-region failover mechanisms

Security Considerations

□ [Back to Top](#)

CDN Security Framework

□ [Back to Top](#)



Testing Strategy

□ [Back to Top](#)

CDN Testing Framework

□ [Back to Top](#)

Comprehensive Testing Approach: - Performance testing under various load conditions
- Geographic testing from multiple regions - Cache invalidation testing - Failover and disaster recovery testing - Security penetration testing

Load Testing

□ [Back to Top](#)

Distributed Testing Environment: - Multi-region load generation - Realistic traffic pattern simulation - Cache behavior validation - Origin shielding effectiveness

Trade-offs and Considerations

□ [Back to Top](#)

Performance vs Cost

□ [Back to Top](#)

- **Edge location density:** Lower latency vs infrastructure cost
- **Cache storage:** Hit rate improvement vs storage expenses
- **Bandwidth:** Performance vs data transfer costs
- **Global reach:** Worldwide coverage vs operational complexity

Consistency vs Availability

□ [Back to Top](#)

- **Cache invalidation:** Consistency vs propagation speed
- **Content updates:** Real-time updates vs cache efficiency
- **Geographic distribution:** Data consistency vs local performance
- **Failover mechanisms:** Availability vs potential inconsistency

Security vs Performance

□ [Back to Top](#)

- **Content validation:** Security thoroughness vs processing speed
- **Access controls:** Security rigor vs user experience

- **DDoS protection:** Attack mitigation vs legitimate traffic impact
- **Encryption overhead:** Security vs performance impact

This content delivery network system provides a comprehensive foundation for global content distribution with features like intelligent caching, geographic routing, content optimization, cache invalidation, and performance analytics while maintaining high performance, security, and reliability standards.