

Ads Click Counter and Aggregator Backend

□ Table of Contents

- Ads Click Counter and Aggregator Backend
 - Requirements Gathering
 - * Functional Requirements
 - * Non-Functional Requirements
 - Traffic Estimation & Capacity Planning
 - * Click Volume Analysis
 - * Storage Calculations
 - * Throughput Requirements
 - Database Schema Design
 - * Click Events Schema
 - * Aggregation Tables Schema
 - * User Demographics Schema
 - System API Design
 - * Click Ingestion APIs
 - * Analytics Query APIs
 - * Aggregation Management APIs
 - High-Level Design (HLD)
 - * System Architecture Overview
 - * Data Flow Architecture
 - * CDN Integration
 - Low-Level Design (LLD)
 - * Click Capture Service
 - * Stream Processing Pipeline
 - * Aggregation Engine
 - Core Algorithms
 - * 1. Click Deduplication Algorithm
 - * 2. Real-time Aggregation Algorithm
 - * 3. Hot Partition Mitigation Algorithm
 - * 4. Data Retention and Archival Algorithm
 - * 5. Fraud Detection Algorithm
 - Performance Optimizations
 - * Partition Strategy
 - * Caching Strategy
 - * Batch Processing Optimization
 - Security Considerations
 - * Click Fraud Prevention
 - * Data Privacy
 - Testing Strategy
 - * Load Testing
 - * Data Consistency Testing
 - Trade-offs and Considerations

- * Consistency vs Availability
- * Real-time vs Batch Processing
- * Storage Cost vs Query Performance

□ [Back to Top](#)

Requirements Gathering

Functional Requirements

Core Click Recording Functionality: - Capture and log advertisement click events in real-time - Record comprehensive metadata including user demographics, ad placement, and context - Support deduplication of duplicate clicks using unique identifiers - Enable complex analytical queries for marketing insights - Provide aggregated metrics across multiple dimensions (time, geography, demographics)

Analytics and Reporting: - Generate real-time dashboards for click performance - Support historical trend analysis and reporting - Enable segmentation by user demographics, ad campaigns, and time periods - Provide conversion funnel analysis capabilities - Support A/B testing metrics and campaign performance comparisons

Data Management: - Handle high-volume click streams with minimal latency - Support long-term data retention (10+ years) with cost optimization - Enable data export capabilities for external analytics tools - Provide data lineage tracking for audit purposes

□ [Back to Top](#)

Non-Functional Requirements

Performance Requirements: - Handle 1 billion clicks per day (peak load: 50,000 clicks/second) - Click logging latency: < 100ms (p99) - Analytics query response time: < 2 seconds for standard reports - Support 99.9% availability with graceful degradation

Scalability Requirements: - Horizontally scalable to handle traffic spikes - Support global deployment across multiple data centers - Auto-scaling capabilities based on traffic patterns - Linear performance scaling with resource addition

Consistency and Durability: - Eventual consistency acceptable for analytics (few minutes delay) - Strong consistency for revenue-critical metrics - Zero data loss for click events (financial impact) - Support for replay capabilities in case of system failures

□ [Back to Top](#)

Traffic Estimation & Capacity Planning

Click Volume Analysis

Daily Traffic Patterns: - 1 billion clicks per day baseline - Peak hours: 3x average traffic (evening hours across time zones) - Average: 11,574 clicks/second - Peak: 50,000 clicks/second during major events or campaigns

Geographic Distribution: - North America: 40% of traffic - Europe: 30% of traffic - Asia-Pacific: 25% of traffic - Other regions: 5% of traffic

Seasonal Variations: - Holiday seasons: 5x normal traffic - Black Friday/Cyber Monday: 10x normal traffic - Regular seasonal fluctuations: $\pm 50\%$ variance

[□ Back to Top](#)

Storage Calculations

Per-Click Data Size: - Basic click metadata: 0.1 KB per click - Extended user demographics: Additional 0.05 KB - Total per click: ~ 0.15 KB

Daily Storage Requirements: - Raw clicks: $1B \times 0.15 \text{ KB} = 150 \text{ GB/day}$ - Aggregated data: $\sim 10 \text{ GB/day}$ (compressed) - Total daily: 160 GB/day

Long-term Storage Projections: - Annual raw data: $160 \text{ GB} \times 365 = \sim 58 \text{ TB/year}$ - 10-year retention: $\sim 580 \text{ TB}$ - With compression and archival: $\sim 200 \text{ TB}$ effective storage

Query Performance Storage: - Hot data (last 30 days): SSD storage for fast queries - Warm data (3-12 months): Standard storage with caching - Cold data (1+ years): Archival storage (S3 Glacier)

[□ Back to Top](#)

Throughput Requirements

Write Throughput: - Peak write: 50,000 writes/second - Average write: 11,574 writes/second - Batch write optimization: 1,000 records per batch

Read Throughput: - Analytics queries: 1,000 concurrent queries - Dashboard updates: 100 real-time connections - Report generation: 50 concurrent heavy queries

Network Bandwidth: - Ingress: $50,000 \times 0.15 \text{ KB} = 7.5 \text{ MB/second}$ peak - Egress: Query results + dashboard data $\sim 10 \text{ MB/second}$ - Internal replication: 2x ingress for redundancy

[□ Back to Top](#)

Database Schema Design

Click Events Schema

Primary Click Events Table: - Click ID (UUID): Unique identifier for deduplication - Ad ID (UUID): Reference to advertisement - User ID (hashed): Privacy-compliant user identifier - Timestamp: Precise click time with timezone - IP Address (hashed): Geographic analysis while preserving privacy - User Agent: Browser and device information - Referrer URL: Source of the click - Click coordinates: X/Y position for heat map analysis

Partitioning Strategy: - Primary partition: Date (daily partitions) - Secondary partition: Geographic region - Enables efficient time-based and geographic queries

□ [Back to Top](#)

Aggregation Tables Schema

Time-based Aggregations: - Minute-level aggregations: Real-time dashboards - Hour-level aggregations: Standard reporting - Daily aggregations: Historical analysis - Monthly aggregations: Long-term trends

Dimensional Aggregations: - By ad campaign: Campaign performance metrics - By geography: Regional performance analysis - By demographics: Audience segment analysis - By device type: Cross-platform performance

Pre-computed Metrics: - Click-through rates (CTR) - Conversion rates - Revenue per click (RPC) - Cost per acquisition (CPA)

□ [Back to Top](#)

User Demographics Schema

User Profile Aggregation: - Age groups: Demographic targeting analysis - Geographic regions: Location-based insights - Device categories: Platform performance - Behavioral segments: Interest-based groupings

Privacy-Compliant Design: - Hashed identifiers only - Aggregated data with k-anonymity - No direct PII storage - GDPR/CCPA compliance built-in

□ [Back to Top](#)

System API Design

Click Ingestion APIs

Bulk Click Ingestion: - High-throughput batch ingestion endpoint - Support for compressed payloads - Idempotent operations with duplicate detection - Asynchronous pro-

cessing with acknowledgment

Real-time Click Capture: - Low-latency single click endpoint - Optimized for CDN edge caching - Minimal payload validation for speed - Fire-and-forget semantics with reliability

Click Validation Service: - Fraud detection integration - Bot traffic filtering - Geographic validation - Rate limiting per user/IP

□ [Back to Top](#)

Analytics Query APIs

Dashboard Data APIs: - Real-time metrics endpoint - Cached aggregation serving - Web-Socket for live updates - Optimized for visualization libraries

Reporting APIs: - Complex query execution engine - Support for custom date ranges - Multi-dimensional filtering - Export capabilities (CSV, JSON)

Trend Analysis APIs: - Time-series data serving - Comparative analysis support - Statistical computation endpoints - Predictive analytics integration

□ [Back to Top](#)

Aggregation Management APIs

Background Processing Control: - Aggregation job scheduling - Processing status monitoring - Error handling and retry logic - Resource utilization tracking

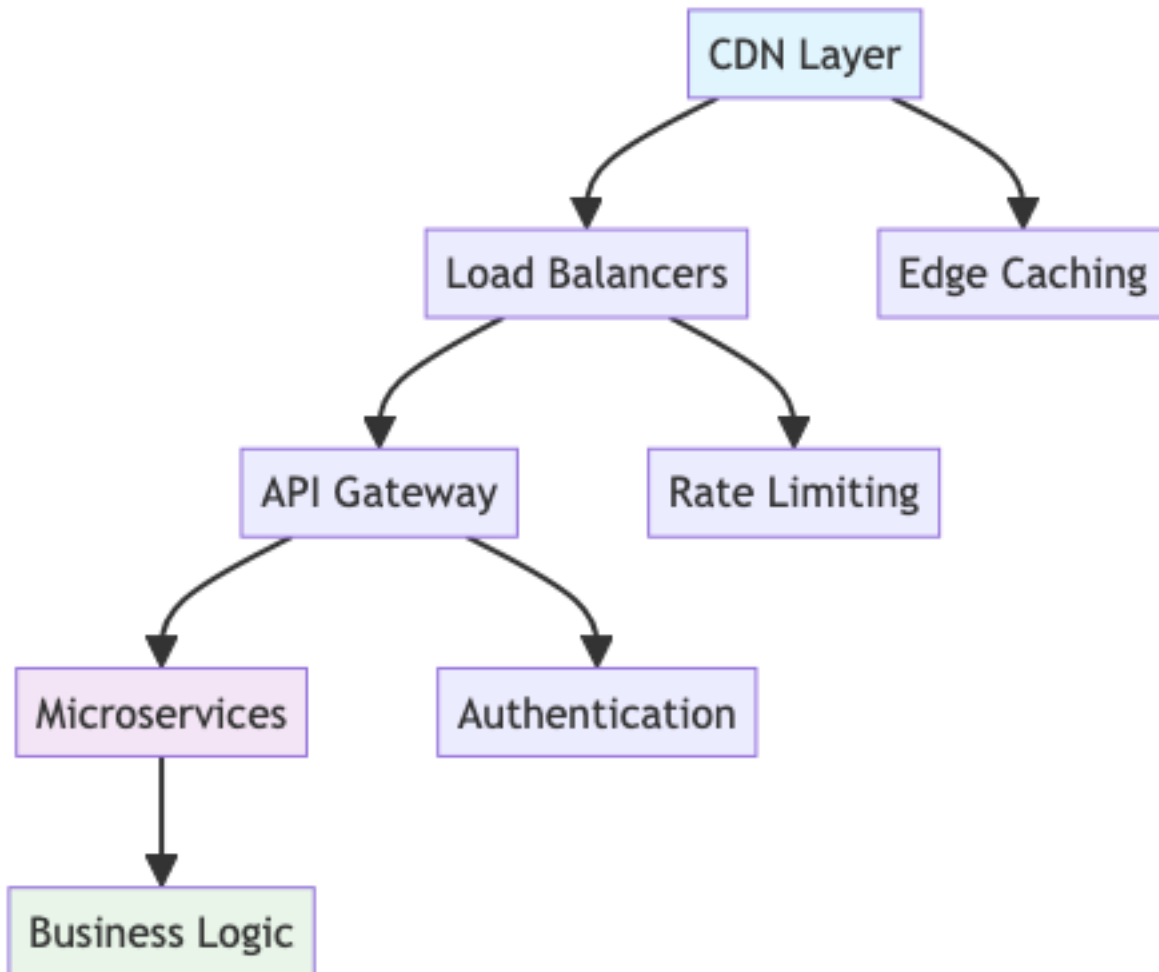
Data Lifecycle Management: - Archive policy configuration - Data retention rule enforcement - Storage tier management - Compliance and audit support

□ [Back to Top](#)

High-Level Design (HLD)

System Architecture Overview

Multi-tier Architecture:




Core Components: - **Click Capture Service:** High-throughput event ingestion - **Stream Processing Engine:** Real-time data transformation - **Aggregation Service:** Batch and real-time metric computation - **Query Engine:** Optimized analytics query processing - **Storage Layer:** Multi-tier data storage with lifecycle management

Geographic Distribution: - Global CDN for click capture with edge processing - Regional data centers for reduced latency - Cross-region replication for disaster recovery - Geo-sharded databases for regulatory compliance

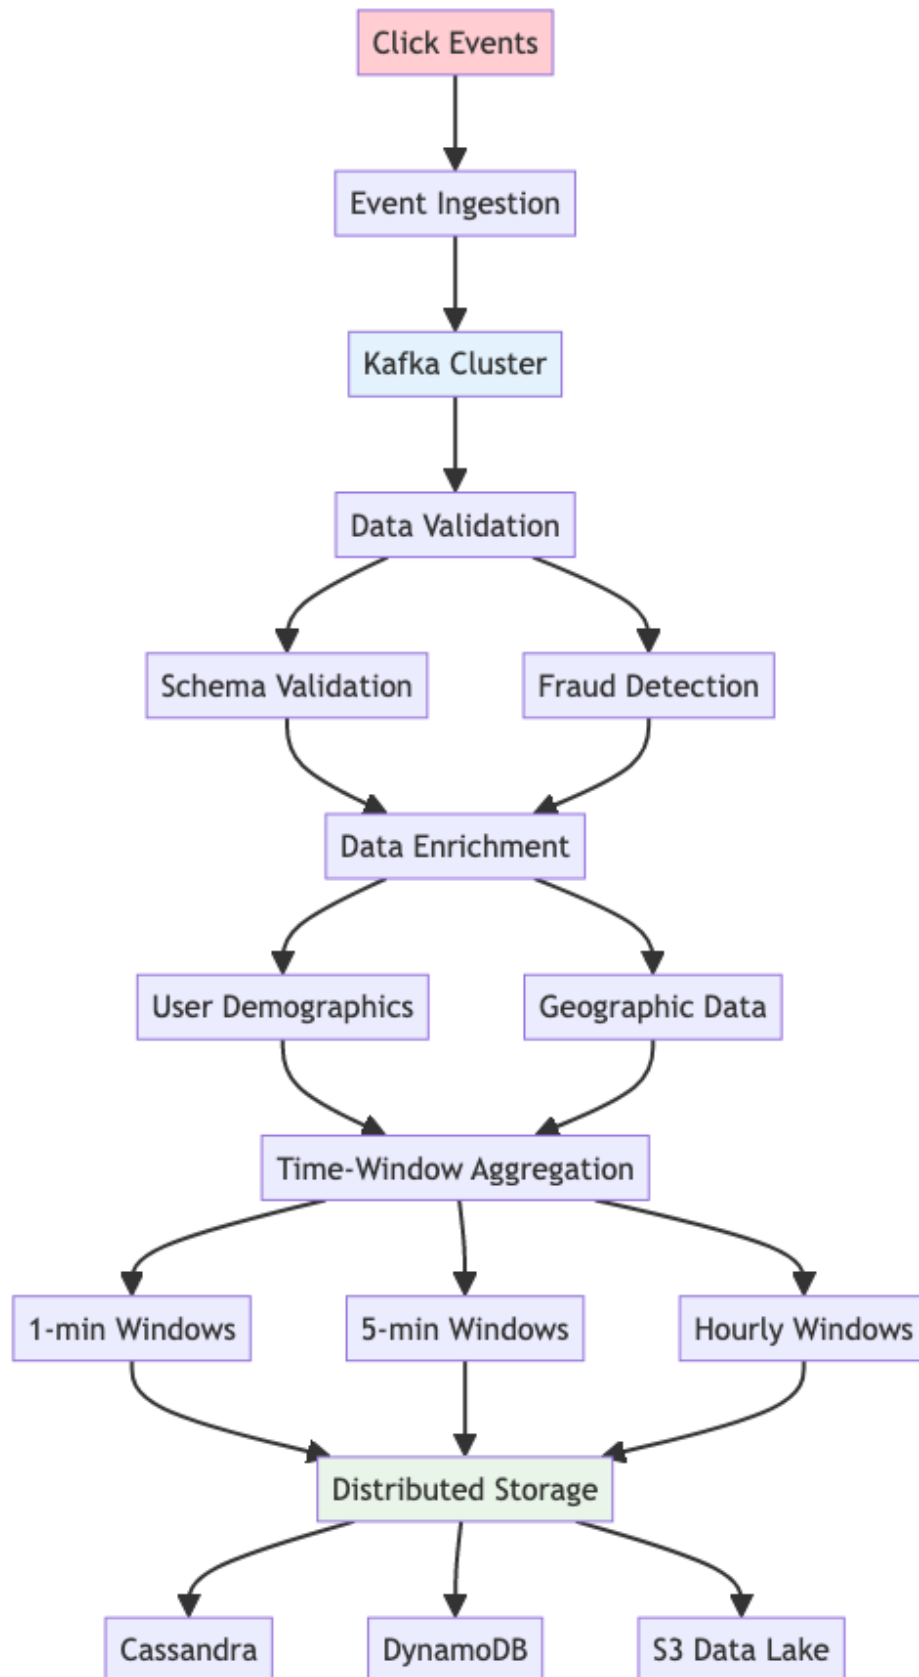
□ [Back to Top](#)

Data Flow Architecture

Lambda Architecture Data Flow:



Stream Processing Pipeline:



Data Lake Architecture: - Raw events stored in object storage (S3/GCS) - Partitioned by date and region for efficient querying - Multiple data formats: Parquet for analytics, JSON for flexibility - Data catalog for schema evolution and discovery

□ [Back to Top](#)

CDN Integration

Edge Processing Capabilities: - Click capture at CDN edge nodes - Basic validation and filtering - Geographic tagging and routing - Reduced latency to sub-50ms globally

Cache Strategy: - Static assets (tracking pixels, JavaScript) cached at edge - Dynamic aggregation results cached with TTL - Geographic distribution based on traffic patterns - Intelligent cache warming for popular queries

□ [Back to Top](#)

Low-Level Design (LLD)

Click Capture Service

High-Throughput Ingestion Design: - Async event processing with message queuing - Connection pooling for database writes - Batch processing for improved throughput - Circuit breaker pattern for fault tolerance

Deduplication Strategy: - Time-window based duplicate detection - Bloom filters for memory-efficient checking - UUID-based unique identification - Configurable deduplication windows

Load Balancing: - Consistent hashing for even distribution - Health check integration - Auto-scaling based on queue depth - Geographic routing for latency optimization

□ [Back to Top](#)

Stream Processing Pipeline

Kafka Configuration: - Topic partitioning by ad campaign and geography - Replication factor of 3 for durability - Configurable retention policies - Producer acknowledgment settings for reliability

Processing Topology: - Parallel processing streams for scalability - Stateful processing for aggregations - Windowing functions for time-based analysis - Error handling with dead letter queues

State Management: - Distributed state stores for stream processing - Checkpoint and recovery mechanisms - State compaction for efficiency - Hot-standby replicas for fast failover

□ [Back to Top](#)

Aggregation Engine

Multi-level Aggregation: - Real-time: Sliding window aggregations (1-minute, 5-minute)
- Near real-time: Tumbling window aggregations (hourly) - Batch: Daily, weekly, monthly aggregations - On-demand: Custom query aggregations

Distributed Computing: - Map-Reduce paradigm for large-scale processing - Spark/Flink for complex analytical computations - Distributed storage for intermediate results - Resource isolation for different workload types

Optimization Techniques: - Pre-aggregation at ingestion time - Materialized views for common queries - Incremental aggregation updates - Parallel processing across multiple dimensions

□ [Back to Top](#)

Core Algorithms

1. Click Deduplication Algorithm

Time-Window Deduplication: - Maintain sliding window of recent clicks (configurable duration) - Use composite key: User ID + Ad ID + Timestamp window - Implement probabilistic data structures (Bloom filters) for memory efficiency - Handle edge cases: clock skew, network delays, retries

UUID-based Approach: - Generate unique identifiers for each ad placement - Client-side UUID generation with server-side validation - Handle UUID collisions with timestamp tiebreakers - Persist deduplication state with configurable TTL

Distributed Deduplication: - Consistent hashing for deduplication service distribution - Shared state across service instances - Eventual consistency model for deduplication decisions - Conflict resolution for concurrent duplicate detection

□ [Back to Top](#)

2. Real-time Aggregation Algorithm

Sliding Window Aggregation: - Maintain multiple overlapping time windows - Incremental updates as new events arrive - Efficient memory management with window expiration - Support for various aggregation functions (sum, count, average, percentiles)

Hierarchical Aggregation: - Bottom-up aggregation: minute □ hour □ day □ month - Incremental computation to avoid reprocessing - Rollup strategies for different granularities - Error propagation and correction mechanisms

Approximate Algorithms: - HyperLogLog for distinct count estimation - Count-Min Sketch for frequency estimation - Reservoir sampling for representative samples - Trade accuracy for performance in high-volume scenarios

[□ Back to Top](#)

3. Hot Partition Mitigation Algorithm

Dynamic Partitioning: - Monitor partition load in real-time - Automatic partition splitting when load exceeds threshold - Consistent hashing with virtual nodes - Load balancing across available partitions

Traffic Shaping: - Rate limiting for individual ad campaigns - Priority queuing for different traffic types - Backpressure handling with queue management - Graceful degradation under extreme load

Consistent Hashing Strategy: - Virtual node mapping for even distribution - Partition rebalancing without service interruption - Replica placement optimization - Hot spot detection and mitigation

[□ Back to Top](#)

4. Data Retention and Archival Algorithm

Lifecycle Management: - Automated data classification by age and access patterns - Tiered storage migration (hot □ warm □ cold) - Compression algorithms for long-term storage - Metadata preservation during archival process

Cost Optimization: - Storage cost analysis and optimization - Intelligent compression based on data characteristics - Deduplication at the storage level - Archive retrieval optimization for occasional access

Compliance and Audit: - Retention policy enforcement - Audit trail for data lifecycle events - Legal hold capabilities for compliance requirements - Secure deletion with verification

[□ Back to Top](#)

5. Fraud Detection Algorithm

Pattern Recognition: - Machine learning models for bot detection - Behavioral analysis for suspicious patterns - Geographic anomaly detection - Time-based pattern analysis

Real-time Scoring: - Risk score calculation for each click - Threshold-based filtering - Adaptive scoring based on historical data - Integration with external fraud detection services

Feedback Loop: - Continuous model training with new data - Human verification integration - False positive reduction strategies - Performance metric tracking and optimization

[□ Back to Top](#)

Performance Optimizations

Partition Strategy

Geographic Partitioning: - Partition data by geographic regions - Optimize for local query patterns - Reduce cross-partition queries - Enable regulatory compliance (data residency)

Time-based Partitioning: - Daily/weekly partitions for time-series data - Partition pruning for historical queries - Parallel processing across time ranges - Efficient archival of old partitions

Hybrid Partitioning: - Multi-dimensional partitioning strategy - Composite partitioning keys - Query pattern optimization - Dynamic repartitioning based on data distribution

□ [Back to Top](#)

Caching Strategy

Multi-level Caching: - L1: Application-level caching for frequently accessed data - L2: Distributed cache (Redis/Memcached) for shared data - L3: CDN caching for static and semi-static content - Cache coherency strategies across levels

Intelligent Cache Warming: - Predictive cache warming based on usage patterns - Priority-based cache population - Cache hit ratio optimization - Memory usage optimization

Cache Invalidation: - Event-driven cache invalidation - TTL-based expiration strategies - Write-through and write-behind patterns - Consistency guarantees across cache layers

□ [Back to Top](#)

Batch Processing Optimization

Parallel Processing: - Data parallelism across processing nodes - Task parallelism for different aggregation types - Resource isolation for batch vs. real-time workloads - Dynamic resource allocation based on workload

I/O Optimization: - Columnar storage formats (Parquet, ORC) - Compression optimization for storage and transfer - Batch size optimization for throughput - Memory mapping for large dataset processing

Processing Pipeline: - ETL pipeline optimization - Incremental processing for efficiency - Error handling and retry mechanisms - Monitoring and alerting for batch jobs

□ [Back to Top](#)

Security Considerations

Click Fraud Prevention

Multi-layered Detection: - Rate limiting per IP address and user - Pattern analysis for automated traffic - Device fingerprinting for bot detection - Geographic and temporal anomaly detection

Machine Learning Integration: - Real-time fraud scoring models - Behavioral analysis algorithms - Adaptive threshold adjustment - Continuous model retraining

Third-party Integration: - External fraud detection services - IP reputation databases - Device intelligence platforms - Threat intelligence feeds

□ [Back to Top](#)

Data Privacy

Privacy by Design: - Data minimization principles - Purpose limitation enforcement - Anonymization and pseudonymization - Consent management integration

Compliance Framework: - GDPR compliance for European users - CCPA compliance for California residents - Data subject rights implementation - Privacy impact assessments

Encryption and Security: - End-to-end encryption for sensitive data - Encryption at rest and in transit - Key management and rotation - Access control and audit logging

□ [Back to Top](#)

Testing Strategy

Load Testing

Performance Testing: - Gradual load increase testing - Peak load simulation - Sustained load testing - Breaking point identification

Stress Testing: - System behavior under extreme conditions - Resource exhaustion scenarios - Recovery testing after failures - Cascading failure prevention

Scalability Testing: - Horizontal scaling validation - Auto-scaling behavior verification - Resource utilization optimization - Performance regression testing

□ [Back to Top](#)

Data Consistency Testing

Eventual Consistency Validation: - Cross-replica consistency checking - Convergence time measurement - Conflict resolution validation - Data integrity verification

Aggregation Accuracy: - End-to-end data flow validation - Aggregation result verification
- Temporal consistency checking - Cross-dimensional consistency

Disaster Recovery Testing: - Data loss prevention validation - Recovery time objective (RTO) testing - Recovery point objective (RPO) validation - Failover and failback procedures

□ [Back to Top](#)

Trade-offs and Considerations

Consistency vs Availability

CAP Theorem Application: - Prioritize availability for click ingestion - Accept eventual consistency for analytics - Strong consistency for revenue-critical operations - Partition tolerance through replication

Consistency Models: - Strong consistency for financial transactions - Eventual consistency for analytics aggregations - Session consistency for user-facing operations - Monotonic consistency for time-series data

□ [Back to Top](#)

Real-time vs Batch Processing

Lambda Architecture Benefits: - Real-time insights for immediate decision making - Batch processing for complex analytics - Cost optimization through appropriate processing choice - Flexibility in handling different data access patterns

Processing Trade-offs: - Real-time: Higher cost, lower latency, approximate results
- Batch: Lower cost, higher latency, exact results - Hybrid approach for optimal cost-performance balance

□ [Back to Top](#)

Storage Cost vs Query Performance

Storage Optimization: - Hot data on fast storage for quick access - Cold data on cheap storage with slower access - Compression for long-term storage cost reduction - Intelligent data lifecycle management

Query Performance: - Indexing strategies for common query patterns - Materialized views for complex aggregations - Caching for frequently accessed data - Query optimization through data modeling

Cost Management: - Storage tier optimization based on access patterns - Compute resource optimization for batch processing - Network cost reduction through geographic distribution - Monitoring and alerting for cost anomalies

□ [Back to Top](#)