

# Privacy & Visibility Controls Backend

## □ Table of Contents

- Privacy & Visibility Controls Backend
  - Requirements Gathering
    - \* Functional Requirements
    - \* Non-Functional Requirements
  - Traffic Estimation & Capacity Planning
    - \* User Activity Analysis
    - \* Permission Check Volume
    - \* Privacy Setting Updates
  - Database Schema Design
    - \* User Profile Schema
    - \* Privacy Settings Schema
    - \* Relationship Graph Schema
  - System API Design
    - \* Privacy Management APIs
    - \* Profile Access APIs
    - \* Relationship Management APIs
  - High-Level Design (HLD)
    - \* System Architecture Overview
    - \* Privacy Control Data Flow
    - \* Authorization Pipeline
  - Low-Level Design (LLD)
    - \* Permissions Service
    - \* Viewer Service
    - \* Relationship Resolution Engine
  - Core Algorithms
    - \* 1. Privacy Resolution Algorithm
    - \* 2. Friend-of-Friend Visibility Algorithm
    - \* 3. Custom Group Access Algorithm
    - \* 4. Batch Permission Update Algorithm
    - \* 5. Privacy Inheritance Algorithm
  - Performance Optimizations
    - \* Caching Strategy
    - \* Pre-computation Optimization
    - \* Database Query Optimization
  - Security Considerations
    - \* Backend Authorization
    - \* Data Leak Prevention
  - Testing Strategy
    - \* Privacy Rule Testing
    - \* Performance Testing
  - Trade-offs and Considerations

- \* Privacy Granularity vs Performance
- \* Real-time vs Cached Permissions
- \* SQL vs NoSQL for Relationship Data

□ [Back to Top](#)

---

## Requirements Gathering

### Functional Requirements

**Core Privacy Control Functionality:** - Enable users to set granular privacy settings for profile fields (birthday, education, photos, etc.) - Support multiple privacy levels: Public, Friends Only, Custom Groups, Only Me - Allow creation and management of custom privacy groups - Provide real-time privacy filtering for profile views and content access - Support bulk privacy updates and inheritance rules

**Profile Visibility Management:** - Control visibility of different profile sections based on viewer relationship - Enable conditional visibility (friends-of-friends, mutual connections) - Support time-based privacy settings (temporary visibility changes) - Provide privacy preview functionality for users to test their settings - Handle blocked user access restrictions

**Relationship-based Access Control:** - Enforce privacy based on friendship status and connection levels - Support custom friend lists and categorization - Handle pending friend requests and limited visibility states - Manage group membership and hierarchical access rights - Enable granular control over mutual friend visibility

**Administrative Features:** - Provide privacy audit trails and access logs - Support privacy policy enforcement and compliance - Enable bulk privacy administration for organizational accounts - Offer privacy analytics and usage insights

□ [Back to Top](#)

### Non-Functional Requirements

**Performance Requirements:** - Support 1 billion user profiles with complex privacy settings - Handle 100 million profile views per hour with real-time privacy filtering - Permission check latency: < 50ms (p99) - Privacy setting updates propagated within 1 second globally - System availability: 99.99% uptime with zero privacy leaks

**Scalability Requirements:** - Horizontally scalable to handle growing user base and relationships - Support for complex relationship graphs with millions of connections per user - Auto-scaling based on privacy check volume and geographic distribution - Linear performance scaling with increased privacy rule complexity

**Security and Privacy:** - Zero tolerance for privacy setting bypass or data leaks - End-to-end encryption for sensitive privacy configuration data - Audit logging for all privacy-related access and changes - GDPR and privacy regulation compliance built-in

**Consistency Requirements:** - Strong consistency for privacy setting updates - Eventual consistency acceptable for relationship graph updates (< 1 second) - Atomic privacy rule changes across multiple profile fields - Consistent privacy enforcement across all system components

□ [Back to Top](#)

---

## Traffic Estimation & Capacity Planning

### User Activity Analysis

**User Base Scale:** - 1 billion registered users with active privacy settings - 500 million daily active users accessing profiles - Average 50 profile fields per user with individual privacy settings - 10 privacy groups per user on average

**Privacy Setting Distribution:** - Public settings: 30% of profile fields - Friends only: 50% of profile fields

- Custom groups: 15% of profile fields - Only me: 5% of profile fields

**Relationship Complexity:** - Average 200 friends per user - Average 5 custom privacy groups per user - Friend-of-friend connections: 20,000 per user average - Group memberships: Average 10 groups per user

□ [Back to Top](#)

### Permission Check Volume

**Read Operations:** - Profile views: 100 million per hour - Permission checks per view: Average 20 fields - Total permission checks: 2 billion per hour (555,000 per second) - Friend relationship lookups: 50 million per hour

**Geographic Distribution:** - North America: 35% of traffic - Europe: 25% of traffic - Asia-Pacific: 30% of traffic - Other regions: 10% of traffic

**Peak Load Scenarios:** - Holiday seasons: 3x normal traffic - Viral content events: 5x normal profile access - Privacy policy updates: 10x permission check volume

□ [Back to Top](#)

### Privacy Setting Updates

**Write Operations:** - Privacy setting updates: 1 million per hour - Bulk privacy changes: 100,000 per hour - Friend relationship changes: 5 million per hour - Group membership updates: 500,000 per hour

**Update Patterns:** - Individual field updates: 70% of changes - Bulk profile updates: 20% of changes - Import/migration operations: 10% of changes

**Propagation Requirements:** - Real-time updates: Privacy setting changes - Near real-time: Relationship changes (< 1 second) - Batch updates: Group membership bulk operations - Background processing: Privacy audit and analytics

□ [Back to Top](#)

---

## Database Schema Design

### User Profile Schema

**Profile Fields Table:** - User ID (Primary Key): Unique user identifier - Field Name: Profile field identifier (birthday, education, etc.) - Field Value: Actual profile data (encrypted for sensitive fields) - Privacy Level: Current privacy setting for this field - Custom Group ID: Reference to custom privacy group if applicable - Last Updated: Timestamp of last privacy or value change

**Field Categories:** - Basic Information: Name, profile picture, cover photo - Contact Information: Email, phone, address - Personal Details: Birthday, relationship status, hometown - Professional Information: Work history, education, skills - Preferences: Interests, favorite content, activities

**Data Fragmentation Strategy:** - Each profile field stored as separate record for granular privacy control - Partition by User ID for efficient retrieval - Secondary indexes on privacy levels for administrative queries - Encryption at rest for PII and sensitive data

□ [Back to Top](#)

### Privacy Settings Schema

**Privacy Rules Table:** - Rule ID (Primary Key): Unique privacy rule identifier - User ID (Foreign Key): Owner of the privacy rule - Field Pattern: Regex or field names this rule applies to - Privacy Level: Public, Friends, Custom Group, Only Me - Group References: List of custom group IDs for custom privacy - Conditions: Additional conditions (time-based, location-based)

**Custom Privacy Groups:** - Group ID (Primary Key): Unique group identifier - User ID (Foreign Key): Group owner - Group Name: User-defined group name - Group Description: Optional group description - Member List: List of user IDs in this privacy group - Created Date: Group creation timestamp

**Privacy Templates:** - Template ID: Predefined privacy setting templates - Template Name: "Conservative", "Open", "Professional", etc. - Field Settings: Default privacy levels for each field type - User Adoption: Tracking of template usage for optimization

□ [Back to Top](#)

## Relationship Graph Schema

**Friendship Relations:** - Relationship ID (Primary Key): Unique relationship identifier - User A ID: First user in the relationship - User B ID: Second user in the relationship - Relationship Type: Friend, Family, Colleague, Blocked, etc. - Status: Confirmed, Pending, Rejected - Created Date: When relationship was established - Mutual Friend Count: Cached count for efficiency

**Connection Levels:** - User ID (Composite Key Part 1): Source user - Target User ID (Composite Key Part 2): Target user - Connection Level: Direct (1), Friend-of-friend (2), Extended (3+) - Path Strength: Weighted strength of connection - Last Computed: When this connection level was calculated - Invalidation Flag: Marks stale data for recomputation

**Adjacency List Optimization:** - User ID (Primary Key): User identifier - Direct Friends: Serialized list of direct friend IDs - Friend Groups: Organized friend lists by categories - Blocked Users: List of blocked user IDs - Cached Metrics: Friend counts, mutual connections stats

□ [Back to Top](#)

---

## System API Design

### Privacy Management APIs

**Privacy Setting Configuration:** - Set field-level privacy settings with custom group support - Bulk privacy updates across multiple profile fields - Privacy template application and customization - Privacy rule validation and conflict resolution

**Custom Group Management:** - Create, update, and delete custom privacy groups - Add/remove members from privacy groups - Group sharing and collaboration features - Group-based permission inheritance

**Privacy Testing and Preview:** - Preview profile appearance for different viewer types - Privacy setting validation and impact analysis - Bulk privacy change preview before application - Privacy recommendation engine integration

□ [Back to Top](#)

### Profile Access APIs

**Privacy-Filtered Profile Retrieval:** - Get user profile with privacy filtering applied - Batch profile retrieval with relationship context - Field-level access checking with detailed reasoning - Profile section access with conditional visibility

**Relationship-Aware Data Access:** - Friend list access with privacy filtering - Mutual connection discovery with privacy respect - Activity feed filtering based on privacy settings - Search result filtering with privacy consideration

**Access Logging and Audit:** - Profile access logging for privacy compliance - Privacy violation detection and alerting - Access pattern analysis for security monitoring - Audit trail generation for regulatory compliance

□ [Back to Top](#)

## Relationship Management APIs

**Friend Connection Management:** - Send, accept, and reject friend requests - Friend categorization and custom list management - Relationship strength calculation and caching - Mutual friend discovery and recommendation

**Connection Graph Queries:** - Friend-of-friend path discovery with privacy filtering - Connection strength analysis and ranking - Network analysis for privacy recommendation - Relationship graph traversal with access control

**Bulk Relationship Operations:** - Import friend lists from external platforms - Bulk friend categorization and privacy application - Relationship cleanup and inactive account handling - Privacy inheritance from relationship changes

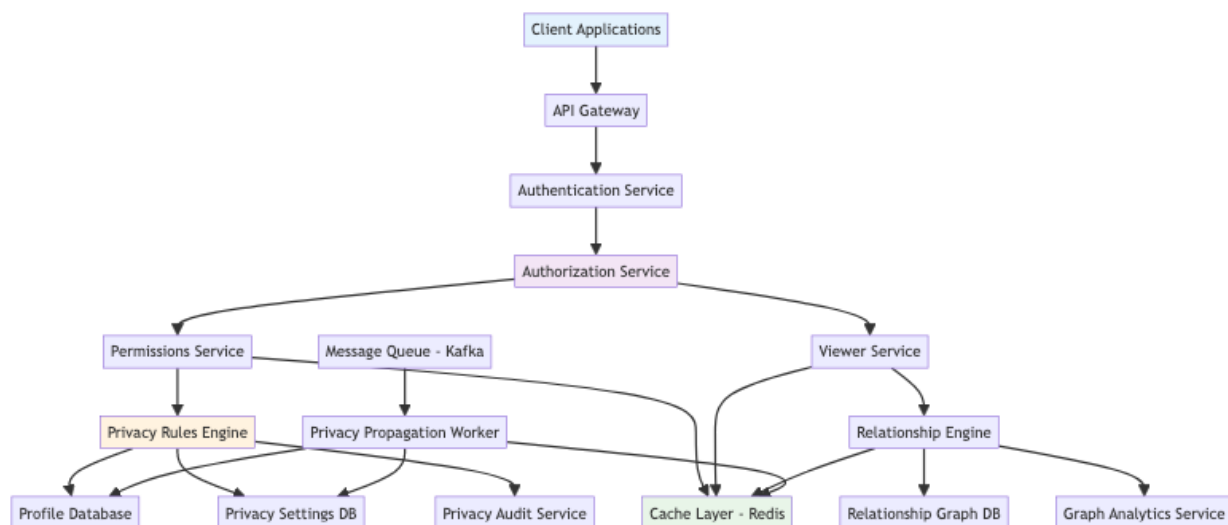
□ [Back to Top](#)

---

## High-Level Design (HLD)

### System Architecture Overview

#### Microservices Privacy Architecture:




**Core Components:** - **Authentication Service:** User identity verification and session management - **Authorization Service:** High-level access control and permission orchestration - **Permissions Service:** Privacy rule evaluation and field-level access control -

**Viewer Service:** Profile viewing with relationship-aware filtering - **Relationship Engine:** Friend graph traversal and connection analysis - **Privacy Rules Engine:** Complex privacy logic and custom group resolution

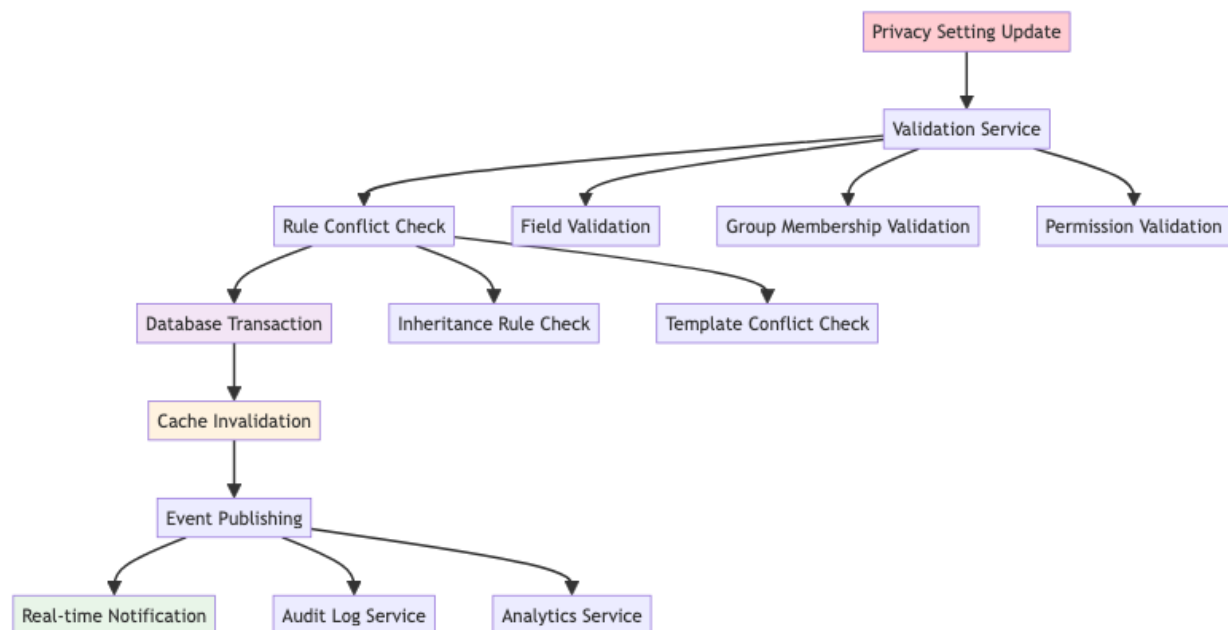
□ Back to Top

## Privacy Control Data Flow

### Profile Access with Privacy Filtering:



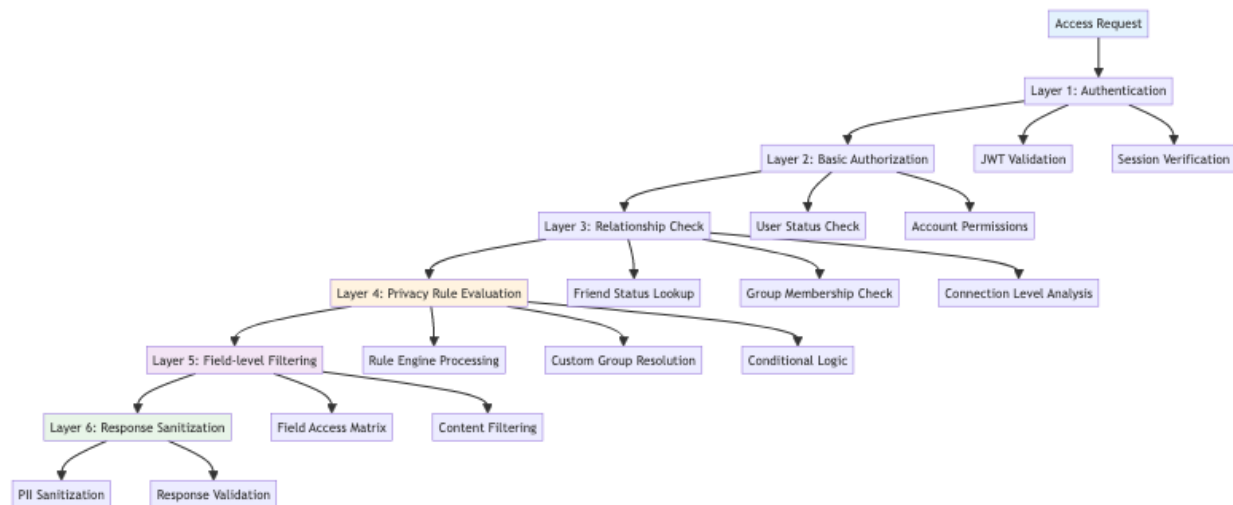
### Privacy Setting Update Flow:



□ Back to Top

## Authorization Pipeline

### Multi-Layer Authorization Architecture:



□ [Back to Top](#)

## Low-Level Design (LLD)

### Permissions Service

**Privacy Rule Engine Design:** - **Rule Parser:** Interprets complex privacy rules with conditions and inheritance - **Context Builder:** Assembles viewer context including relationships and group memberships - **Decision Engine:** Evaluates rules against context using configurable rule precedence - **Result Caching:** Intelligent caching of permission decisions with invalidation strategies

**Custom Group Resolution:** - **Group Membership Cache:** In-memory cache of user group memberships for fast lookup - **Hierarchical Groups:** Support for nested groups and inheritance relationships - **Dynamic Groups:** Groups based on computed criteria (mutual friends, location, etc.) - **Group Validation:** Ensures group membership consistency and prevents circular dependencies

**Performance Optimization:** - **Batch Processing:** Bulk permission evaluation for multiple fields or users - **Parallel Evaluation:** Concurrent processing of independent privacy rules - **Rule Compilation:** Pre-compiled rule structures for faster evaluation - **Memory Pooling:** Efficient memory management for high-frequency operations

□ [Back to Top](#)

### Viewer Service

**Profile Filtering Engine:** - **Field-level Filtering:** Granular control over individual profile field visibility - **Contextual Filtering:** Dynamic filtering based on viewer relationship and context - **Bulk Profile Processing:** Efficient filtering of multiple profiles in single request - **Progressive Disclosure:** Staged revelation of information based on interaction level



**Relationship Context Assembly:** - **Connection Graph Traversal:** Efficient graph algorithms for relationship discovery - **Relationship Strength Calculation:** Weighted relationship scoring for nuanced privacy - **Mutual Connection Analysis:** Fast mutual friend discovery and intersection operations - **Social Graph Caching:** Strategic caching of relationship data for performance

**Response Optimization:** - **Selective Field Loading:** Load only fields that pass privacy checks - **Response Compression:** Efficient serialization of filtered profile data - **Streaming Responses:** Progressive loading for large profile datasets - **Error Handling:** Graceful degradation when privacy checks fail

□ [Back to Top](#)

## Relationship Resolution Engine

**Graph Traversal Optimization:** - **Adjacency List Caching:** Memory-optimized representation of user connections - **Path Finding Algorithms:** BFS/DFS optimization for friend-of-friend discovery - **Graph Partitioning:** Distributed graph storage for scalability - **Connection Strength Metrics:** Real-time calculation of relationship weights

**Relationship State Management:** - **State Synchronization:** Consistent relationship state across distributed services - **Event-Driven Updates:** Real-time relationship change propagation - **Conflict Resolution:** Handling simultaneous relationship state changes - **Garbage Collection:** Cleanup of stale relationship data and cached connections

**Scalability Architecture:** - **Sharded Graph Storage:** Distributed relationship data across multiple nodes - **Replication Strategy:** Multi-region relationship data replication - **Load Balancing:** Intelligent routing of relationship queries - **Hot Spot Mitigation:** Handling popular users with massive connection counts

□ [Back to Top](#)

---

## Core Algorithms

### 1. Privacy Resolution Algorithm

**Multi-Stage Privacy Decision Tree:** - Evaluate user authentication and basic access permissions - Resolve viewer-target relationship status and connection level - Look up applicable privacy rules based on field and relationship context - Process custom group memberships and hierarchical permissions - Apply conditional logic (time-based, location-based, interaction-based rules) - Generate final access decision with detailed reasoning for audit

**Rule Precedence and Conflict Resolution:** - Explicit deny rules override all allow rules (security-first approach) - User-specific rules override group-based or template rules -

More restrictive settings take precedence in conflicts - Recent rule changes override older conflicting rules - Administrative override capabilities for compliance scenarios

**Optimization Strategies:** - Short-circuit evaluation for common deny scenarios - Rule compilation and caching for frequently accessed patterns - Batch evaluation for multiple fields with shared context - Probabilistic early termination for complex rule sets

□ [Back to Top](#)

## 2. Friend-of-Friend Visibility Algorithm

**Multi-Hop Connection Discovery:** - Build adjacency lists for each user with direct connections - Implement bidirectional BFS for efficient path discovery - Cache intermediate results for common connection patterns - Apply privacy filtering at each hop to respect friend visibility settings - Limit search depth to prevent performance degradation

**Connection Strength Calculation:** - Weight edges based on interaction frequency and recency - Consider mutual friend counts and shared group memberships - Apply decay functions for inactive or weak connections - Normalize scores across different user activity levels - Update strengths asynchronously to maintain performance

**Fan-out Optimization:** - Pre-compute friend-of-friend lists for active users - Use probabilistic data structures for large-scale approximations - Implement incremental updates when friendship networks change - Balance accuracy vs. performance based on use case requirements - Provide fallback to real-time computation for cache misses

□ [Back to Top](#)

## 3. Custom Group Access Algorithm

**Dynamic Group Membership Resolution:** - Support static membership lists with explicit user inclusion - Implement computed groups based on relationship criteria - Handle nested groups with inheritance and override rules - Manage temporal group memberships with expiration dates - Provide group template systems for common access patterns

**Group-based Permission Inheritance:** - Resolve group hierarchies with proper precedence ordering - Handle circular dependency detection and prevention - Implement permission aggregation across multiple group memberships - Support group-specific overrides for individual members - Maintain audit trails for group-based access decisions

**Scalability and Performance:** - Cache group membership data with smart invalidation - Implement lazy loading for large groups with thousands of members - Use bloom filters for fast negative membership checks - Batch group resolution for multiple access checks - Optimize storage for sparse group membership matrices

□ [Back to Top](#)

## 4. Batch Permission Update Algorithm

**Atomic Privacy Rule Changes:** - Group related privacy changes into atomic transactions - Implement two-phase commit for distributed privacy updates - Provide rollback capabilities for failed batch operations - Ensure consistency across multiple database partitions - Handle partial failures with appropriate error recovery

**Propagation and Cache Invalidation:** - Identify all affected cached data from privacy rule changes - Implement intelligent cache invalidation strategies - Use message queues for asynchronous propagation - Provide progress tracking for long-running batch operations - Ensure global consistency with eventual propagation guarantees

**Performance Optimization:** - Batch similar privacy changes for efficient database operations - Implement priority queues for urgent privacy updates - Use parallel processing for independent privacy rule changes - Optimize database queries with proper indexing strategies - Provide real-time progress feedback for user experience

□ [Back to Top](#)

## 5. Privacy Inheritance Algorithm

**Template-based Privacy Configuration:** - Provide predefined privacy templates for common use cases - Allow customization of templates while maintaining base structure - Implement template versioning for updates and rollbacks - Support organization-wide templates for corporate accounts - Enable smart recommendations based on user behavior patterns

**Field-level Inheritance Rules:** - Support inheritance from profile sections to individual fields - Implement override capabilities for specific field requirements - Handle inheritance conflicts with clear precedence rules - Provide bulk inheritance application with selective overrides - Maintain inheritance history for audit and troubleshooting

**Dynamic Inheritance Updates:** - Automatically apply template updates to derived privacy settings - Allow users to opt-out of automatic inheritance updates - Implement cascading updates with impact analysis - Provide preview capabilities before applying inheritance changes - Handle inheritance conflicts with user notification and choice

□ [Back to Top](#)

---

## Performance Optimizations

### Caching Strategy

**Multi-Level Caching Architecture:** - **L1 Cache:** Application-level caching for hot privacy rules and relationships - **L2 Cache:** Distributed cache (Redis) for cross-service privacy data sharing - **L3 Cache:** Database query result caching for complex relationship queries - **L4 Cache:** CDN caching for public profile data and static privacy templates

**Intelligent Cache Management:** - **Smart Invalidation:** Selective cache invalidation based on privacy change impact - **Predictive Warming:** Pre-populate cache with likely-accessed privacy data - **Hierarchical Expiration:** Different TTL values based on data volatility - **Cache Coherency:** Maintain consistency across distributed cache layers

**Performance Metrics:** - Cache hit ratios > 95% for privacy rule lookups - Sub-10ms cache response times for critical privacy decisions - Automatic cache sizing based on usage patterns - Real-time cache performance monitoring and alerting

□ [Back to Top](#)

## Pre-computation Optimization

**Friend-of-Friend Pre-computation:** - **Batch Processing:** Nightly computation of friend-of-friend relationships - **Incremental Updates:** Real-time updates for new friend connections - **Materialized Views:** Pre-computed privacy-filtered friend lists - **Connection Strength:** Pre-calculated relationship weights and scores

**Privacy Rule Compilation:** - **Rule Optimization:** Compile complex privacy rules into efficient decision trees - **Pattern Recognition:** Identify common privacy patterns for optimization - **Index Generation:** Create specialized indexes for frequent privacy queries - **Result Memoization:** Cache privacy decisions for repeated access patterns

**Graph Analytics:** - **Community Detection:** Identify user clusters for optimized privacy grouping - **Influence Scoring:** Calculate user influence for privacy recommendation - **Anomaly Detection:** Identify unusual privacy access patterns - **Trend Analysis:** Track privacy setting changes for recommendation improvement

□ [Back to Top](#)

## Database Query Optimization

**Partition Strategy:** - **User-based Partitioning:** Partition privacy data by user ID for locality - **Temporal Partitioning:** Separate current and historical privacy data - **Relationship Sharding:** Distribute relationship graph across multiple shards - **Hot-Cold Data Separation:** Optimize storage for frequently vs. rarely accessed data

**Index Optimization:** - **Composite Indexes:** Multi-field indexes for complex privacy queries - **Partial Indexes:** Indexes on subset of data based on common query patterns - **Covering Indexes:** Include all needed fields to avoid table lookups - **Adaptive Indexing:** Dynamic index creation based on query patterns

**Query Performance:** - **Query Plan Optimization:** Analyze and optimize expensive privacy queries - **Parallel Execution:** Concurrent processing of independent privacy checks - **Connection Pooling:** Efficient database connection management - **Read Replicas:** Distribute read queries across multiple database replicas

□ [Back to Top](#)

---

## Security Considerations

### Backend Authorization

**Defense in Depth:** - **Multiple Authorization Layers:** Authentication, authorization, privacy, and field-level checks - **Principle of Least Privilege:** Grant minimum necessary access for each request - **Zero Trust Architecture:** Verify every access request regardless of source - **Continuous Validation:** Re-validate permissions for long-running sessions

**Authorization Enforcement:** - **Server-side Filtering:** All privacy filtering performed on backend to prevent bypass - **API Gateway Controls:** Centralized authorization enforcement at entry points - **Service Mesh Security:** Inter-service communication with mutual TLS and authorization - **Database-level Security:** Row-level security and column encryption for sensitive data

**Audit and Monitoring:** - **Access Logging:** Comprehensive logging of all privacy-related access attempts - **Anomaly Detection:** Machine learning-based detection of unusual access patterns - **Real-time Alerting:** Immediate alerts for potential privacy violations - **Compliance Reporting:** Automated reports for regulatory compliance requirements

□ [Back to Top](#)

### Data Leak Prevention

**Frontend Security:** - **Data Minimization:** Send only approved data to frontend applications - **Response Sanitization:** Remove sensitive metadata from API responses - **Client-side Validation:** Prevent UI manipulation from bypassing privacy controls - **Session Management:** Secure session handling with automatic expiration

**API Security:** - **Rate Limiting:** Prevent privacy probing through automated requests - **Request Validation:** Strict validation of all privacy-related API requests - **Error Message Security:** Avoid leaking privacy information through error responses - **API Versioning:** Maintain security across different API versions

**Data Protection:** - **Encryption at Rest:** Encrypt sensitive privacy data in databases - **Encryption in Transit:** TLS encryption for all privacy-related communications - **Key Management:** Secure key storage and rotation for encryption systems - **Data Anonymization:** Remove or pseudonymize personal identifiers where possible

□ [Back to Top](#)

---

## Testing Strategy

### Privacy Rule Testing

**Comprehensive Test Coverage:** - **Unit Tests:** Individual privacy rule evaluation with edge cases - **Integration Tests:** End-to-end privacy filtering across multiple services - **Regression Tests:** Ensure privacy rule changes don't break existing functionality - **Performance Tests:** Validate privacy rule evaluation performance under load

**Privacy Scenario Testing:** - **Relationship Testing:** Test all combinations of friendship and relationship states - **Group Membership Testing:** Validate custom group access across different scenarios - **Inheritance Testing:** Ensure proper privacy rule inheritance and override behavior - **Edge Case Testing:** Handle unusual relationship configurations and privacy settings

**Compliance Testing:** - **Regulatory Testing:** Validate compliance with GDPR, CCPA, and other privacy regulations - **Policy Testing:** Ensure privacy policies are correctly implemented and enforced - **Audit Testing:** Verify audit trail completeness and accuracy - **Security Testing:** Penetration testing for privacy bypass vulnerabilities

□ [Back to Top](#)

### Performance Testing

**Load Testing:** - **Privacy Check Volume:** Test system under peak privacy evaluation loads - **Concurrent Users:** Validate performance with millions of simultaneous users - **Database Performance:** Test privacy database queries under heavy load - **Cache Performance:** Validate cache effectiveness under various load patterns

**Stress Testing:** - **Resource Exhaustion:** Test system behavior when resources are fully utilized - **Cascade Failures:** Validate graceful degradation when components fail - **Recovery Testing:** Test system recovery after privacy service outages - **Data Consistency:** Ensure privacy consistency during high-stress scenarios

**Scalability Testing:** - **Horizontal Scaling:** Validate privacy system scaling across multiple nodes - **Database Scaling:** Test privacy database scaling strategies - **Cache Scaling:** Validate distributed cache scaling for privacy data - **Geographic Distribution:** Test privacy consistency across global deployments

□ [Back to Top](#)

---

## Trade-offs and Considerations

### Privacy Granularity vs Performance

**Granularity Benefits:** - Enhanced user control over personal information sharing - Compliance with strict privacy regulations and user expectations - Competitive advantage

through superior privacy features - Reduced privacy-related support requests through clear controls

**Performance Impact:** - Increased complexity in privacy rule evaluation and caching - Higher database storage requirements for detailed privacy settings - More complex cache invalidation strategies - Potential latency increase for profile access operations

**Optimization Strategies:** - Intelligent defaults that cover majority of use cases - Performance-aware privacy rule design and evaluation - Strategic pre-computation of common privacy scenarios - Tiered privacy levels balancing granularity with performance

□ [Back to Top](#)

## Real-time vs Cached Permissions

**Real-time Benefits:** - Immediate privacy setting updates without propagation delays - Guaranteed consistency for privacy-critical operations - Accurate permission evaluation for edge cases - Complete audit trail for real-time access decisions

**Caching Advantages:** - Significantly improved response times for profile access - Reduced database load and improved overall system scalability - Better user experience through faster page loads - Cost savings through reduced computational requirements

**Hybrid Approach:** - Real-time evaluation for privacy-critical operations - Cached evaluation for performance-critical read operations - Smart cache invalidation based on privacy change impact - Fallback to real-time when cache inconsistency is detected

□ [Back to Top](#)

## SQL vs NoSQL for Relationship Data

**SQL Database Benefits:** - ACID transactions for consistent relationship and privacy data - Complex JOIN operations for relationship graph queries - Mature indexing and query optimization capabilities - Strong consistency guarantees for critical privacy operations

**NoSQL Advantages:** - Horizontal scaling for massive relationship graphs - Flexible schema for evolving privacy and relationship models - Better performance for simple relationship lookups - Natural fit for graph-like relationship data structures

**Hybrid Architecture:** - SQL for transactional privacy settings and user data - Graph databases (Neo4j) for complex relationship traversal - NoSQL (DynamoDB) for high-volume relationship caching - Event streaming (Kafka) for maintaining consistency across storage systems

**Technology Selection:** - PostgreSQL for privacy settings with JSONB for flexible rules - Neo4j for complex relationship graph analysis and traversal - Redis for high-performance caching of privacy decisions - Cassandra for time-series privacy audit logs and analytics

□ [Back to Top](#)